

AWS CSAA Practice Tests

[Home](#) / [My courses](#) / [AWS CSAA Practice Tests](#) / [FULL TEST\(S\)](#) / [New Practice Test V](#)

Started on Friday, 9 February 2018, 11:37 PM

State Finished

Completed on Friday, 9 February 2018, 11:38 PM

Time taken 9 secs

Grade 0.00 out of 60.00 (0%)

Result FAIL

QUESTION 1

NOT ANSWERED

[Submit Feedback](#)

A company AWS account consist of approximately 300 IAM users. Now there is a mandate that an access change is required for 100 IAM users to have unlimited privileges to S3.As a system administrator, how can you implement this effectively so that there is no need to apply the policy at the individual user level?

Please select :

- ☐ A. Create a new role and add each user to the IAM role
- ☐ B. Use the IAM groups and add users, based upon their role, to different groups and apply the policy to group
- ☐ C. Create a policy and apply it to multiple users using a JSON script
- ☐ D. Create an S3 bucket policy with unlimited access which includes each user's AWS account ID

Your answer is incorrect.

Answer:- B

An IAM group is used to collectively manage users who need the same set of permissions. By having groups, it becomes easier to manage permissions. So if you change the permissions on the group scale, it will affect all the users in that group. Please find the steps below for the group creation.

Step 1) Go to IAM and click on the Groups section. Click on Create New Group.

Dashboard

Groups

Users

Roles

Policies

Identity providers

Filter



Group Name ↕

Users

No records found.

Step 2) Provide a name for the Group

Set Group Name

Specify a group name. Group names can be edited any time.

Group Name:

Demo




Example: Developers or ProjectAlpha
Maximum 128 characters

Step 3) Next you need to attach a policy. Since the question asks that this group needs full access to S3 , choose the AmazonS3FullAccess role.


Attach Policy

Select one or more policies to attach. Each group can have up to 10 policies

Filter: Policy Type ▼ S3

| | | Policy Name ↕ | Attached Entities ↕ |
|-------------------------------------|---|-------------------------|---------------------|
| <input type="checkbox"/> |  | AmazonDMSRedshiftS3Role | 0 |
| <input checked="" type="checkbox"/> |  | AmazonS3FullAccess | 0 |
| <input type="checkbox"/> |  | AmazonS3ReadOnlyAccess | 0 |

Step 4) Once the group is created, you can then add the 50 users to the group



For more information on users and groups, please visit the url -

<http://docs.aws.amazon.com/IAM/latest/UserGuide/id.html>

The correct answer is: Use the IAM groups and add users, based upon their role, to different groups and apply the policy to group

[Feedback about this question and answer](#)

| QUESTION 2 | NOT ANSWERED | Submit Feedback |
|------------|--------------|---------------------------------|
|------------|--------------|---------------------------------|

You are a consultant tasked with migrating an on-premise application architecture to AWS. During your design process you have to give consideration to current on-premise security and determine which security attributes you are responsible for on AWS. Which of the following does AWS provide for you as part of the shared responsibility model? Choose the 2 correct options

Please select :

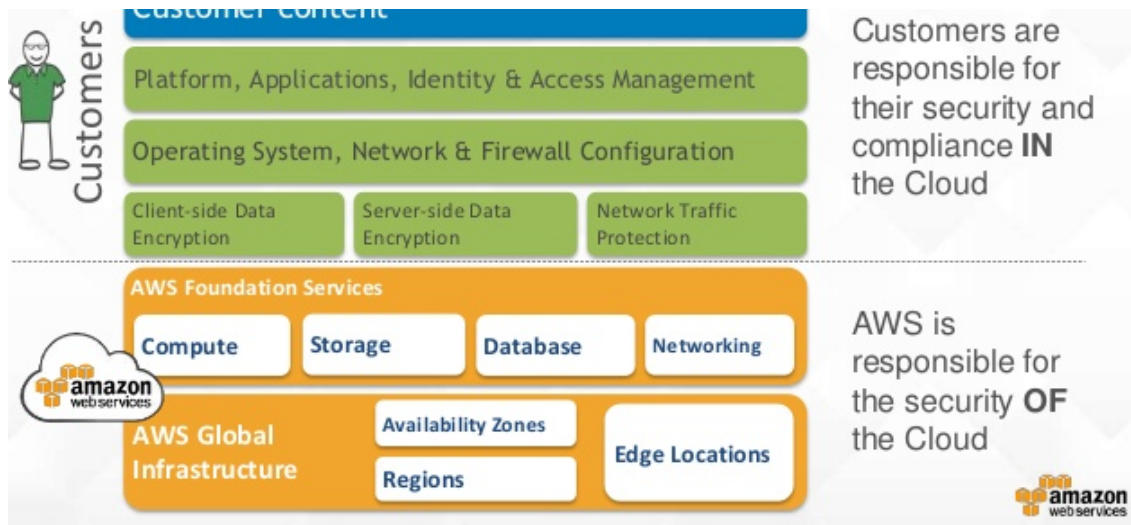
- ☐ A. EC2 Instance security
- ☐ B. Physical network infrastructure
- ☐ C. User access to the AWS environment via IAM.
- ☐ D. Virtualization infrastructure

Your answer is incorrect.

Answer:- B and D.

As per the shared responsibility shown below, the users are required to control the EC2 security via security groups and network access control layers.

Also it is the user's responsibility model, aws takes care of the physical components and the infrastructure to provide Virtualization.



For more information on aws shared responsibility model, please visit the link - <https://aws.amazon.com/blogs/security/tag/shared-responsibility-model/>

The correct answers are: Physical network infrastructure, Virtualization infrastructure

[Feedback about this question and answer](#)

| | | |
|------------|--------------|---------------------------------|
| QUESTION 3 | NOT ANSWERED | Submit Feedback |
|------------|--------------|---------------------------------|

There is a requirement to host an application in aws that requires access to a NoSQL database. But there are no human resources available who can take care of the database infrastructure. Which Amazon service provides a fully-managed and highly available NoSQL service? Choose the correct option

Please select :

- ☐ A. DynamoDB
- ☐ B. ElasticMap Reduce
- ☐ C. Amazon RDS
- ☐ D. SimpleDB

Your answer is incorrect.

Answer:- A

DynamoDB is an aws service that provides a NoSQL database option to users. DynamoDB is a hosted solution by aws , there is no requirement to manage the environment for DynamoDB. And the question clearly states there are no resources in place to manage the DynamoDB environment.

DynamoDB lets you offload the administrative burdens of operating and scaling a distributed database, so that you don't have to worry about hardware provisioning, setup and configuration, replication, software patching, or cluster scaling.

For more information on DynamoDB, please refer the link

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>

The correct answer is: DynamoDB

[Feedback about this question and answer](#)

QUESTION 4

NOT ANSWERED

[Submit Feedback](#)

As a AWS Solution architect , you have been tasked to put the organization data on the cloud. But there is a concern from a security perspective on what can be put on the cloud. What are the best security options from the ones listed below which can be used from a security perspective. Please choose the 3 correct answers from the below options.

Please select :

- ☐ A. Enable EBS Encyption
- ☐ B. Enable S3 Encryption
- ☐ C. Encrypt the file system on an EBS volume using Linux tools
- ☐ D. In AWS , you dont need to worry as it enrypts all data

Your answer is incorrect.

Answer:- A, B and C

Encryption in aws needs to be done by the users and can be done on different levels.

For EBS , we can enable encryption at the volume level. This can be done when the volume is created, this is shown in the screenshot below.

| | | |
|--------------------------|----------------|--|
| Volume Type | i | General Purpose SSD (GP2) ▼ |
| Size (GiB) | i | 100 (Min: 1 GiB, Max: 16384 GiB) |
| IOPS | i | 300 / 3000 (Baseline of 3 IOPS per GiB with a minimum of 100 IOPS, burstable to 3000 IOPS) |
| Throughput (MB/s) | i | Not Applicable |
| Availability Zone | i | us-east-1a ▼ |
| Snapshot ID | i | Search (case-insensitive) |
| Encryption | i | <input checked="" type="checkbox"/> Encrypt this volume |
| Master Key | i | (default) aws/ebs ▼ |

Key Details

| | |
|--------------------|--|
| Description | Default master key that protects my EBS volumes when no other key is defined |
| Account | This account (085363624145) |
| KMS Key ID | ca4e8a5f-8b4e-43f0-97f7-2c3583d57dd0 |
| KMS Key ARN | arn:aws:kms:us-east-1:085363624145:key/ca4e8a5f-8b4e-43f0-97f7-2c3583d57dd0 |

[Cancel](#) [Create](#)

On S3, For any object you can enable server side encryption by going to the Permissions section of the object in S3 and enable the server side encryption option.

Object: aws-programmatic-access-test-object ×

| | |
|-------------------------|---|
| Bucket: | shakinstev |
| Name: | aws-programmatic-access-test-object |
| Link: | https://s3-ap-southeast-1.amazonaws.com/shakinstev/aws-programmatic-access-test-object |
| Size: | 4 |
| Last Modified: | Wed Jul 27 12:40:02 GMT+400 2016 |
| ETag: | 098f6bcd4621d373cade4e832627b4f6 |
| Expiry Date: | None |
| Expiration Rule: | N/A |

▼ Details

Storage Class: ☒ Standard ☐ Standard - Infrequent Access ☐ Reduced Redundancy

Server Side Encryption: ☐ None ☒ AES-256

And finally , one can use Linux based tools to Encrypt a volume if it is not encrypted.

The correct answers are: Enable EBS Encryption, Enable S3 Encryption, Encrypt the file system on an EBS volume using Linux tools

[Feedback about this question and answer](#)

AWS provides a storage option known as Amazon Glacier. What is this aws service designed for. Please specify 2 correct options.

Please select :

- ☐ A. Cached session data
- ☐ B. Infrequently accessed data
- ☐ C. Data archives
- ☐ D. Active database storage

Your answer is incorrect.

Answer – B and C

Amazon Glacier is an extremely low-cost storage service that provides secure, durable, and flexible storage for data backup and archival. So Amazon glacier is used for Infrequently accessed data and Data archives.

For Cached Data Session , the service provided by aws is known as elastic cache. So Amazon glacier is the wrong option.

For Active database storage , this is done via EBS volumes , so this option is also incorrect.

For more information on Amazon Glacier , please visit the link -

<https://aws.amazon.com/glacier/faqs/>

The correct answers are: Infrequently accessed data, Data archives

[Feedback about this question and answer](#)

| | | |
|------------|--------------|---------------------------------|
| QUESTION 6 | NOT ANSWERED | Submit Feedback |
|------------|--------------|---------------------------------|

There is a requirement for a user to modify the configuration of one of your Elastic Load Balancers (ELB). This access is just required one time only. Which of the following choices would be the best way to allow this access?

Please select :

- ☐ A. Open up whichever port ELB uses in a security group and give the user access to that security group via a policy
- ☐ B. Create an IAM Role and attach a policy allowing modification access to the ELB
- ☐ C. Create a new IAM user who only has access to the ELB resources and delete that user when the work is completed.
- ☐ D. Give them temporary access to the root account for 12 hours only and change the password once the activity is completed

Your answer is incorrect.

Answer – B

Step 1) Create a role which has the required ELB access

Set Role Name

Enter a role name. You cannot edit the role name after the role is created.

Role Name

ELBAccess



Maximum 64 characters. Use alphanumeric and '+,=, @, _' characters

Step 2) You need to provide permissions to the underlying EC2 instances in the Elastic Load Balancer

Select Role Type

AWS Service Roles

Amazon EC2
Allows EC2 instances to call AWS services on your behalf. Select

| | | | | | |
|-------------------------------------|---|------------------------------|---|---------------------------|---------------------------|
| <input type="checkbox"/> |  | AmazonEC2ContainerService... | 0 | 2015-04-09 20:14 UTC+0400 | 2016-08-11 17:08 UTC+0400 |
| <input checked="" type="checkbox"/> |  | AmazonEC2FullAccess | 0 | 2015-02-06 22:40 UTC+0400 | 2015-02-06 22:40 UTC+0400 |

For the best practises on IAM policies, please visit the link:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>
- http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_create_for-user.html

The correct answer is: Create an IAM Role and attach a policy allowing modification access to the ELB

[Feedback about this question and answer](#)

QUESTION 7

NOT ANSWERED

[Submit Feedback](#)

You are an AWS Solution Architect and architecting an application environment on AWS. Which services or service features might you enable to take advantage of monitoring to ensure auditing the environment for compliance is easy and follows the strict security compliance requirements? Choose the correct option

Please select :

- ☐ A. CloudTrail for security logs
- ☐ B. S3 logging
- ☐ C. Encrypted data storage

Your answer is incorrect.

Answer – A

AWS Cloudtrail is the defacto service provided by aws for monitoring all API calls to AWS and is used for logging and monitoring purposes for compliance purposes. Amazon cloudtrail detects every call made to aws and creates a log which can then be further used for analysis.

For more information on Amazon Cloudtrail , please visit the link –

<https://aws.amazon.com/cloudtrail/>

The correct answer is: CloudTrail for security logs

[Feedback about this question and answer](#)

QUESTION 8

NOT ANSWERED

[Submit Feedback](#)

An application has been migrated from on-premise to AWS in your company and you will not be responsible for the ongoing maintenance of packages. Which of the below services allows for access to the underlying infrastructure. Choose the 2 correct options

Please select :

- ☐ A. Elastic Beanstalk
- ☐ B. EC2
- ☐ C. DynamoDB
- ☐ D. RDS

Your answer is incorrect.

Answer – A and B

EC2 and Elastic Beanstalk are aws services that allow the developer access to the underlying infrastructure.

When you create an Elastic beanstalk environment as shown below , you will have access to the underlying EC2 instance. So in the below example , for the Elastic beanstalk environment , you will have access to the Windows Server 2012 environment.



Health

Grey

Causes

Running Version

Sample Application

Upload and Deploy

 Windows Server 2012

Configuration

64bit Windows Server 2012 R2

v1.2.0 running IIS 8.5

Change

DynamoDB and RDS are services provided and the infrastructure is managed by aws.

The correct answers are: Elastic Beanstalk, EC2

[Feedback about this question and answer](#)

QUESTION 9

NOT ANSWERED

[Submit Feedback](#)

To protect S3 data from both accidental deletion and accidental overwriting, you should

Please select :

- ☐ A. Enable Multi-Factor Authentication (MFA) protected access
- ☐ B. Disable S3 delete using an IAM bucket policy
- ☐ C. Access S3 data using only signed URLs
- ☐ D. Enable S3 versioning on the bucket

Your answer is incorrect.

Answer- D

To protect objects in S3 from both accidental deletion and accidental overwriting, the methodology adopted by aws is to Enable versioning on the bucket. Versioning allows to store every version of an object , so that if by mistake there is a version deleted , you can recover other versions, because the entire object is not deleted.

Enable Multi-Factor Authentication (MFA) protected access on S3 is only used to add an additional security layer to S3. So that users who are authenticated properly before having access to the bucket. But this is not what the question is asking.

To enable versioning on S3 , you need to go to the bucket , and in the properties , you can enable versioning.

► Logging

► Events

▼ Versioning

Versioning allows you to preserve, retrieve, and restore every version of every object stored in this bucket. This provides an additional level of protection by providing a means of recovery for accidental overwrites or expirations. Versioning-enabled buckets store all versions of your objects by default.

You can use Lifecycle rules to manage all versions of your objects as well as their associated costs. Lifecycle rules enable you to automatically archive your objects to the Glacier Storage Class and/or remove them after a specified time period.

Once enabled, Versioning cannot be disabled, only suspended.

Versioning is currently not enabled on this bucket.

Enable Versioning

The correct answer is: Enable S3 versioning on the bucket

[Feedback about this question and answer](#)

QUESTION 10

NOT ANSWERED

[Submit Feedback](#)

By default is data in S3 encrypted?

Please select :

- ☐ A. Yes, S3 always encrypts data for security purposes.
- ☐ B. Yes, but only in government cloud data centers
- ☐ C. No, but it can be when the right APIs are called for SSE
- ☐ D. No, it must be encrypted before upload of any data to S3.

Your answer is incorrect.

Answer- C

Please note that, no , by default , Encryption is not enabled. So option A and B are incorrect.

Also note that it is not necessary to encrypt before every upload.

For any object you can enable server side encryption by going to the Permissions section of the object in S3 and enable the server side encryption option.

Bucket: shakinstev
Name: aws-programmatic-access-test-object
Link: <https://s3-ap-southeast-1.amazonaws.com/shakinstev/aws-programmatic-access-test-object>
Size: 4
Last Modified: Wed Jul 27 12:40:02 GMT+400 2016
ETag: 098f6bcd4621d373cade4e832627b4f6
Expiry Date: None
Expiration Rule: N/A

▼ Details

Storage Class: ☒ Standard ☐ Standard - Infrequent Access ☐ Reduced Redundancy

Server Side Encryption: ☐ None ☒ AES-256

For more information on Encryption for S3 , please refer to the link -
<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingEncryption.html>

The correct answer is: No, but it can be when the right APIs are called for SSE

[Feedback about this question and answer](#)

QUESTION 11

NOT ANSWERED

[Submit Feedback](#)

Your AWS environment contains several reserved EC2 instances dedicated to a project that has just been cancelled. You need to stop incurring charges for the reserved instances immediately. What steps would you take to avoid taking the hit on the charge for these reserved instances? Choose 2 correct options

Please select :

- ☐ A. Stop the instances as soon as possible.
- ☐ B. Contact AWS and explain the situation to try and recover the costs.
- ☐ C. Sell the reserved instances on the AWS Reserved Instance Marketplace
- ☐ D. Terminate the instances as soon as possible.

Your answer is incorrect.

Answer: - C and D.

Reserved Instances provide a significant discount (up to 75%) compared to On-Demand instance pricing. There is a fixed quote of reserved power that it is given to the account. You have the Flexibility to change families, OS types, and tenancies while benefitting from Reserved Instance pricing.

Now reserved instances are bought upfront for a specified duration. Unlike On-demand instances, there is no cost difference if you stop the instances, so option A is incorrect. Since you have already bought the reserved instances, you cannot ask aws to recover the costs.

The only 2 options available are to terminate the instances immediately and sell them on AWS Reserved Instance Marketplace for a specified price.

You can purchase reserved instances from the reserved instances section in the EC2 dashboard.

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Spot Requests

Reserved Instances

Scheduled Instances

Dedicated Hosts

Purchase Reserved Instances

Actions ▾

Filter by attributes or search by keyword

To buy a Re:

In the next screen, you can choose the reserved instance to buy. But you are making an upfront commitment to buy the instances.

| Platform Linux/UNIX ▾ | | Tenancy Default ▾ | | Offering Class Any ▾ | | | | | |
|--------------------------|-----------|-------------------|-----------------|----------------------|------------------|-------------------|----------------------|--------------------------------|------------------------|
| Instance Type t2.micro ▾ | | Term Any ▾ | | Payment Option Any ▾ | | <div>Search</div> | | | |
| Seller ▾ | Term ▾ | Effective Rate ▾ | Upfront Price ▾ | Hourly Rate ▾ | Payment Option ▾ | Offering Class ▾ | Quantity Available ▾ | Desired Quantity | |
| AWS | 36 months | \$0.008 | \$0.00 | \$0.008 | No Upfront | convertible | Unlimited | <input type="text" value="1"/> | <div>Add to Cart</div> |
| AWS | 12 months | \$0.009 | \$0.00 | \$0.009 | No Upfront | standard | Unlimited | <input type="text" value="1"/> | <div>Add to Cart</div> |
| AWS | 12 months | \$0.009 | \$51.00 | \$0.003 | Partial Upfront | standard | Unlimited | <input type="text" value="1"/> | <div>Add to Cart</div> |
| AWS | 12 months | \$0.009 | \$75.00 | \$0.000 | All Upfront | standard | Unlimited | <input type="text" value="1"/> | <div>Add to Cart</div> |

For more information on reserved instances please follow the link -

<https://aws.amazon.com/ec2/pricing/reserved-instances/>

The correct answers are: Sell the reserved instances on the AWS Reserved Instance Marketplace, Terminate the instances as soon as possible.

[Feedback about this question and answer](#)

for your data? Please choose 2 options.

Please select :

- ☐ A. Encrypt the data locally using your own encryption keys, then copy the data to Amazon S3 over HTTPS endpoints
- ☐ B. Store the data on EBS volumes with encryption enabled instead of using Amazon S3
- ☐ C. Store the data in encrypted EBS snapshots
- ☐ D. Enable SSE on an S3 bucket to make use of AES-256 encryption.

Your answer is incorrect.

Answer:- A and D.

The question asks for Encryption at rest for S3, so any answer related to EBS encryption does not correspond to the right answer.

For any object you can enable server side encryption by going to the Permissions section of the object in S3 and enable the server side encryption option.

Object: aws-programmatic-access-test-object ×

Bucket: shakinstev
Name: aws-programmatic-access-test-object
Link: <https://s3-ap-southeast-1.amazonaws.com/shakinstev/aws-programmatic-access-test-object>
Size: 4
Last Modified: Wed Jul 27 12:40:02 GMT+400 2016
ETag: 098f6bcd4621d373cade4e832627b4f6
Expiry Date: None
Expiration Rule: N/A

▼ Details

Storage Class: ☒ Standard ☐ Standard - Infrequent Access ☐ Reduced Redundancy

Server Side Encryption: ☐ None ☒ AES-256

And then for client side encryption, you can encrypt the object and send it to S3 when you program your application.

For the entire detailed description on Encryption strategies, please visit the link -

<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingEncryption.html>

The correct answers are: Encrypt the data locally using your own encryption keys, then copy the data to Amazon S3 over

HTTPS endpoints, Enable SSE on an S3 bucket to make use of AES-256 encryption.

[Feedback about this question and answer](#)

QUESTION 13

NOT ANSWERED

[Submit Feedback](#)

Which of the following is true of an SQS message? Choose the correct option

- ☐ B. SQS messages must be less than 32 KB in size
- ☐ C. SQS messages must be in JSON format
- ☐ D. SQS messages can live in the queue up to thirty days

Your answer is incorrect.

Answer - A

If you look at the SQS FAQ, it is clearly mentioned that SQS messages are guaranteed to be delivered at least once.

Q: Does Amazon SQS guarantee delivery of messages?

Standard queues provide at-least-once delivery, which means that each message is delivered at least once.

The message size for SQS can be 256KB in size.

The message formats can be in XML, JSON and unformatted text.

The messages can live in the queue for a maximum of 14 days.

For more information on SQS messages, please follow the link

<https://aws.amazon.com/sqs/faqs/>

The correct answer is: SQS messages are guaranteed to be delivered at least once

[Feedback about this question and answer](#)

QUESTION 14

NOT ANSWERED

[Submit Feedback](#)

An EC2 instance has been running and data has been stored on the instance's volumes. The instance was shutdown over the weekend to save costs. The next week, after starting the instance, you notice that all data is lost and is no longer available on the EC2 instance. What might be the cause of this?

Please select :

- ☐ A. The EC2 instance was using instance store volumes, which are ephemeral and only lives for the life of the instance
- ☐ B. The EC2 instance was using EBS backed root volumes, which are ephemeral and only lives for the life of the instance
- ☐ C. The EBS volume was not big enough to handle all of the processing data.
- ☐ D. The instance has been compromised

Your answer is incorrect.

Answer - A

Anything that is stored on an instance store volume is destroyed when the instance is shutdown. Instance store volumes are ephemeral, which means that they only survive when the instance is active.

EBS backed Volumes are not ephemeral and exists even if the instance is stopped and started, so Option B is wrong.

Even if EBS volume is not big enough, it does not mean that it will not be present when the instance is stopped and started, so Option C is wrong.

The correct answer is: The EC2 instance was using instance store volumes, which are ephemeral and only lives for the life of the instance

[Feedback about this question and answer](#)

QUESTION 15

NOT ANSWERED

[Submit Feedback](#)

What database services are provided by aws from the list mentioned below. Choose the 3 correct answers.

Please select :

- ☐ A. Aurora
- ☐ B. MariaDB
- ☐ C. MySQL
- ☐ D. DB2

Your answer is incorrect.

Answer – A, B and C

DB2 is the only database service not provided by AWS.

- For the list of DB services , please visit the link - <https://aws.amazon.com/rds/>
- For more information on Aurora , please visit the link - <https://aws.amazon.com/rds/aurora/>
- For more information on mySQL , please visit the link - <https://aws.amazon.com/rds/mysql/>
- For more information on MariaDB , please visit the link - <https://aws.amazon.com/rds/mariadb/>

The correct answers are: Aurora, MariaDB, MySQL

[Feedback about this question and answer](#)

QUESTION 16

NOT ANSWERED

[Submit Feedback](#)

There is a requirement to move 10 TB data warehouse to the cloud. With the current bandwidth allocation it would take 2 months to transfer the data. Which service would allow you to quickly get ther data into AWS? Choose the correct option.

Please select :

- ☐ A. Amazon Import/Export
- ☐ B. Amazon Direct Connect
- ☐ C. Amazon S3 MultiPart Upload
- ☐ D. Amazon S3 Connector

Your answer is incorrect.

Answer – A

For Amazon S3 multipart uploads, there are the following restrictions, so choose better to use Amazon Import/Export.

| Item | Specification |
|--|---------------------------------------|
| Maximum object size | 5 TB |
| Maximum number of parts per upload | 10,000 |
| Part numbers | 1 to 10,000 (inclusive) |
| Part size | 5 MB to 5 GB, last part can be < 5 MB |
| Maximum number of parts returned for a list parts request | 1000 |
| Maximum number of multipart uploads returned in a list multipart uploads request | 1000 |

For more information on aws import/export, please visit the link -

<https://aws.amazon.com/snowball/>

Amazon Direct Connect is used as a connection between AWS and On-premise so this is the wrong option.

The correct answer is: Amazon Import/Export

[Feedback about this question and answer](#)

QUESTION 17

NOT ANSWERED

[Submit Feedback](#)

What is the difference between an availability zone and an edge location? Choose the correct option

Please select :

- ☐ A. An availability zone is a grouping of AWS resources in a specific region; an edge location is a specific resource within the AWS region
- ☐ B. An availability zone is an **isolated location** within an AWS region, whereas an edge location will deliver cached content to the closest location to reduce latency
- ☐ C. Edge locations are used as control stations for AWS resources
- ☐ D. None of the above

Your answer is incorrect.

Answer – B

In aws , there are regions with each region separated in a separate geographic area. Each region has multiple, isolated locations known as Availability Zones. An availability zone is used to host resources in a specific region.

For more information on Regions and availability zone, please visit the url -

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

integrity of data is maintained across all edge locations.

For more information on Edge locations, please visit the url –

<https://aws.amazon.com/about-aws/global-infrastructure/>

The correct answer is: An availability zone is an **isolated location** within an AWS region, whereas an edge location will deliver cached content to the closest location to reduce latency

[Feedback about this question and answer](#)

QUESTION 18

NOT ANSWERED

[Submit Feedback](#)

An order processing website is using EC2 instances to process messages from an SQS queue. A user reported an issue that their order was processed twice and hence charged twice. What action would you recommend to ensure this does not happen again? Choose the correct option

Please select :

- ☐ A. Insert code into the application to delete messages after processing
- ☐ B. Increase the visibility timeout for the queue
- ☐ C. Modify the order process to use SWF
- ☐ D. Use long polling rather than short polling

Your answer is incorrect.

Answer – C

This is a tricky question and note that Options A,B and D can be used to decrease the likelihood of duplicate messages , but cannot remove the change entirely.

For option A , even if the code is inserted , which should be the case already , if the EC2 instance goes down , the same issue can occur again. The message will not be deleted and when it comes in the SQS queue , it will be processed again.

For option B , even if you increase the visibility timeout , if the process has taken the message but not deleted the message , after the visibility timeout expires , the EC2 instance will again process the message and the same issue will happen again.

If you use long polling instead of short polling , you still have the same problem with Option A and B.

For more information on SQS , please visit the link – <https://aws.amazon.com/sqs/faqs/>

The correct answer is: Modify the order process to use SWF

[Feedback about this question and answer](#)

QUESTION 19

NOT ANSWERED

[Submit Feedback](#)

with an attached internet gateway. The instances also belong to the same security group. However, one of the instances is not able to send or receive traffic like the other three. You see that there is no OS level issue and the instance is working as it should.

What could be the possible issue? Choose the correct option.

Please select :

- ☐ A. A proper route table configuration that sends traffic from the instance to the Internet through the internet gateway
- ☐ B. The EC2 instance is running in an availability zone that does not support Internet gateways
- ☐ C. The EC2 instance is not a member of the same Auto Scaling group/policy
- ☐ D. The EC2 instance does not have a public IP address associated with it

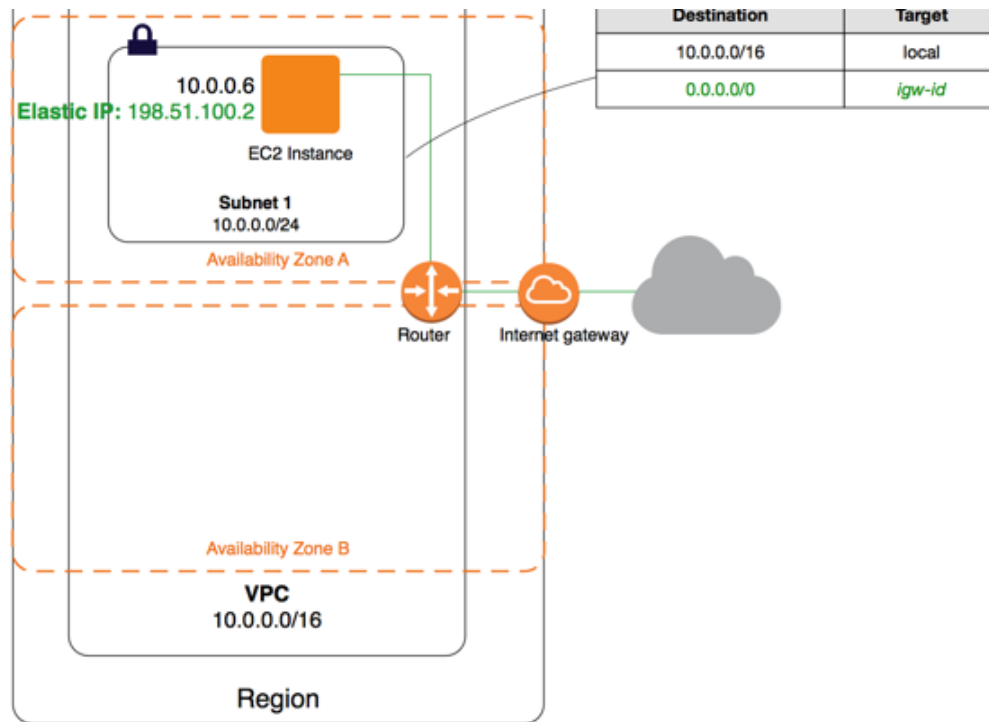
Your answer is incorrect.

Answer – D

Below is a sample VPC from the aws VPC guides. For an instance to be available from the internet , you need to ensure

- 1) The Internet gateway is in place – This has been confirmed in the question.
- 2) There is a route entry for the internet gateway – This should be in place , because out of the 2 instances , one is working.
- 3) The EC2 instance should have a public or Elastic IP – From the question , there is no mention of one being allocated to the problem instance.

Hence option D is the right answer.



For more information on VPC public subnets , please visit the url - http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html

The correct answer is: The EC2 instance does not have a public IP address associated with it
[Feedback about this question and answer](#)

QUESTION 20

NOT ANSWERED

[Submit Feedback](#)

Which of the following best describes what "bastion hosts" are? Choose the correct option.

Please select :

- ☐ A. Bastion hosts are instances that sit within a private subnet and are typically accessed using SSH or RDP. Once remote connectivity has been established with the bastion host, it then acts as a 'jump' server, allowing you to use SSH or RDP to log into other instances (within public subnets) deeper within your network.
- ☐ B. Bastion hosts are instances that sit within your public subnet and are typically accessed using SSH or RDP. Once remote connectivity has been established with the bastion host, it then acts as a 'jump' server, allowing you to use HTTPS to log into other instances (within private subnets) deeper within your network.
- ☐ C. Bastion hosts are instances that sit within your public subnet and are typically accessed using SSH or RDP. Once remote connectivity has been established with a bastion host, it then acts as a 'jump' server, allowing you to use SSH or RDP to log into other instances (within private subnets) deeper within your network.
- ☐ D. Bastion hosts are instances that sit within your private subnet and are typically accessed using SSH or RDP. Once remote connectivity has been established with the bastion host, it then acts as a 'jump' server, allowing you to use HTTPS to log into other instances (within public subnets) deeper within your network.

Your answer is incorrect.

However, the computer generally hosts a single application, for example a proxy server, and all other services are removed or limited to reduce the threat to the computer.

In AWS, A bastion host is kept on a public subnet. Users log on to the bastion host via SSH or RDP and then use that session to manage other hosts in the private subnets.

This is a security practise adopted by many organization to secure the assets in their private subnets.

The correct answer is: Bastion hosts are instances that sit within your public subnet and are typically accessed using SSH or RDP. Once remote connectivity has been established with a bastion host, it then acts as a 'jump' server, allowing you to use SSH or RDP to log into other instances (within private subnets) deeper within your network.

[Feedback about this question and answer](#)

QUESTION 21

NOT ANSWERED

[Submit Feedback](#)

A web application is hosted on EC2 instances and using SQS. Requests are saved as messages in the SQS queue. The SQS queue is configured with the maximum message retention period. After 10 days you notice that the application was in a hung state and 2000 messages are still lying in the queue unprocessed. You are going to resolve the issue but you need to send a communication to the users on the issue. What information will you provide?

Choose the correct option.

Please select :

- ☐ A. An apology for the delay in processing requests and telling them that unfortunately they have to resubmit all the requests.
- ☐ B. An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that requests greater than five days old will need to be resubmitted.
- ☐ C. An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that all received requests will be processed at that time.

Your answer is incorrect.

Answer – C

Since the question states that the SQS is configured with the maximum retention period , it means that messages can last for 14 days.

- So option A is invalid , since the messages will still be in the queue even after 10 days
- Option B is invalid for the same reason noted in Option B
- Option D is invalid because a queue can have up to 120,000 messages.

Q: How long can I keep my messages in Amazon SQS message queues?

Longer message retention provides greater flexibility to allow for longer intervals between message production and consumption.

You can configure the Amazon SQS message retention period to a value from 1 minute to 14 days. The default is 4 days. Once the message retention limit is reached, your messages are automatically deleted.

For more information on SQS , please visit the link:

- <https://aws.amazon.com/sqs/>

The correct answer is: An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that all received requests will be processed at that time.

[Feedback about this question and answer](#)

QUESTION 22

NOT ANSWERED

[Submit Feedback](#)

A Company provides an online service that utilizes SQS to decouple system components for scalability. The SQS consumer's EC2 instances poll the queue as often as possible to keep end-to-end throughput as high as possible. However, it is noticed that polling in tight loops is burning CPU cycles and increasing costs with empty responses. What can be done to reduce the number of empty responses? Choose the correct option.

Please select :

- ☐ A. Scale the component making the request using Auto Scaling based off the number of messages in the queue
- ☐ B. Enable long polling by setting the ReceiveMessageWaitTimeSeconds to a number > 0
- ☐ C. Enable short polling on the SQS queue by setting the ReceiveMessageWaitTimeSeconds to a number > 0
- ☐ D. Enable short polling on the SQS message by setting the ReceiveMessageWaitTimeSeconds to a number = 0

Your answer is incorrect.

Answer - B

enabling long polling. And this can be done by setting the `ReceiveMessageWaitTimeSeconds` attribute of the queue to a value greater than 0.

You can do this by changing the queue attributes as shown below

Queue Attributes

| | | | |
|-------------------------------------|----------------------------------|----------------------|---|
| Default Visibility Timeout ⓘ | <input type="text" value="30"/> | <div>seconds ▾</div> | Value must be between 0 seconds and 12 hours. |
| Message Retention Period ⓘ | <input type="text" value="4"/> | <div>days ▾</div> | Value must be between 1 minute and 14 days. |
| Maximum Message Size ⓘ | <input type="text" value="256"/> | KB | Value must be between 1 and 256 KB. |
| Delivery Delay ⓘ | <input type="text" value="0"/> | <div>seconds ▾</div> | Value must be between 0 seconds and 15 minutes. |
| Receive Message Wait Time ⓘ | <input type="text" value="10"/> | seconds | Value must be between 0 and 20 seconds. |

Answer - B

For more information on polling, please visit the link -

<http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-long-polling.html>

The correct answer is: Enable long polling by setting the `ReceiveMessageWaitTimeSeconds` to a number > 0

[Feedback about this question and answer](#)

QUESTION 23

NOT ANSWERED

[Submit Feedback](#)

Your company has resources set up on the AWS Cloud. Your company is now going through a set of scheduled audits by an external auditing firm. Which of the following services can be utilized to help ensure the right information is present for auditing purposes.

Please select :

- ☐ A. AWS CloudTrail
- ☐ B. AWS VPC
- ☐ C. AWS EC2
- ☐ D. AWS Cloudwatch

Your answer is incorrect.

Answer – A

The AWS Documentation mentions the following on Cloudtrail

AWS CloudTrail is a service that enables governance, compliance, operational auditing, and risk auditing of your AWS account. With CloudTrail, you can log, continuously monitor, and retain account activity related to actions across your AWS infrastructure. CloudTrail provides event history of your AWS account activity, including actions taken

For more information on CloudTrail, please visit the below URL:

- <https://aws.amazon.com/cloudtrail/>

The correct answer is: AWS CloudTrail

[Feedback about this question and answer](#)

| QUESTION 24 | NOT ANSWERED | Submit Feedback |
|---|--------------|-----------------|
| <p>Which of the following will incur a cost when working with AWS resources. Choose 2 answers from the options given below</p> <p>Please select :</p> <ul style="list-style-type: none"><input type="checkbox"/> A. A running EC2 Instance<input type="checkbox"/> B. A stopped EC2 Instance<input type="checkbox"/> C. EBS Volumes attached to stopped EC2 Instances<input type="checkbox"/> D. Using an Amazon VPC | | |
| <p>Your answer is incorrect.</p> <p>Answer – A and C</p> <p>The AWS Documentation clearly mentions the cost to EC2 Instances</p> <p>Billing commences when Amazon EC2 initiates the boot sequence of an AMI instance. Billing ends when the instance terminates, which could occur through a web services command, by running "shutdown -h", or through instance failure. When you stop an instance, we shut it down but don't charge hourly usage for a stopped instance, or data transfer fees, but we do charge for the storage for any Amazon EBS volumes.</p> <p>For more information, please visit the below URL:</p> <ul style="list-style-type: none">• https://aws.amazon.com/ec2/faqs/ <p>The AWS Documentation clearly mentions the cost with regards to the VPC</p> <p>There are no additional charges for creating and using the VPC itself.</p> <p>For more information, please visit the below URL:</p> <ul style="list-style-type: none">• https://aws.amazon.com/vpc/faqs <p>The correct answers are: A running EC2 Instance, EBS Volumes attached to stopped EC2 Instances</p> <p>Feedback about this question and answer</p> | | |

| QUESTION 25 | NOT ANSWERED | Submit Feedback |
|--|--------------|-----------------|
| <p>As part of your application architecture requirements, the company you are working for has requested the ability to run analytics against all combined log files from the Elastic Load Balancer. Which services are used together to collect logs and process log file analysis in an AWS environment? Choose the correct option.</p> | | |

- ☐ B. Amazon EC2 for storing and processing the log files
- ☐ C. Amazon S3 for storing the ELB log files and EC2 for processing the log files in analysis
- ☐ D. Amazon S3 for storing ELB log files and Amazon EMR for processing the log files in analysis

Your answer is incorrect.

Answer – D

This question is not that complicated, even though if you don't understand the options. By default when you see "collection of logs and processing of logs", directly think of AWS EMR.

Amazon EMR provides a managed Hadoop framework that makes it easy, fast, and cost-effective to process vast amounts of data across dynamically scalable Amazon EC2 instances. You can also run other popular distributed frameworks such as Apache Spark, HBase, Presto, and Flink in Amazon EMR, and interact with data in other AWS data stores such as Amazon S3 and Amazon DynamoDB.

Amazon EMR securely and reliably handles a broad set of big data use cases, including log analysis, web indexing, data transformations (ETL), machine learning, financial analysis, scientific simulation, and bioinformatics.

For more information on EMR, please visit the link –

<https://aws.amazon.com/emr/>

The correct answer is: Amazon S3 for storing ELB log files and Amazon EMR for processing the log files in analysis

[Feedback about this question and answer](#)

QUESTION 26

NOT ANSWERED

[Submit Feedback](#)

You have been told that you need to set up a bastion host by your manager in the cheapest, most secure way, and that you should be the only person that can access it via SSH. Which of the following setups would satisfy your manager's request?

Choose the correct option

Please select :

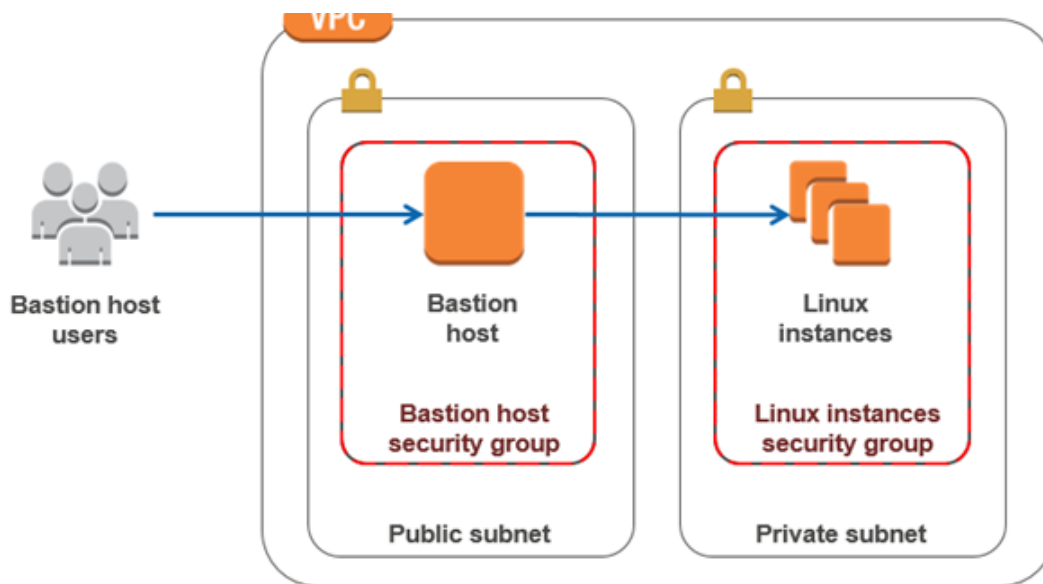
- ☐ A. A large EC2 instance and a security group which only allows access on port 22
- ☐ B. A large EC2 instance and a security group which only allows access on port 22 via your IP address
- ☐ C. A small EC2 instance and a security group which only allows access on port 22
- ☐ D. A small EC2 instance and a security group which only allows access on port 22 via your IP address

Your answer is incorrect.

Answer – D

A bastion host should always be a small EC2 instance, because there is no requirement of applications to run on it. Also you should only open port 22 from your IP address and no other IP Address.

In AWS , A bastion host is kept on a public subnet. Users log on to the bastion host via SSH or RDP and then use that session to manage other hosts in the private subnets.



The correct answer is: A small EC2 instance and a security group which only allows access on port 22 via your IP address

[Feedback about this question and answer](#)

QUESTION 27

NOT ANSWERED

[Submit Feedback](#)

You have a web application hosted in AWS on EC2 Instances. The application provides newspaper content to users around the world. Off late , the load on the web application has increased and is subsequently increasing the response time for the application for end users. Which of the below services can be used to alleviate this problem. Choose 2 answers from the options given below

Please select :

- ☐ A. Use Cloudfront and use the web application as the origin
- ☐ B. Use AWS Storage gateways to distribute the content across multiple storage devices for better read throughput.
- ☐ C. Use Elastic cache behind of the web application.
- ☐ D. Consider using SQS to process some of the user requests

Your answer is incorrect.

Answer – A and C

The AWS Documentation provides the following information on Cloudfront and Elastic Cache

Amazon CloudFront is a web service that speeds up distribution of your static and dynamic web content, such as .html, .css, .php, and image files, to your users. CloudFront delivers your content through a worldwide network of data centers called edge locations. When a user requests content that you're serving with CloudFront, the user is routed to the edge location that provides the lowest latency (time delay), so that content is delivered with the best possible performance.

For more information on AWS Cloudfront, please visit the below URL:

or cache in the cloud. The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases.

For more information on AWS Elastic Cache, please visit the below URL:

- <https://aws.amazon.com/elasticache/>

The correct answers are: Use Cloudfront and use the web application as the origin, Use Elastic cache behind of the web application.

[Feedback about this question and answer](#)

QUESTION 28

NOT ANSWERED

[Submit Feedback](#)

Your supervisor asks you to create a highly available website which serves static content from EC2 instances. Which of the following is not a requirement to accomplish this goal? Choose the correct option

Please select :

- ☐ A. Multiple Availability Zones
- ☐ B. Multiple subnets
- ☐ C. An SQS queue
- ☐ D. An auto scaling group to recover from EC2 instance failures


Your answer is incorrect.

Answer – C

For highly available websites, yes Multiple Availability Zones and Multiple subnets are required. Below is a simple architecture of a highly available website consisting of an ELB and 2 AZ's. BY default each AZ should be located in a different subnet.

Also auto-scaling is used to add additional EC2 instances for fault tolerance.

SQS is not an option, because SQS is only used to decouple components in an architecture, it is not necessary for a high available web site.



The correct answer is: An SQS queue

[Feedback about this question and answer](#)

QUESTION 29

NOT ANSWERED

[Submit Feedback](#)

What is the maximum object size allowed for Multi-part file upload for S3.

Please select :

- ☐ A. 10 TB
- ☐ B. 5 TB
- ☐ C. 1 TB
- ☐ D. 5 GB

Your answer is incorrect.

Answer – B

Please refer to the below table which gives the restrictions for the Multi-part file upload for S3. From here it clearly shows that the right answer is 5TB.

| Item | Specification |
|--|---------------------------------------|
| Maximum object size | 5 TB |
| Maximum number of parts per upload | 10,000 |
| Part numbers | 1 to 10,000 (inclusive) |
| Part size | 5 MB to 5 GB, last part can be < 5 MB |
| Maximum number of parts returned for a list parts request | 1000 |
| Maximum number of multipart uploads returned in a list multipart uploads request | 1000 |

[Feedback about this question and answer](#)

QUESTION 30

NOT ANSWERED

[Submit Feedback](#)

Which of the following statements about S3 are true. Please choose 2 options

Please select :

- ☐ A. The total volume of data and number of objects you can store are unlimited
- ☐ B. Individual Amazon S3 objects can range in size from a minimum of 0 bytes to a maximum of 1 terabytes
- ☐ C. You can use Multi-Object Delete to delete large numbers of objects from Amazon S3
- ☐ D. You can store only objects of a particular format in S3

Your answer is incorrect.

Answer – A and C

The below screenshots are from the S3 FAQ's of AWS.

The total volume of data and number of objects you can store are unlimited. Individual Amazon S3 objects can range in size from a minimum of 0 bytes to a maximum of 5 terabytes. The largest object that can be uploaded in a single PUT is 5 gigabytes. For objects larger than 100 megabytes, customers should consider using the [Multipart Upload](#) capability.

You can use [Multi-Object Delete](#) to delete large numbers of objects from Amazon S3. This feature allows you to send multiple object keys in a single request to speed up your deletes. Amazon does not charge you for using Multi-Object Delete.

Option B is incorrect, because as per the S3 definition, the maximum size of objects can be 5 TB

Option D is incorrect because you can virtually store objects of any type

For more information on S3 , please visit the URL –

<https://aws.amazon.com/s3/faqs/>

The correct answers are: The total volume of data and number of objects you can store are unlimited, You can use Multi-Object Delete to delete large numbers of objects from Amazon S3

[Feedback about this question and answer](#)

QUESTION 31

NOT ANSWERED

[Submit Feedback](#)

What is a document that provides a formal statement of one or more permissions?

Please select :

- ☐ A. Policy
- ☐ B. Permission

Your answer is incorrect.

Answer – A

A policy is a JSON document that specifies what a user can do on AWS. This document consists of

- Actions: what actions you will allow. Each AWS service has its own set of actions.
- Resources: which resources you allow the action on.
- Effect: what the effect will be when the user requests access—either allow or deny.

Below is a sample snippet of a policy document that allows access to all users to Describe EC2 Instances. You can clearly see the Actions, Resources and Effect which define the policy document.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "ec2:Describe*",
      "Resource": "*"
    }
  ],
}
```

For more information on policies, please visit the url -

http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html

The correct answer is: Policy

[Feedback about this question and answer](#)

QUESTION 32

NOT ANSWERED

[Submit Feedback](#)

Which of the following is not required to ensure that you can SSH into a Linux instance hosted in a VPC from the internet.

Please select :

- ☐ A. Private IP Address
- ☐ B. Public IP Address
- ☐ C. Internet gateway attached to the VPC
- ☐ D. Elastic IP

Your answer is incorrect.

Answer – A

The AWS Documentation provides the following information

A private IPv4 address is an IP address that's not reachable over the Internet. You can use private IPv4 addresses for communication between instances in the same network (EC2-Classical or a VPC).

For more information on AWS IP Addressing, please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-instance-addressing.html>

QUESTION 33

NOT ANSWERED

[Submit Feedback](#)

What are the two permission types used by AWS?

Please select :

- ☐ A. Resource-based and Product-based
- ☐ B. Product-based and Service-based
- ☐ C. Service-based
- ☐ D. User-based and Resource-based

Your answer is incorrect.

Answer – D

Permissions are defined via policies which consist of the following elements

- Actions: what actions you will allow. Each AWS service has its own set of actions.
- Resources: which resources you allow the action on.
- Effect: what the effect will be when the user requests access—either allow or deny.

Below is a sample snippet of a policy document that allows access to all users to Describe EC2 Instances. You can clearly see the Actions, Resources and Effect which define the policy document.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "ec2:Describe*",
      "Resource": "*"
    }
  ],
}
```

For more information on policies, please visit the url:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html

The correct answer is: User-based and Resource-based

[Feedback about this question and answer](#)

QUESTION 34

NOT ANSWERED

[Submit Feedback](#)

A company has resources hosted both on their on-premise network and in AWS. They want their IT administrators to access resources in both environments using their on-premise credentials which is stored in Active Directory. Which of the following can be used to fulfil this requirement?

Please select :

- ☐ A. Use Web Identity Federation

☐ D. Use AWS VPC

Your answer is incorrect.

Answer – B

The AWS Documentation provides the following information on SAML Federation

AWS supports identity federation with SAML 2.0 (Security Assertion Markup Language 2.0), an open standard that many identity providers (IdPs) use. This feature enables federated single sign-on (SSO), so users can log into the AWS Management Console or call the AWS APIs without you having to create an IAM user for everyone in your organization. By using SAML, you can simplify the process of configuring federation with AWS, because you can use the IdP's service instead of writing custom identity proxy code.

For more information on SAML Federation, please visit the below URL:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_providers_saml.html

The correct answer is: Use SAML Federation

[Feedback about this question and answer](#)

QUESTION 35

NOT ANSWERED

[Submit Feedback](#)

Amazon RDS DB snapshots and automated backups are stored in

Please select :

- ☐ A. Amazon S3
- ☐ B. Amazon ECS Volume
- ☐ C. Amazon RDS
- ☐ D. Amazon EMR

Your answer is incorrect.

Answer – A

Automated backups automatically back up your DB instance during a specific, user-definable backup window. Amazon RDS keeps these backups for a limited period that you can specify. You can later recover your database to any point in time during this backup retention period. And all of these backups get stored to S3 by default.

Option B is not correct, because that is used to store data for EC2 instances

Option C is not correct because an RDS cannot be used to store snapshots

Option D because EMR is used for storing and processing logs.

For more information on DB instance backup's , go to the url -

<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Overview.BackingUpAndRestoringAmazonRDSInstances.html>

The correct answer is: Amazon S3

[Feedback about this question and answer](#)

Using Amazon CloudWatch's Free Tier, what is the frequency of metric updates which you receive?

Please select :

- ☐ A. 5 minutes
- ☐ B. 500 milliseconds.
- ☐ C. 30 seconds
- ☐ D. 1 minute

Your answer is incorrect.

Answer – A

AWS free tier gives you access to the basic metrics for Cloudwatch and by default the basic package gives 5 minutes of aggregation of Cloudwatch metrics. If you need a further shorter interval, the you need to pay extra.

For more information on the free tier , please feel free to visit the url - <https://aws.amazon.com/free/>

The correct answer is: 5 minutes

[Feedback about this question and answer](#)

QUESTION 37

NOT ANSWERED

[Submit Feedback](#)

What option from the below lets you categorize your EC2 resources in different ways, for example, by purpose, owner, or environment.

Please select :

- ☐ A. wildcards
- ☐ B. pointers
- ☐ C. Tags
- ☐ D. special filters

Your answer is incorrect.

Answer – C

Please note that this is an important concept, if you are pursuing further certifications in AWS. Tags in aws are used to segregate resources in aws , which can also be used for cost reporting and billing purposes. In EC2 dashboard , there is a separate section for Tags.

The correct answer is: Tags

[Feedback about this question and answer](#)

QUESTION 38

NOT ANSWERED

[Submit Feedback](#)

What acts as a firewall that controls the traffic allowed to reach one or more instances?

Please select :

- ☐ A. Security group
- ☐ B. ACL
- ☐ C. IAM
- ☐ D. Private IP Addresses

Your answer is incorrect.

Answer – A

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. When you launch an instance in a VPC, you can assign the instance to up to five security groups. Security groups act at the instance level. Below is an example of a security group which has inbound rules. The below rule states that users can only SSH into EC2 instances that are attached to this security group.

Security Group: sg-74229a09

Description

Inbound

Outbound

Tags

Edit

| Type ⓘ | Protocol ⓘ | Port Range ⓘ | Source ⓘ |
|--------|------------|--------------|-----------|
| SSH | TCP | 22 | 0.0.0.0/0 |

For more information on Security Groups , please visit the url - http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html

The correct answer is: Security group

[Feedback about this question and answer](#)

QUESTION 39

NOT ANSWERED

[Submit Feedback](#)

You design an application that checks for new items in an S3 bucket once per hour. If new items exist, a message is added to an SQS queue. You have several EC2 instances which retrieve messages from the SQS queue, parse the file, and send you an email containing the relevant information from the file. You upload one test file to the bucket, wait a couple hours and find that you have hundreds of emails from the application. What is the most likely cause for this volume of email? Choose the correct answer from the options given below

Please select :

- ☐ A. This is expected behavior when using short polling because SQS does not guarantee that there will not be duplicate messages processed
- ☐ B. You can only have one EC2 instance polling the SQS queue at a time
- ☐ C. This is expected behavior when using long polling because SQS does not guarantee that there will not be duplicate messages processed
- ☐ D. Your application does not issue a delete command to the SQS queue after processing the message

Your answer is incorrect.

Answer – D

You need to ensure that after a message is processed in SQS, the message is deleted.

For more information on SQS, please visit the below url

<https://aws.amazon.com/sqs/faqs/>

Feedback about this question and answer

QUESTION 40

NOT ANSWERED

[Submit Feedback](#)

An application requires a minimum of 4 instances to run to ensure that it can cater to its users. You want to ensure fault tolerance and high availability. Which of the following is the best option.

Please select :

- ☐ A. Deploy 2 instances in each of 3 Availability Zones, add a load balancer and an Auto Scaling group to launch more instances if required.
- ☐ B. Deploy 2 instances in each of 2 Availability Zones, add a load balancer and an Auto Scaling group to launch more instances if required.
- ☐ C. Deploy 4 instances in one Availability Zone, add a load balancer and an Auto Scaling group to launch more instances if required.
- ☐ D. Deploy 1 instance in each of 3 Availability Zones, add a load balancer and an Auto Scaling group to launch more instances if required.

Your answer is incorrect.

Answer – A

Since there is a minimum of 4 instances required to run, if you deploy them in 3 AZ's and even if one AZ goes down , you will have at least 4 instances running.

For more information on fault tolerance and high availability, please visit the below URL:

- https://media.amazonwebservices.com/architecturecenter/AWS_ac_ra_ftha_04.pdf

The correct answer is: Deploy 2 instances in each of 3 Availability Zones, add a load balancer and an Auto Scaling group to launch more instances if required.

[Feedback about this question and answer](#)

QUESTION 41

NOT ANSWERED

[Submit Feedback](#)

Which of the following is true when it comes to hosting a database in VPC's using the AWS RDS service.

Please select :

- ☐ A. The VPC must have at least one subnet
- ☐ B. The VPC must have at least one subnet in one Availability Zone
- ☐ C. Your VPC must have at least one subnet in at least two of the Availability Zones
- ☐ D. None of the above

Your answer is incorrect.

Answer - C

deploy your DB instance. A subnet is a segment of a VPC's IP address range that you can specify and that lets you group instances based on your security and operational needs.

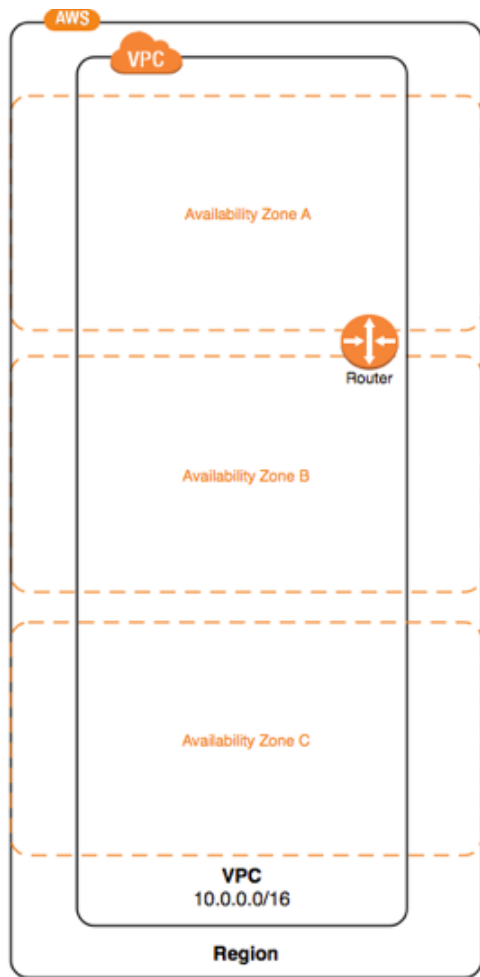
Working with a DB Instance in a VPC

Here are some tips on working with a DB instance in a VPC:

- Your VPC must have at least one subnet in at least two of the Availability Zones in the region where you want to deploy your DB instance. A subnet is a segment of a VPC's IP address range that you can specify and that lets you group instances based on your security and operational needs.
- If you want your DB instance in the VPC to be publicly accessible, you must enable the VPC attributes *DNS hostnames* and *DNS resolution*.
- Your VPC must have a DB subnet group that you create (for more information, see the next section). You create a DB subnet group by specifying the subnets you created. Amazon RDS uses that DB subnet group and your preferred Availability Zone to select a subnet and an IP address within that subnet to assign to your DB instance.
- Your VPC must have a VPC security group that allows access to the DB instance.
- The CIDR blocks in each of your subnets must be large enough to accommodate spare IP addresses for Amazon RDS to use during maintenance activities, including failover and compute scaling.

Few important points about VPC:

- When you create a VPC, it spans all the Availability Zones in the region.
- After creating a VPC, you can add one or more subnets in each Availability Zone.



Main route table

| Destination | Target |
|-------------|--------|
| 10.0.0.0/16 | local |

For more information on working with RDS instances , please refer to the below link:

- http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_VPC.WorkingWithRDSInstanceinaVPC.html

The correct answer is: Your VPC must have at least one subnet in at least two of the Availability Zones

[Feedback about this question and answer](#)

QUESTION 42

NOT ANSWERED

[Submit Feedback](#)

There are currently multiple applications hosted in a VPC. During monitoring it has been noticed that multiple port scans are coming in from a specific IP Address block. The internal security team has requested that all offending IP Addresses be denied for the next 24 hours. Which of the following is the best method to quickly and temporarily deny access from the specified IP Address's.

Please select :

- ☐ A. Create an AD policy to modify the Windows Firewall settings on all hosts in the VPC to deny access from the IP Address block.
- ☐ B. Modify the Network ACLs associated with all public subnets in the VPC to deny access from the IP Address block.
- ☐ C. Add a rule to all of the VPC Security Groups to deny access from the IP Address block.

Your answer is incorrect.

Answer – B

A network access control list (ACL) is an optional layer of security for your VPC that acts as a firewall for controlling traffic in and out of one or more subnets.

Option A and D are wrong because this is a tedious task and it only works for Windows systems. You need something that will work for Linux systems as well.

Option C is only adequate for EC2 instances, but you need rules that will apply to the whole subnet. Otherwise the task of having this done for all servers becomes a tedious task.

For more information on Network ACL's, please visit the URL

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLS.html

The correct answer is: Modify the Network ACLs associated with all public subnets in the VPC to deny access from the IP Address block.

[Feedback about this question and answer](#)

QUESTION 43

NOT ANSWERED

[Submit Feedback](#)

Which service allows one to issue temporary credentials in AWS? Choose one answer from the options below.

Please select :

- ☐ A. AWS SQS
- ☐ B. AWS STS
- ☐ C. AWS SES
- ☐ D. None of the above. You need to use a third party software to achieve this.

Your answer is incorrect.

Answer – B

You can use the AWS Security Token Service (AWS STS) to create and provide trusted users with temporary security credentials that can control access to your AWS resources. Temporary security credentials work almost identically to the long-term access key credentials that your IAM users can use.

Option A is wrong because this is the queuing service provided by AWS.

Option C is wrong because this is the emailing service provided by AWS.

Option D is wrong because there is a service which exists from AWS.

For more information on STS, please visit the below URL

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_temp.html

The correct answer is: AWS STS

[Feedback about this question and answer](#)

You have two Elastic Compute Cloud (EC2) instances inside a Virtual Private Cloud (VPC) in the same Availability Zone (AZ) but in different subnets. One instance is running a database and the other instance an application that will interface with the database. You want to confirm that they can talk to each other for your application to work properly. Which two things do we need to confirm in the VPC settings so that these EC2 instances can communicate inside the VPC? Choose 2 answers.

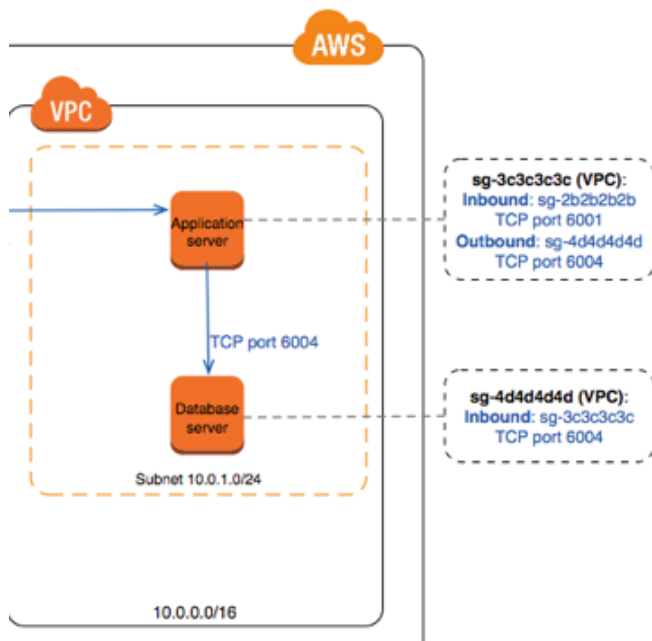
Please select :

- ☐ A. A network ACL that allows communication between the two subnets.
- ☐ B. Both instances are the same instance class and using the same Key-pair.
- ☐ C. That the default route is set to a NAT instance or internet Gateway (IGW) for them to communicate.
- ☐ D. Security groups are set to allow the application host to talk to the database on the right port/protocol.

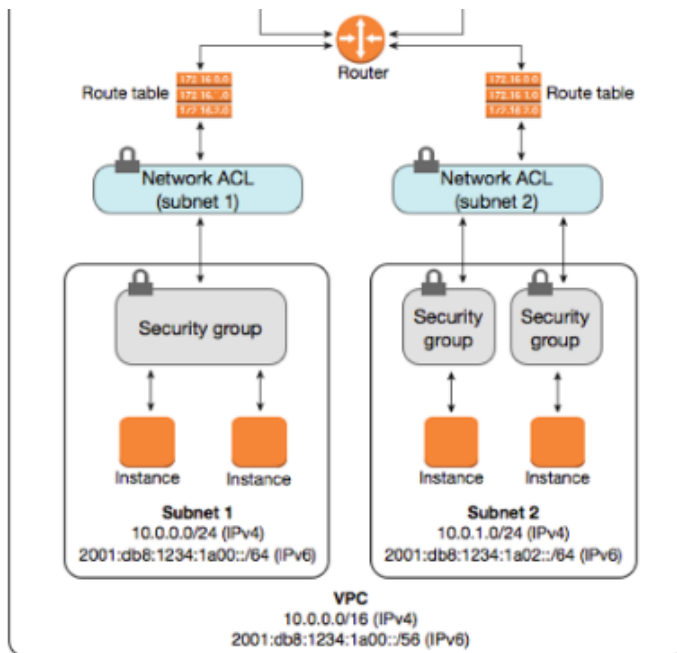
Your answer is incorrect.

Answer - A and D

When you design a web server and database server, the security groups must be defined so that the web server can talk to the database server. An example image from the AWS documentation is given below



Also when communicating between subnets you need to have the NACL's defined



Option B is wrong since the EC2 instances need not be of the same class or same key pair to communicate to each other.

Option C is wrong since there the NAT and Internet gateway is used for the subnet to communicate to the internet.

For more information on VPC and Subnets, please visit the below URL

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html

The correct answers are: A network ACL that allows communication between the two subnets., Security groups are set to allow the application host to talk to the database on the right port/protocol.

[Feedback about this question and answer](#)

| | | |
|-------------|--------------|---------------------------------|
| QUESTION 45 | NOT ANSWERED | Submit Feedback |
|-------------|--------------|---------------------------------|

As a solution architect, you have been asked to design a cloud service based on AWS and choose to use RRS on S3 instead of S3 standard storage type. In such a case what type of trade-offs do you have to build your application around?

Please select :

- ☐ A. With RRS you have to copy data and extract data which can take up to 3 hours.
- ☐ B. RRS only has 99.99% availability
- ☐ C. With RRS, you don't need to worry since AWS will take care of the durability of RRS.
- ☐ D. RRS only has 99.99% durability and you have to design automation around replacing lost objects

Your answer is incorrect.

Answer:- D

| | Standard | Standard - Infrequent Access | Reduced Redundancy Storage |
|------------|---------------|------------------------------|----------------------------|
| Durability | 99.999999999% | 99.999999999% | 99.99% |

Even though RRS has 99.99% availability, all storage types have the same availability, so it does not answer the question on the specific trade-offs for RRS.

| | Standard | Standard - Infrequent Access | Reduced Redundancy Storage |
|--------------|---------------|------------------------------|----------------------------|
| Durability | 99.999999999% | 99.999999999% | 99.99% |
| Availability | 99.99% | 99.9% | 99.99% |

For more information on RRS , visit the link

<https://aws.amazon.com/s3/reduced-redundancy/>

The correct answer is: RRS only has 99.99% durability and you have to design automation around replacing lost objects

[Feedback about this question and answer](#)

| | | |
|-------------|--------------|---------------------------------|
| QUESTION 46 | NOT ANSWERED | Submit Feedback |
|-------------|--------------|---------------------------------|

You are running a web-application on AWS consisting of the following components an Elastic Load Balancer (ELB) an Auto-Scaling Group of EC2 instances running Linux/PHP/Apache, and Relational DataBase Service (RDS) MySQL. Which security measures fall into AWS's responsibility?

Please select :

- ☐ A. Protect the EC2 instances against unsolicited access by enforcing the principle of least-privilege access
- ☐ B. Protect against IP spoofing or packet sniffing
- ☐ C. Assure all communication between EC2 instances and ELB is encrypted
- ☐ D. Install latest security patches on ELB, RDS and EC2 instances

Your answer is incorrect.

Answer – B

As per the shared responsibility shown below, the users are required to control the EC2 security via security groups and network access control layers.

The correct answer is: Protect against IP spoofing or packet sniffing

[Feedback about this question and answer](#)

QUESTION 47

NOT ANSWERED

[Submit Feedback](#)

What would happen to an RDS (Relational Database Service) multi-Availability Zone deployment of the primary DB instance fails?

Please select :

- ☐ A. The IP of the primary DB instance is switched to the standby DB instance
- ☐ B. The RDS (Relational Database Service) DB instance reboots
- ☐ C. A new DB instance is created in the standby availability zone
- ☐ D. The canonical name record (CNAME) is changed from primary to standby

Your answer is incorrect.

Answer – D

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure, and is engineered to be highly reliable. In case of an infrastructure failure (for example, instance hardware failure, storage failure, or network disruption), Amazon RDS performs an automatic failover to the standby, so that you can resume database operations as soon as the failover is complete.

And as per the AWS documentation, the cname is changed to the standby DB when the primary one fails.

Q: What happens during Multi-AZ failover and how long does it take?

Failover is automatically handled by Amazon RDS so that you can resume database operations as quickly as possible without administrative intervention. When failing over, Amazon RDS simply flips the canonical name record (CNAME) for your DB Instance to point at the standby, which is in turn promoted to become the new primary. We encourage you to follow best practices and implement database connection retry at the application layer.

For more information on Multi-AZ RDS, please visit the link –

<https://aws.amazon.com/rds/details/multi-az/>

The correct answer is: The canonical name record (CNAME) is changed from primary to standby

[Feedback about this question and answer](#)

QUESTION 48

NOT ANSWERED

[Submit Feedback](#)

installable from S3 and setup the ELB. which of the below mentioned AWS services meets the requirement for making an orderly deployment of the software?

Please select :

- ☐ A. AWS Elastic Beanstalk
- ☐ B. AWS Cloudfront
- ☐ C. AWS Cloudformation
- ☐ D. AWS DevOps

Your answer is incorrect.

Answer – A

The Elastic Beanstalk is an easy-to-use service for deploying and scaling web applications and services

We can simply upload code and Elastic Beanstalk automatically handles the deployment, from capacity provisioning, load balancing, auto-scaling to application health monitoring. Meanwhile we can retain full control over the AWS resources used in the application and can access the underlying resources at any time.

Launch LAMP stack with Elastic Beanstalk:

- <https://aws.amazon.com/getting-started/projects/launch-lamp-web-app/>

We can do it on AWS CloudFormation as well in a harder way and it will be less Native:

- <https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/deploying.applications.html>

The correct answer is: AWS Elastic Beanstalk

[Feedback about this question and answer](#)

QUESTION 49

NOT ANSWERED

[Submit Feedback](#)

In which of the following ways can you manage lambda functions. Choose all 3 correct answers.

Please select :

- ☐ A. Console
- ☐ B. CLI
- ☐ C. SDK
- ☐ D. EC2 Instances

Your answer is incorrect.

Answer – A,B and C

You can easily list, delete, update, and monitor your Lambda functions using the dashboard in the AWS Lambda console. You can also use the AWS CLI and AWS SDK to manage your Lambda functions. Visit the [Lambda Developers Guide](#) to learn more.

For more information on AWS Lambda, please visit the link:

- <https://aws.amazon.com/lambda/faqs/>

The correct answers are: Console, CLI, SDK

[Feedback about this question and answer](#)

| QUESTION 50 | NOT ANSWERED | Submit Feedback |
|-------------|--------------|-----------------|
|-------------|--------------|-----------------|

What is the maximum execution time for a Lambda function?

Please select :

- ☐ A. 3 seconds
- ☐ B. 300 seconds
- ☐ C. 24 hours
- ☐ D. No limit

Your answer is incorrect.

Answer – B

This is given in the aws documentation

Q: How long can an AWS Lambda function execute?

All calls made to AWS Lambda must complete execution within 300 seconds. The default timeout is 3 seconds, but you can set the timeout to any value between 1 and 300 seconds.

For more information on AWS Lambda, please visit the link

<https://aws.amazon.com/lambda/faqs/>

The correct answer is: 300 seconds

[Feedback about this question and answer](#)

| QUESTION 51 | NOT ANSWERED | Submit Feedback |
|-------------|--------------|-----------------|
|-------------|--------------|-----------------|

If you want to point a domain name to an AWS elastic load balancer in Route 53, how would you need to configure the record set? Choose the correct answer from the options below

Please select :

- ☐ C. Alias with a type "CNAME" record set
- ☐ D. Alias with a type "A" record set

Your answer is incorrect.

Answer – D

Thanks! for asking clarification. Yes, You need to configure ALIAS record for ELB but it should point to A record. You can find details in below AWS document

<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-to-elb-load-balancer.html>

This is given in the aws documentation

Type

The DNS record type. For more information, see [Supported DNS Resource Record Types](#).

Select the applicable value based on the AWS resource for which you're creating a resource record set:

CloudFront distribution

Select **A — IPv4 address**.

If IPv6 is enabled for the distribution, create two resource record sets, one with a value of **A — IPv4 address** for **Type**, and one with a value of **AAAA — IPv6 address**.

Elastic Beanstalk environment that has regionalized subdomains

Select **A — IPv4 address**

ELB load balancer

Select **A — IPv4 address** or **AAAA — IPv6 address**

For more information on Route53, please visit the link

<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/resource-record-sets-values-alias.html>

The correct answer is: Alias with a type "A" record set

[Feedback about this question and answer](#)

QUESTION 52

NOT ANSWERED

[Submit Feedback](#)

Which of the following statement is false with regards to the AWS Simple Queue Service?

Please select :

- ☐ A. Standard queues provide at-least-once delivery, which means that each message is delivered at least once
- ☐ B. Both FIFO queues and Standard queues preserve the order of messages
- ☐ C. Amazon SQS can help you build a distributed application with decoupled components
- ☐ D. FIFO queues provide exactly-once processing

Your answer is incorrect.

For more information on standard queues, please visit the below URL

- <https://aws.amazon.com/sqs/faqs/>

The correct answer is: Both FIFO queues and Standard queues preserve the order of messages

[Feedback about this question and answer](#)

QUESTION 53

NOT ANSWERED

[Submit Feedback](#)

A company wants to implement a hybrid architecture where it wants to connect VPC's in its account to its on-premise architecture. Which of the following can be used to create a private connection between the Company's on-premise architecture and the VPC's hosted in AWS

Please select :

- ☐ A. AWS VPN
- ☐ B. Route53
- ☐ C. ClassicLink
- ☐ D. AWS Direct Link

Your answer is incorrect.

Answer – A

This is provided in the AWS documentation

You can connect your VPC to remote networks by using a VPN connection. The following are some of the connectivity options available to you.

| VPN connectivity option | Description |
|-------------------------|--|
| AWS hardware VPN | You can create an IPsec, hardware VPN connection between your VPC and your remote network. On the AWS side of the VPN connection, a <i>virtual private gateway</i> provides two VPN endpoints for automatic failover. You configure your <i>customer gateway</i> , which is the physical device or software application on the remote side of the VPN connection. For more information, see Adding a Hardware Virtual Private Gateway to Your VPC , and the Amazon VPC Network Administrator Guide . |

For more information on VPN Connections, please visit the below URL:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html>

The correct answer is: AWS VPN

[Feedback about this question and answer](#)

Which of the following statements are true when it comes to EBS volumes and snapshots. Choose all that apply.

Please select :

- ☐ A. You can change the size of an EBS volume.
- ☐ B. If you have an unencrypted volume, you can still create an encrypted snapshot from it.
- ☐ C. The volume change size can also happen when it is attached to an instance.
- ☐ D. The volume change size can only happen if the volume is detached from an instance.

Your answer is incorrect.

Answer – A and C

If your Amazon EBS volume is attached to a current generation EC2 instance type, you can increase its size, change its volume type, or (for an io1 volume) adjust its IOPS performance, all without detaching it.

For more information on changing the volume size, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-expand-volume.html>

For more information on changing the EBS encryption, please visit the link

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html>

The correct answers are: You can change the size of an EBS volume., The volume change size can also happen when it is attached to an instance.

The correct answers are: You can change the size of an EBS volume., The volume change size can also happen when it is attached to an instance.

[Feedback about this question and answer](#)

QUESTION 55

NOT ANSWERED

[Submit Feedback](#)

You have set up a CloudFront distribution but find that instead of each edge location serving up objects that should be cached, your application's origins are being hit for each request. What could be a possible cause of this behavior? Choose the correct answer from the options below

Please select :

- ☐ A. The requested content has never been requested before
- ☐ B. The objects file size are too large
- ☐ C. The cache expiration time is set to a low value
- ☐ D. You didn't configure the objects with a X.509 certificate

You can control how long your objects stay in a CloudFront cache before CloudFront forwards another request to your origin. Reducing the duration allows you to serve dynamic content. Increasing the duration means your users get better performance because your objects are more likely to be served directly from the edge cache. A longer duration also reduces the load on your origin.

For more information on changing the volume encryption, please visit the link

<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html>

The correct answer is: The cache expiration time is set to a low value

[Feedback about this question and answer](#)

QUESTION 56

NOT ANSWERED

[Submit Feedback](#)

You are using an EC2 instance that is backed by an S3-based AMI. You are planning on terminating that instance. When the instance is terminated, what happens to the data on the root volume?

Please select :

- ☐ A. Data is automatically saved as an EBS snapshot.
- ☐ B. Data is automatically saved as an EBS volume.
- ☐ C. Data is unavailable until the instance is restarted.
- ☐ D. Data is automatically deleted.

Your answer is incorrect.

Answer – D

The AWS documentation mentions the following

The data in an instance store persists only during the lifetime of its associated instance. If an instance reboots (intentionally or unintentionally), data in the instance store persists. However, data in the instance store is lost under the following circumstances:

- The underlying disk drive fails
- The instance stops
- The instance terminates

For more information on Instance store AMI's, please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html>

The correct answer is: Data is automatically deleted.

[Feedback about this question and answer](#)

QUESTION 57

NOT ANSWERED

[Submit Feedback](#)

would you use? Choose an answer from the options given below

Please select :

- ☐ A. Latency
- ☐ B. Failover
- ☐ C. Weighted
- ☐ D. Geolocation

Your answer is incorrect.

Answer – C

Use the weighted routing policy when you have multiple resources that perform the same function (for example, web servers that serve the same website) and you want Amazon Route 53 to route traffic to those resources in proportions that you specify (for example, one quarter to one server and three quarters to the other).

For more information on the routing policy, please visit the link

<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>

The correct answer is: Weighted

[Feedback about this question and answer](#)

QUESTION 58

NOT ANSWERED

[Submit Feedback](#)

An application in AWS is currently running in the Singapore region. You have been asked to implement disaster recovery. So if the application goes down in the Singapore region, it has to be started in the Asia region. Your application relies on pre-built AMIs. As part of your disaster recovery strategy, which of the below points should you consider.

Please select :

- ☐ A. Nothing, because all AMI's de default are available in any region as long as it is created within the same account
- ☐ B. Copy the AMI from the Singapore region to the Asia region. Modify the Auto Scaling groups in the backup region to use the new AMI ID in the backup region
- ☐ C. Modify the image permissions and share the AMI to the Asia region.
- ☐ D. Modify the image permissions to share the AMI with another account, then set the default region to the backup region

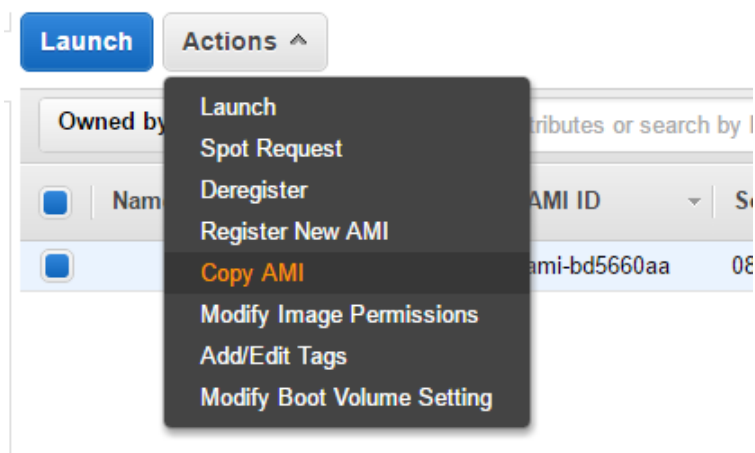
Your answer is incorrect.

Answer:- B

If you need an AMI across multiple regions , then you have to copy the AMI across regions. Note that by default AMI's that you have created will not be available across all regions. So option A is automatically invalid.

Next you can share AMI's with other users, but they will not be available across regions. So option C and D is invalid. You have to copy the AMI across regions.

Step 2) Once the Image has been created, go to the AMI section in the EC2 dashboard and click on the Copy AMI option.



Step 3) In the next screen , you can specify where to copy the AMI to.

AMI ami-bd5660aa will be copied to a new AMI. Set the new AMI settings below.

Destination region* Asia Pacific (Tokyo) ▼

Name Demo

Description [Copied ami-bd5660aa from us-east-1] Demo

Encryption ☐ Encrypt target EBS snapshots ⓘ

Cancel Copy AMI

For the entire details to copy AMI's , please visit the link -

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>

The correct answer is: Copy the AMI from the Singapore region to the Asia region. Modify the Auto Scaling groups in the backup region to use the new AMI ID in the backup region

[Feedback about this question and answer](#)

QUESTION 59

NOT ANSWERED

[Submit Feedback](#)

In the basic monitoring package for RDS, Amazon CloudWatch provides the following metrics. Choose three correct option.

Please select :

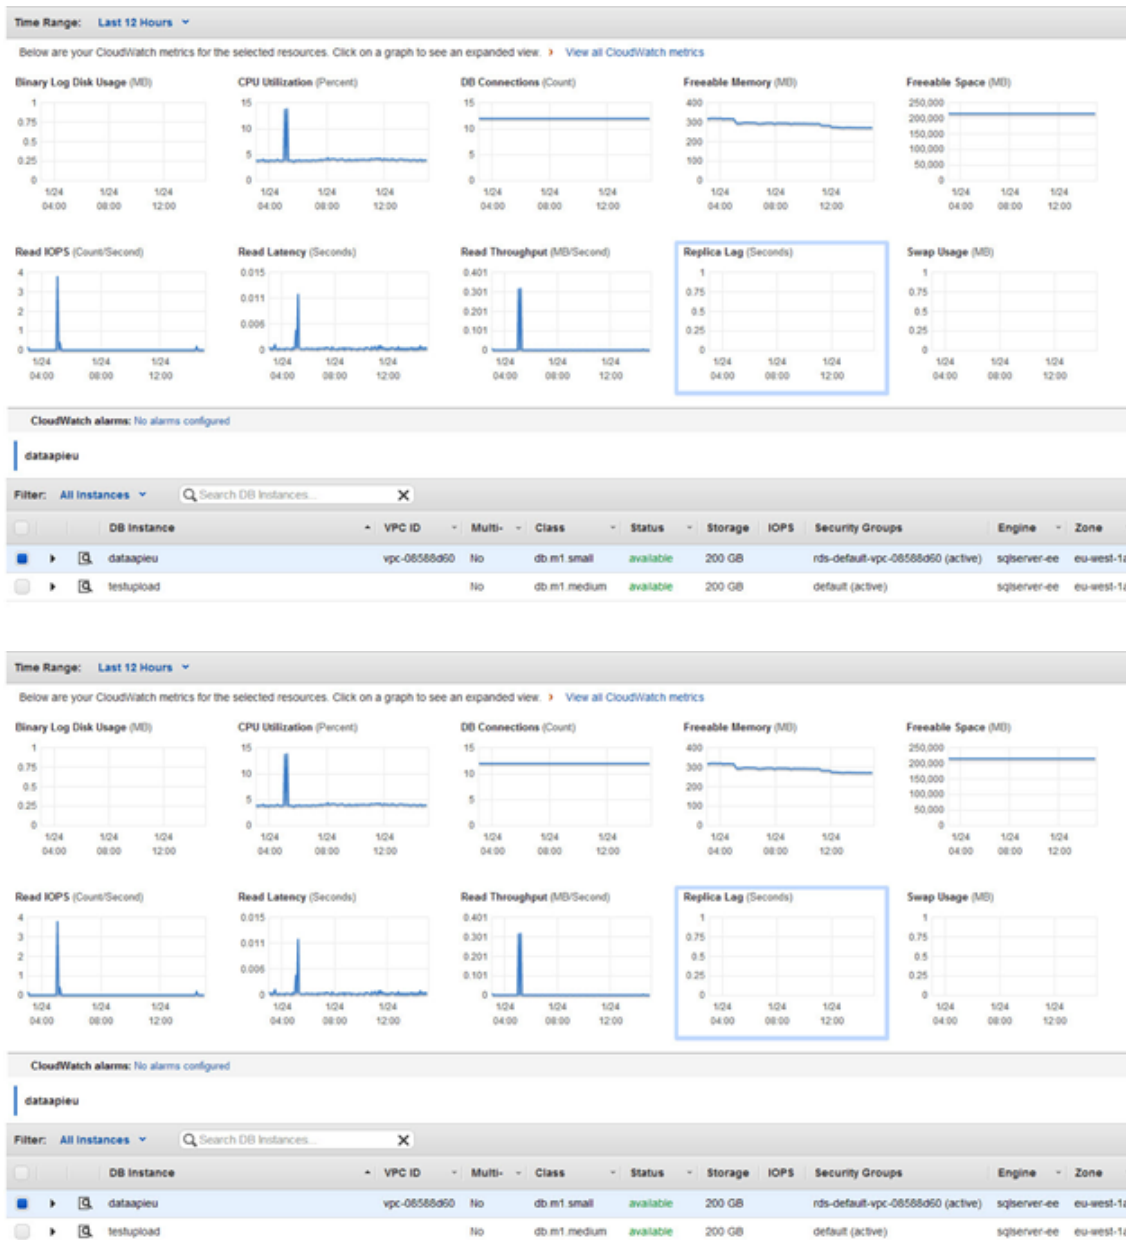
- ☐ A. Database visible metrics such as number of connections
- ☐ B. Disk OPS metrics
- ☐ C. Database memory usage
- ☐ D. Web service visible metrics such as number failed transaction requests

Your answer is incorrect.

Answer:- A, B and C

As RDS Instance is completely managed by AWS and user doesn't have access Operating System metrics, So it is logical for AWS to provide us those metrics.

Please refer to AWS documentation http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Monitoring.html



The correct answers are: Database visible metrics such as number of connections, Disk OPS metrics, Database memory usage

Feedback about this question and answer

QUESTION 60

NOT ANSWERED

[Submit Feedback](#)

What is the maximum size of an EBS Provisioned IOPS SSD volume? Choose the correct option.

Please select :

- ☐ C. 4Gib
- ☐ D. 16TB

Your answer is incorrect.

Answer:- B

The maximum size for EBS provisioned IOPS volume allowed is 16384 GiB which 16 TiB. See error while trying to create volume more than available size :

Create Volume

Volume Type

Provisioned IOPS SSD (IO1) ▼



Size (GiB)

20000

(Min: 4 GiB, Max: 16384 GiB)



⚠ Size can't exceed 16384 GiB

IOPS

100

(Min: 100 IOPS, Max: 20000 IOPS)



Availability Zone*

us-east-1a ▼



Throughput (MB/s)

Not applicable



Snapshot ID

Select a snapshot ▼



Encryption

☐

Encrypt this volume



For more information on EBS volume types, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>

The correct answer is: 16TiB

[Feedback about this question and answer](#)

Company

[About Us](#)

[Discussions](#)

[Blog](#)

Follow Us



Support

contact us @ support@whizlabs.com