

## AWS CSAA Practice Tests

[Home](#) / [My courses](#) / [AWS CSAA Practice Tests](#) / [FULL TEST\(S\)](#) / [New Practice Test VI](#)

**Started on** Saturday, 10 February 2018, 12:26 AM

**State** Finished

**Completed on** Saturday, 10 February 2018, 12:26 AM

**Time taken** 16 secs

**Grade** 0.00 out of 60.00 (0%)

**Result** FAIL

QUESTION 1

NOT ANSWERED

[Submit Feedback](#)

A company wants to store data that is not frequently accessed. What is the best and cost efficient solution that should be considered?

Please select :

- ☐ A. Amazon Storage Gateway
- ☐ B. Amazon Glacier
- ☐ C. Amazon EBS
- ☐ D. Amazon S3

**Your answer is incorrect.**

Answer: B – Since the data is not required to be accessed frequently, the data can be stored on Amazon glacier for cheaper storage. Remember that the recovery time for getting data from Glacier is from 3-5 hours.

You can look at the FAQ section of aws glacier - <https://aws.amazon.com/glacier/faqs/>

customers to offload the administrative burdens of operating and scaling storage to AWS, so that they don't have to worry about capacity planning, hardware provisioning, data replication, hardware failure detection and repair, or time-consuming hardware migrations.

The correct answer is: Amazon Glacier

[Feedback about this question and answer](#)

QUESTION 2	NOT ANSWERED	<a href="#">Submit Feedback</a>
------------	--------------	---------------------------------

You have an EC2 instance that is transferring data from S3 in the same region. The project sponsor is worried about the cost of the infrastructure. What can you do to convince him that you have a cost effective solution.

Please select :

- ☐ A. You are going to be hosting only 4 instances, so you are minimizing on cost.
- ☐ B. There is no cost for transferring data from EC2 to S3 if they are in the same region.
- ☐ C. AWS provides a discount if you transfer data from EC2 to S3 if they are in the same region.
- ☐ D. Both EC2 and S3 are in the same availability zone, so you can save via consolidated billing.

**Your answer is incorrect.**

Answer – B – Please note that there is no cost when data is transferred from EC2 to S3 if they are in the same region. This is very important for an AWS Solution Architect to know.

The correct answer is: There is no cost for transferring data from EC2 to S3 if they are in the same region.

[Feedback about this question and answer](#)

QUESTION 3	NOT ANSWERED	<a href="#">Submit Feedback</a>
------------	--------------	---------------------------------

Which services allow the customer to retain full administrative privileges of the underlying EC2 instances?

Please select :

- ☐ A. Amazon Relational Database Service
- ☐ B. Amazon Elastic Map Reduce
- ☐ C. Amazon ElastiCache
- ☐ D. Amazon DynamoDB

In Amazon EMR , you have the ability to work with the underlying instances wherein the EMR service allows you to associate the EC2 Key pair with the launched instances. This is also given in the AWS documentation.

### Create an Amazon EC2 Key Pair and PEM File

Amazon EMR uses an Amazon EC2 key pair to ensure that you alone have access to the instances that you launch. The PEM file associated with this key pair is required to `ssh` directly to the master node of the cluster.

For more information on the access to EMR nodes , please visit the below URL

<http://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-access-ssh.html>

The correct answer is: Amazon Elastic Map Reduce

[Feedback about this question and answer](#)

QUESTION 4

NOT ANSWERED

[Submit Feedback](#)

Bucket names must be unique across all S3.

Please select :

- ☐ A. True
- ☐ B. False

**Your answer is incorrect.**

Answer – A

Bucket names must be unique across all regions. Let's say you have created a bucket names devtoolslogging in the Singapore region.

A bucket is a container for objects stored in Amazon S3. When creating a bucket, you can choose a Region to optimize for latency, minimize costs, or address regulatory requirements. For more information regarding bucket naming conventions, please visit the [Amazon S3 documentation](#).

Bucket Name: devtoolslogging

Region: Singapore

Set Up Logging >

Create

Cancel

Now if you want to create a bucket of the same name in the Oregon region, you will get an error that the bucket already exists.

### Create a Bucket - Select a Bucket Name and Region

Cancel

**That bucket already exists, you own it!**

A bucket is a container for objects stored in Amazon S3. When creating a bucket, you can choose a Region to optimize for latency, minimize costs, or address regulatory requirements. For more information regarding bucket naming conventions, please visit the [Amazon S3 documentation](#).

Bucket Name: devtoolslogging

Region: Oregon

Set Up Logging >

Create

Cancel

The correct answer is: True

[Feedback about this question and answer](#)

QUESTION 5

NOT ANSWERED

[Submit Feedback](#)

A customer has enabled website hosting on a bucket named "devtoolslogging" in the Singapore region. What website URL is assigned to your bucket?

Please select :

- ☐ A. devtoolslogging.s3-website-ap-southeast-1.amazonaws.com

- ☐ D. devtoolslogging.ap-southeast-1.amazonaws.com

**Your answer is incorrect.**

Answer – A – You have the chance to enable static web site hosting for S3 buckets. This can be done via the properties option for the bucket. The end point of the bucket for static hosting will also be configured.

The correct answer is: devtoolslogging.s3-website-ap-southeast-1.amazonaws.com

[Feedback about this question and answer](#)

**QUESTION 6**

**NOT ANSWERED**

[Submit Feedback](#)

As a solutions architect, it is your job to design for high availability and fault tolerance. Company-A is utilizing Amazon S3 to store large amounts of file data. What steps would you take to ensure that if an availability zone was lost due to a natural disaster your files would still be in place and accessible

Please select :

- ☐ A. Copy the S3 bucket to an EBS optimized backed EC2 instance
- ☐ B. Amazon S3 is highly available and fault tolerant by design and requires no additional configuration
- ☐ C. Enable AWS Storage Gateway using gateway-stored setup
- ☐ D. None of the above

**Your answer is incorrect.**

Answer – B – AWS S3 is already highly available and fault tolerant. This is very clearly mentioned in its FAQ's - <https://aws.amazon.com/s3/faqs/>

Amazon S3 Standard and Standard - IA are designed to provide 99.999999999% durability of objects over a given year. This durability level corresponds to an average annual expected loss of 0.000000001% of objects. For example, if you store 10,000 objects with Amazon S3, you can on average expect to incur a loss of a single object once every 10,000,000 years. In addition, Amazon S3 is designed to sustain the concurrent loss of data in two facilities.

As with any environments, the best practice is to have a backup and to put in place safeguards against malicious or accidental users errors. For S3 data, that best practice includes secure access permissions, Cross-Region Replication, versioning and a functioning, regularly tested backup.

**Q: How is Amazon S3 designed to achieve 99.999999999% durability?**

Amazon S3 Standard and Standard - IA redundantly stores your objects on multiple devices across multiple facilities in an Amazon S3 Region. The service is designed to sustain concurrent device failures by quickly detecting and repairing any lost redundancy. When processing a request to store data, the service will redundantly store your object across multiple facilities before returning SUCCESS. Amazon S3 also regularly verifies the integrity of your data using checksums.

The correct answer is: Amazon S3 is highly available and fault tolerant by design and requires no additional configuration

[Feedback about this question and answer](#)

QUESTION 7

NOT ANSWERED

[Submit Feedback](#)

Which of the following are characteristics of Amazon VPC subnets? Choose 2 answers from the options given below

Please select :

- ☐ A. Each subnet spans at least 2 Availability Zones to provide a high-availability environment.
- ☐ B. Each subnet maps to a single Availability Zone.
- ☐ C. CIDR block mask of /25 is the smallest range supported.
- ☐ D. By default, all subnets can route between each other, whether they are private or public.
- ☐ E. Instances in a private subnet can communicate with the Internet only if they have an Elastic IP.

**Your answer is incorrect.**

Answer – B, D. Please see the below for further justification.

A subnet can only map to one availability zone. So from options A and B, B is correct.

When you create a CIDR block, the least allowable is /28, so option C is wrong.

Create Subnet

Use the CIDR format to specify your subnet's IP address block (e.g., 10.0.0.0/24). Note that block sizes must be between a /16 netmask and /28 netmask. Also, note that a subnet can be the same size as your VPC.

Name tag

VPC

vpc-6dcc550a

Availability Zone

us-east-1c

CIDR block

Cancel

Yes, Create

Option E is wrong because EC2 instances in a private subnet will not be able to route anything on the internet even if they have an elastic IP.

The correct answers are: Each subnet maps to a single Availability Zone., By default, all subnets can route between each other, whether they are private or public.

[Feedback about this question and answer](#)

QUESTION 8

NOT ANSWERED

[Submit Feedback](#)

What are the different options available when creating a VPC using the VPC wizard? Please choose all options that apply.

Please select :

- ☐ C. VPC with Public and Private Subnets and Hardware VPN Access
- ☐ D. VPC with default settings

**Your answer is incorrect.**

Answer – B and C. When you launch the VPC wizard, you will get the below options in the VPC wizard.

### Step 1: Select a VPC Configuration

**VPC with a Single Public Subnet**

VPC with Public and Private Subnets

VPC with Public and Private Subnets and Hardware VPN Access

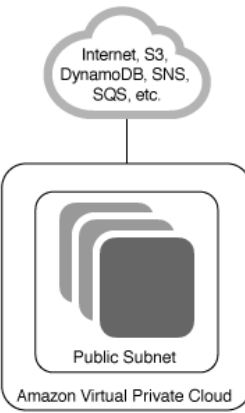
VPC with a Private Subnet Only and Hardware VPN Access

Your instances run in a private, isolated section of the AWS cloud with direct access to the Internet. Network access control lists and security groups can be used to provide strict control over inbound and outbound network traffic to your instances.

**Creates:**

A /16 network with a /24 subnet. Public subnet instances use Elastic IPs or Public IPs to access the Internet.

Select



The diagram illustrates a VPC configuration. At the top, a cloud icon represents the Internet, with text listing services like S3, DynamoDB, SNS, and SQS. A line connects this cloud to a box labeled 'Public Subnet'. This box is part of a larger container labeled 'Amazon Virtual Private Cloud'.

The correct answers are: VPC with Public and Private Subnets, VPC with Public and Private Subnets and Hardware VPN Access

[Feedback about this question and answer](#)

**QUESTION 9**

**NOT ANSWERED**

[Submit Feedback](#)

In AWS, which security aspects are the customer's responsibility? Choose 4 answers

Please select :

- ☐ A. Security Group and ACL (Access Control List) settings
- ☐ B. Decommissioning storage devices
- ☐ C. Patch management on the EC2 instance's operating system
- ☐ D. Life-cycle management of IAM credentials
- ☐ E. Controlling physical access to compute resources
- ☐ F. Encryption of EBS (Elastic Block Storage) volumes

**Your answer is incorrect.**

Answer – A, C, D and F.



The correct answers are: Security Group and ACL (Access Control List) settings, Patch management on the EC2 instance's operating system, Life-cycle management of IAM credentials, Encryption of EBS (Elastic Block Storage) volumes

[Feedback about this question and answer](#)

QUESTION 10

NOT ANSWERED

[Submit Feedback](#)

When an EC2 EBS-backed (EBS root) instance is stopped, what happens to the data on any ephemeral store volumes?

Please select :

- ☐ A. Data is automatically saved in an EBS volume.
- ☐ B. Data is unavailable until the instance is restarted.
- ☐ C. Data will be deleted and will no longer be accessible.
- ☐ D. Data is automatically saved as an EBS snapshot.

**Your answer is incorrect.**

Answer – C – ephemeral is temporary storage that is always deleted when an instance is restarted in aws.

When you stop or terminate an instance, every block of storage in the instance store is reset. Therefore, your data cannot be accessed through the instance store of another instance.

Data on the EBS volume is LOST only if the Root Volume is EBS backed and the Delete On Termination flag is checked (Checked by default)

The correct answer is: Data will be deleted and will no longer be accessible.

[Feedback about this question and answer](#)

QUESTION 11

NOT ANSWERED

[Submit Feedback](#)

SQS provides a timeout which is a period of time during which Amazon SQS prevents other consuming components from receiving and processing. What is this time period called?

Please select :

- ☐ A. Component Timeout
- ☐ B. Visibility Timeout
- ☐ C. Processing Timeout
- ☐ D. Receiving Timeout

**Your answer is incorrect.**

Answer - B

Please refer to the AWS SQS FAQ section - <https://aws.amazon.com/sqs/faqs/>

The correct answer is: Visibility Timeout

[Feedback about this question and answer](#)

QUESTION 12

NOT ANSWERED

[Submit Feedback](#)

You are currently hosting an infrastructure and most of the EC2 instances are near 90 – 100% utilized. What is the type of EC2 instances you would utilize to ensure costs are minimized? Assume that the EC2 instance will be running continuously throughout the year.

Please select :

- ☐ A. Reserved instances

- ☐ D. Regular instances

**Your answer is incorrect.**

Answer: A – When you have instances that will be used continuously and throughout the year, the best option is to buy reserved instances. By buying reserved instances, you are actually allocated an instance for the entire year or the duration you specify with a reduced cost.

To understand more on reserved instances, please visit the link:

- <https://aws.amazon.com/ec2/pricing/reserved-instances/>
- <https://blog.cloudability.com/maximizing-cost-savings-aws-reserved-instances/>
- <https://awsinsider.net/articles/2017/03/21/controlling-aws-costs.aspx>

The correct answer is: Reserved instances

[Feedback about this question and answer](#)

QUESTION 13

NOT ANSWERED

[Submit Feedback](#)

What is the ability provided by AWS to enable fast, easy, and secure transfers of files over long distances between your client and your Amazon S3 bucket.

Please select :

- ☐ A. File Transfer
- ☐ B. HTTP Transfer
- ☐ C. Transfer Acceleration
- ☐ D. S3 Acceleration

**Your answer is incorrect.**

Answer - C

Please refer to the AWS S3 FAQ section - <https://aws.amazon.com/s3/faqs/>

**Q. What is Transfer Acceleration?**

Amazon S3 Transfer Acceleration enables fast, easy, and secure transfers of files over long distances between your client and your Amazon S3 bucket. Transfer Acceleration leverages Amazon CloudFront's globally distributed AWS Edge Locations. As data arrives at an AWS Edge Location, data is routed to your Amazon S3 bucket over an optimized network path.

The correct answer is: Transfer Acceleration

## QUESTION 14

NOT ANSWERED

[Submit Feedback](#)

What is one key difference between an Amazon EBS-backed and an instance-store backed instance?

Please select :

- ☐ A. Amazon EBS-backed instances can be stopped and restarted.
- ☐ B. Instance-store backed instances can be stopped and restarted.
- ☐ C. Auto scaling requires using Amazon EBS-backed instances.
- ☐ D. Virtual Private Cloud requires EBS backed instances.

**Your answer is incorrect.**

Answer A

Amazon EBS-backed instances can be stopped and restarted. So we can say Instance-store backed instances cannot be restarted.

- Please see the url for the key differences between EBS and instance store volumes - <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>

The correct answer is: Amazon EBS-backed instances can be stopped and restarted.

[Feedback about this question and answer](#)

## QUESTION 15

NOT ANSWERED

[Submit Feedback](#)

You have an application hosted in AWS. The logs from the application are sent to Cloudwatch. The application has recently been encountering some errors. A patch needs to be developed for the error to be rectified. For the moment you need to automate the restart of the server whenever the error occurs. How can you achieve this?

Please select :

- ☐ A. Check the Cloudwatch logs for the error keywords , create an alarm and then restart the server
- ☐ B. Create a cloudwatch metric which looks at the CPU utilization and then restarts the server
- ☐ C. Create a cloudwatch metric which looks at the Memory utilization and then restarts the server
- ☐ D. Check the Cloudwatch logs for the error keywords, then send a notification to SQS to restart the server

**Your answer is incorrect.**

You can use CloudWatch Logs to monitor applications and systems using log data. For example, CloudWatch Logs can track the number of errors that occur in your application logs and send you a notification whenever the rate of errors exceeds a threshold you specify. CloudWatch Logs uses your log data for monitoring; so, no code changes are required. For example, you can monitor application logs for specific literal terms (such as "NullPointerException") or count the number of occurrences of a literal term at a particular position in log data (such as "404" status codes in an Apache access log). When the term you are searching for is found, CloudWatch Logs reports the data to a CloudWatch metric that you specify. Log data is encrypted while in transit and while it is at rest.

For more information on Cloudwatch logs , please visit the below URL:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/WhatIsCloudWatchLogs.html>

### Create Alarms to Stop, Terminate, Reboot, or Recover an Instance

Using Amazon CloudWatch alarm actions, you can create alarms that automatically stop, terminate, reboot, or recover your EC2 instances. You can use the stop or terminate actions to help you save money when you no longer need an instance to be running. You can use the reboot and recover actions to automatically reboot those instances or recover them onto new hardware if a system impairment occurs.

There are a number of scenarios in which you might want to automatically stop or terminate your instance. For example, you might have instances dedicated to batch payroll processing jobs or scientific computing tasks that run for a period of time and then complete their work. Rather than letting those instances sit idle (and accrue charges), you can stop or terminate them which can help you to save money. The main difference between using the stop and the terminate alarm actions is that you can easily restart a stopped instance if you need to run it again later, and you can keep the same instance ID and root volume. However, you cannot restart a terminated instance. Instead, you must launch a new instance.

You can add the stop, terminate, reboot, or recover actions to any alarm that is set on an Amazon EC2 per-instance metric, including basic and detailed monitoring metrics provided by Amazon CloudWatch (in the AWS/EC2 namespace), in addition to any custom metrics that include the "InstanceId=" dimension, as long as the InstanceId value refers to a valid running Amazon EC2 instance.

To set up a CloudWatch alarm action that can reboot, stop, or terminate an instance, you must use a service-linked IAM role, *AWSServiceRoleForCloudWatchEvents*. The *AWSServiceRoleForCloudWatchEvents* IAM role enables AWS to perform alarm actions on your behalf.

The correct answer is: Check the Cloudwatch logs for the error keywords , create an alarm and then restart the server

[Feedback about this question and answer](#)

QUESTION 16

NOT ANSWERED

[Submit Feedback](#)

A company wants to utilize aws storage. For them low storage cost is paramount, the data is rarely retrieved, and data retrieval times of several hours are acceptable for them. What is the best storage option to use?

Please select :

- ☐ A. Glacier
- ☐ B. Reduced Redundancy Storage
- ☐ C. EBS backed storage connected to EC2

Your answer is incorrect.

Answer – A.

With the above requirements, the best option is to opt for Amazon Glacier. Please refer to the Glacier FAQ's

<https://aws.amazon.com/glacier/faqs/>

**Q: How should I choose between Amazon Glacier and Amazon Simple Storage Service (Amazon S3)?**

Amazon S3 is a durable, secure, simple, and fast storage service designed to make web-scale computing easier for developers. Use Amazon S3 if you need low latency or frequent access to your data. Use Amazon Glacier if low storage cost is paramount, your data is rarely retrieved, and data retrieval times of several hours are acceptable.

The correct answer is: Glacier

[Feedback about this question and answer](#)

QUESTION 17

NOT ANSWERED

[Submit Feedback](#)

A client application requires operating system privileges on a relational database server. What is an appropriate configuration for a highly available database architecture?

Please select :

- ☐ A. Standalone Amazon EC2 instance
- ☐ B. Amazon RDS in a Multi-AZ configuration
- ☐ C. Amazon EC2 instances in a replication configuration utilizing a Single Availability Zone
- ☐ D. Amazon EC2 instances in a replication configuration utilizing two different Availability Zones

Your answer is incorrect.

Answer – D

You can not access OS of RDS Databases, as RDS is fully managed service by AWS. In case a customer wants to have access to OS for their Database for more granular control or other compliance reason, then they can install their Database engine in EC2 instance. In choice D , DB needs to be installed in EC2 for OS access with replication to support failover.

Please follow below link for reference, which shows steps to install and configure Oracle in EC2 instance <https://oracle-base.com/articles/vm/aws-ec2-installation-of-oracle>

Since the client wants privilege on the RDS, option B is not valid. Since there is a requirement for highly availability, you cannot have just one AZ and one EC2 instance. Hence D is the right answer.

For more information, please read the below link:

[http://docs.aws.amazon.com/dms/latest/userguide/CHAP\\_Introduction.ReplicationInstance.html](http://docs.aws.amazon.com/dms/latest/userguide/CHAP_Introduction.ReplicationInstance.html)

The correct answer is: Amazon EC2 instances in a replication configuration utilizing two different Availability Zones

[Feedback about this question and answer](#)

QUESTION 18

NOT ANSWERED

[Submit Feedback](#)

Per the AWS Acceptable Use Policy, penetration testing of EC2 instances:

Please select :

- ☐ A. May be performed by AWS, and will be performed by AWS upon customer request.
- ☐ B. May be performed by AWS, and is periodically performed by AWS.
- ☐ C. Are expressly prohibited under all circumstances.
- ☐ D. May be performed by the customer on their own instances with prior authorization from AWS.
- ☐ E. May be performed by the customer on their own instances, only if performed from EC2 instances.

**Your answer is incorrect.**

Answer – D. You need to take prior authorization from AWS before doing a penetration test on EC2 Instances. Please refer to the below url for more details.

<https://aws.amazon.com/security/penetration-testing/>

The correct answer is: May be performed by the customer on their own instances with prior authorization from AWS.

[Feedback about this question and answer](#)

QUESTION 19

NOT ANSWERED

[Submit Feedback](#)

How can you secure data at rest on an EBS volume?

Please select :

- ☐ A. Attach the volume to an instance using EC2's SSL interface.
- ☐ B. Write the data randomly instead of sequentially.
- ☐ C. Encrypt the volume using the S3 server-side encryption service.

**Your answer is incorrect.**

Answer – E. In order to secure data at rest on an EBS volume, you either have to encrypt the volume when it is being created or encrypt the data after the volume is created.

The correct answer is: Use an encrypted file system on top of the EBS volume.

[Feedback about this question and answer](#)

**QUESTION 20**

**NOT ANSWERED**

[Submit Feedback](#)

Which aws service is used to monitor all API calls to AWS

Please select :

- ☐ A. Amazon SES
- ☐ B. Amazon Cloudtrail
- ☐ C. Amazon CloudFront
- ☐ D. Amazon S3

**Your answer is incorrect.**

Answer – B

Please refer to the product description for AWS Cloutrail at the URL - <https://aws.amazon.com/cloudtrail/>

The correct answer is: Amazon Cloudtrail

[Feedback about this question and answer](#)

**QUESTION 21**

**NOT ANSWERED**

[Submit Feedback](#)

If you want to launch Amazon Elastic Compute Cloud (EC2) instances and assign each instance a private IP address in a predetermined range, you should: (choose one of the correct answer below)

Please select :

- ☐ A. Launch the instance from a private Amazon Machine Image (AMI).
- ☐ B. Assign a group of sequential Elastic IP address to the instances.
- ☐ C. Launch the instances in the Amazon Virtual Private Cloud (VPC).



aready.

**Your answer is incorrect.**

Answer – C.

This is the default reason for a VPC to host your own subnet and have EC2 instances have a private IP when it is launched in a VPC. Below is an example of an EC2 instance having a Private IP.

Instance ID	i-794961f7
Instance state	running
Instance type	t2.micro
Private DNS	ip-10-0-0-68.ap-southeast-1.compute.internal
Private IPs	10.0.0.68
Secondary private IPs	
VPC ID	vpc-39ddd55c
Subnet ID	subnet-858373f3
Network interfaces	eth0
Source/dest. check	False

For more information on private IP addresses, please refer the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-instance-addressing.html>

The correct answer is: Launch the instances in the Amazon Virtual Private Cloud (VPC).

[Feedback about this question and answer](#)

QUESTION 22

NOT ANSWERED

[Submit Feedback](#)

A company needs to deploy virtual desktops to its customers in a virtual private cloud, leveraging existing security controls. Which set of AWS services and features will meet the company's requirements?

Please select :

- ☐ A. Virtual Private Network connection. AWS Directory Services, and ClassicLink
- ☐ B. Virtual Private Network connection. AWS Directory Services, and Amazon Workspaces
- ☐ C. AWS Directory Service, Amazon Workspaces, and AWS Identity and Access Management
- ☐ D. Amazon Elastic Compute Cloud, and AWS Identity and Access Management

**Your answer is incorrect.**

Answer – B

Option A is incorrect because a ClassicLink, within the same region, allows us to link an EC2-Classical instance to a VPC in our account.

Option C is incorrect because AWS Directory service needs a VPN connection to interact with an On-premise AD directory.

Option D is incorrect because we need WorkSpaces for virtual desktops.

The correct answer is: Virtual Private Network connection. AWS Directory Services, and Amazon Workspaces

[Feedback about this question and answer](#)

QUESTION 23

NOT ANSWERED

[Submit Feedback](#)

Which of the following statements are true about Amazon Reduced Redundancy Storage (RRS)? Choose the correct 3 answers from the below options.

Please select :

- ☐ A. RRS has the ability to provide eleven nines availability.
- ☐ B. RRS has the ability to provide 99.99% availability.
- ☐ C. RRS has the ability to provide 99.99% durability.
- ☐ D. If there is a requirement to store data that is easily reproducible or durably stored elsewhere, then RRS is the ideal option.

**Your answer is incorrect.**

Answer: B,C and D. The Durability and availability are given in the aws site for RRS.

- Reduced Redundancy Storage (RRS) is an Amazon S3 storage option that enables customers to store noncritical, reproducible data at lower levels of redundancy than Amazon S3's standard storage.
- S3 is the most reliable and durable storage service from Amazon. Where as if you have data that is non-critical and can be easily reproducible if lost, then that can be stored in RRS to reduce the cost of your storage.
- The RRS option stores objects on multiple devices across multiple facilities, providing 400 times the durability of a typical disk drive, but does not replicate objects as many times as standard Amazon S3 storage.

You can read more about RRS in the below link:

- <https://aws.amazon.com/s3/reduced-redundancy/>

Here is the snapshot of durability and availability of RRS compared to other standard storage options:

The correct answers are: RRS has the ability to provide 99.99% availability., RRS has the ability to provide 99.99% durability., If there is a requirement to store data that is easily reproducible or durably stored elsewhere, then RRS is the ideal option.

[Feedback about this question and answer](#)

**QUESTION 24**

**NOT ANSWERED**

**[Submit Feedback](#)**

After creating a new IAM user which of the following must be done before they can successfully make API calls?

- ☐ B. Enable Multi-Factor Authentication for the user.
- ☐ C. Assign a Password Policy to the user.
- ☐ D. Create a set of Access Keys for the user.

**Your answer is incorrect.**

Answer – D

In IAM , when you create a user , you need to download the Access Key ID and Secret access key so that the user can access aws.

The correct answer is: Create a set of Access Keys for the user.

[Feedback about this question and answer](#)

**QUESTION 25**

**NOT ANSWERED**

[Submit Feedback](#)

What is the AWS service provided which provides a fully managed NoSQL database service that provides fast and predictable performance with seamless scalability.

Please select :

- ☐ A. AWS RDS
- ☐ B. DynamoDB
- ☐ C. Oracle RDS
- ☐ D. Elastic Map Reduce

**Your answer is incorrect.**

Answer: - B

The correct answer is: DynamoDB

[Feedback about this question and answer](#)

QUESTION 26

NOT ANSWERED

[Submit Feedback](#)

A company's application is intending to use Auto Scaling and has the requirement to store user state information.

Which of the following AWS services provides a shared data store with durability and low latency?

Please select :

- ☐ A. AWS ElastiCache Memcached
- ☐ B. Amazon Simple Storage Service
- ☐ C. Amazon EC2 instance storage
- ☐ D. Amazon DynamoDB

**Your answer is incorrect.**

Answer – D

Amazon Dynamo DB is used for storing small amounts of data such as user state information. And this service offer's durability and low latency. Below is a snapshot of when to use S3 and DynamoDB from the DynamoDB FAQ's - <https://aws.amazon.com/dynamodb/faqs/>

The correct answer is: Amazon DynamoDB

[Feedback about this question and answer](#)

QUESTION 27

NOT ANSWERED

[Submit Feedback](#)

Which of the following requires a custom CloudWatch metric to monitor?

- ☐ B. CPU Utilization of an EC2 instance
- ☐ C. Disk usage activity of an EC2 instance
- ☐ D. Data transfer of an EC2 instance

**Your answer is incorrect.**

Answer – A

Memory Utilization is a metric not offered directly by Cloudwatch.

So when you view the Cloudwatch metrics for your EC2 instance, you can see CPU Utilization and Disk Read Operations metrics.


**CloudWatch metrics:** Basic monitoring. [Enable Detailed Monitoring](#)

Showing data for: **Last Hour**

Below are your CloudWatch metrics for the selected resources (a maximum of 10). Click on a graph to see an expanded view. All times shown are in UTC.

[View all CloudWatch metrics](#)

**CPU Utilization (Percent)**






**Disk Reads (Bytes)**



**Disk Read Operations (Operations)**



You can also see Network statistics for Data transfer, but you will not be able to see Memory Utilization. This will be a custom Cloudwatch metric.



The correct answer is: Memory Utilization of an EC2 instance

[Feedback about this question and answer](#)

QUESTION 28	NOT ANSWERED	<a href="#">Submit Feedback</a>
-------------	--------------	---------------------------------

Which of the following notification endpoints or clients are supported by Amazon Simple Notification Service? Choose 2 answers from the options below.

Please select :

- ☐ A. Email
- ☐ B. CloudFront distribution
- ☐ C. File Transfer Protocol
- ☐ D. Short Message Service
- ☐ E. Simple Network Management Protocol

**Your answer is incorrect.**

Answer – A and D

When you create a subscription in SNS , below are the protocols available.

The correct answers are: Email, Short Message Service

[Feedback about this question and answer](#)

QUESTION 29

NOT ANSWERED

[Submit Feedback](#)

You have a read intensive application hosted in AWS. The application is currently using the MySQL RDS feature in AWS. Which of the following can be used to reduce the read throughput on the MySQL database

Please select :

- ☐ A. Enable the Multi-AZ on the MySQL RDS
- ☐ B. Use Cold Storage Volumes for the MySQL RDS
- ☐ C. Enable Read Replica's and offload the reads to the replica's
- ☐ D. Use SQS to queue up the reads

**Your answer is incorrect.**

Answer – C

The AWS documentation mentions the following on Read Replica's



DB Instance and serve high-volume application read traffic from multiple copies of your data, thereby increasing aggregate read throughput. Read replicas can also be promoted when needed to become standalone DB instances.

For more information on Read Replica's , please visit the below URL:

- <https://aws.amazon.com/rds/details/read-replicas/>

The correct answer is: Enable Read Replica's and offload the reads to the replica's

[Feedback about this question and answer](#)

QUESTION 30

NOT ANSWERED

[Submit Feedback](#)

The Trusted Advisor service provides insight regarding which four categories of an AWS account?

Please select :

- ☐ A. Security, fault tolerance, high availability, and connectivity
- ☐ B. Security, access control, high availability, and performance
- ☐ C. Performance, cost optimization, security, and fault tolerance
- ☐ D. Performance, cost optimization, access control, and connectivity

**Your answer is incorrect.**

Answer – C

Below is the screenshot of what services the Trusted Advisor Dashboard offers.

**QUESTION 31****NOT ANSWERED**[Submit Feedback](#)

A company is deploying a two-tier, highly available web application to AWS. Which service provides durable storage for static content while utilizing lower Overall CPU resources for the web tier?

Please select :

- ☐ A. Amazon EBS volume
- ☐ B. Amazon S3
- ☐ C. Amazon EC2 instance store
- ☐ D. Amazon RDS instance

**Your answer is incorrect.**

Answer – B

Amazon S3 is the default storage service that should be considered for companies. It provides durable storage for all static content.

The correct answer is: Amazon S3

[Feedback about this question and answer](#)

**QUESTION 32****NOT ANSWERED**[Submit Feedback](#)

When will you incur costs with an Elastic IP address (EIP)?

Please select :

- ☐ A. When an EIP is allocated.
- ☐ B. When it is allocated and associated with a running instance.
- ☐ C. When it is allocated and associated with a stopped instance.
- ☐ D. Costs are incurred regardless of whether the EIP is associated with a running instance.

**Your answer is incorrect.**

Answer- C

AWS doesn't want you waste the static public IP's . You will be charged for elastic IP

1 - If EIP is created but not allocated to any instance.

2 - If EIP is attached to a stop instance.

<https://aws.amazon.com/premiumsupport/knowledge-center/elastic-ip-charges/>

The correct answer is: When it is allocated and associated with a stopped instance.

[Feedback about this question and answer](#)

QUESTION 33

NOT ANSWERED

[Submit Feedback](#)

Which of the following instance types are available as Amazon EBS-backed only?  
Choose 2 answers from the options below.

Please select :

- ☐ A. General purpose T2
- ☐ B. General purpose M3
- ☐ C. Compute-optimized C4
- ☐ D. Compute-optimized C3
- ☐ E. Storage-optimized I2

**Your answer is incorrect.**

Answer – A and C

For details for all instance types, please visit the url - <https://aws.amazon.com/ec2/instance-types/>

The below screenshots show the details for T2 and C4 instances which have a storage type of EBS.

t2.nano	1	3	0.5	EBS-Only
t2.micro	1	6	1	EBS-Only
t2.small	1	12	2	EBS-Only
t2.medium	2	24	4	EBS-Only
t2.large	2	36	8	EBS-Only

Model	vCPU	Mem (GiB)	Storage	Dedicated EBS Bandwidth (Mbps)
c4.large	2	3.75	EBS-Only	500
c4.xlarge	4	7.5	EBS-Only	750
c4.2xlarge	8	15	EBS-Only	1,000
c4.4xlarge	16	30	EBS-Only	2,000
c4.8xlarge	36	60	EBS-Only	4,000

The correct answers are: General purpose T2, Compute-optimized C4

[Feedback about this question and answer](#)

QUESTION 34

NOT ANSWERED

[Submit Feedback](#)

How many availability zones are mapped to a subnet?

- ☐ B. 2
- ☐ C. Depends on aws at the time of creating a subnet
- ☐ D. Depends on the number of instances you are going to host in the subnet.

**Your answer is incorrect.**

Answer: A – Remember that when a subnet is created, it is always mapped to one availability zone.

When you go the VPC dashboard, and go to the Subnet section, you can click on Create Subnet

When you create the subnet, you can only attach one AZ to the subnet.

The correct answer is: 1

[Feedback about this question and answer](#)

QUESTION 35

NOT ANSWERED

[Submit Feedback](#)

A company is building a service in which Amazon EC2 worker instances process an uploaded audio file and generate a text file. You must store both of these files in the same durable storage until the text file is retrieved. You do not know what the storage capacity requirements are. Which storage option is both cost-efficient and scalable?

Please select :

- ☐ C. A single Amazon S3 bucket
- ☐ D. Multiple instance stores

**Your answer is incorrect.**

Answer – C

For any sort of storage for file based system, it must be done in Amazon S3.

The correct answer is: A single Amazon S3 bucket

[Feedback about this question and answer](#)

**QUESTION 36**

**NOT ANSWERED**

[Submit Feedback](#)

A custom script needs to be passed to a new Amazon Linux instances created in your Auto Scaling group. Which feature allows you to accomplish this?

Please select :

- ☐ A. User data
- ☐ B. EC2Config service
- ☐ C. IAM roles
- ☐ D. AWS Config

**Your answer is incorrect.**

Answer – A

When you configure an instance during creation, you can add custom scripts to the User data section.

So in Step 3 of creating an instance, in the Advanced Details section, we can enter custom scripts in the User Data section. The below script installs Perl during the instance creation of the EC2 instance.

Tenancy ⓘ

Shared - Run a shared hardware instance ▼

[Additional charges will apply for dedicated tenancy.](#)

### ▼ Network interfaces ⓘ

Device	Network Interface	Subnet	Primary IP	Secondary IP addresses
eth0	New network interface ▼	subnet-95ed8dd1 ▼	Auto-assign	<a href="#">Add IP</a>

Add Device

### ▼ Advanced Details

User data ⓘ

☒ As text ☐ As file ☐ Input is already base64 encoded

```
yum install perl-Switch perl-DateTime perl-Sys-Syslog perl-LWP-Protocol-https  
-y
```

The correct answer is: User data

[Feedback about this question and answer](#)

QUESTION 37

NOT ANSWERED

[Submit Feedback](#)

A company is building software on AWS that require access to various AWS services. Which configuration should be used to ensure that AWS Credentials like Access Keys and Secret access keys are not compromised? (Choose Two Options)

Please select :

- ☐ A. Enable Multi-Factor Authentication for your AWS root account.
- ☐ B. Assign an IAM role to the Amazon EC2 instance.
- ☐ C. Store the AWS Access Key ID/Seret Access Key combination in software comments.
- ☐ D. Assign an IAM user to the Amazon EC2 Instance.

**Your answer is incorrect.**

Answer – A and B

The correct answer is: Enable Multi-Factor Authentication for your AWS root account and Assign an IAM role to the Amazon EC2 instance.

The correct answers are: Enable Multi-Factor Authentication for your AWS root account., Assign an IAM role to the Amazon EC2 instance.

[Feedback about this question and answer](#)

QUESTION 38

NOT ANSWERED

[Submit Feedback](#)

A company has the requirement to store data using AWS storage services. The data is not frequently accessed. If data recovery time not an issue, Which of the below is the best and cost efficient solution to fulfil this requirement ?

Please select :

- ☐ A. S3 Standard
- ☐ B. S3 Standard - IA (Infrequently Accessed)
- ☐ C. Glacier
- ☐ D. Reduced Redundancy Storage

**Your answer is incorrect.**

Answer – C

The default time interval is one minute.

Note: Answer can also be B. S3 Standard - IA (Infrequently Accessed). However since other details are mentioned in question. we can say C. Glacier is most effective way of cost saving in this case.

Reference link: <https://aws.amazon.com/products/storage/>

The correct answer is: Glacier

[Feedback about this question and answer](#)

QUESTION 39

NOT ANSWERED

[Submit Feedback](#)

Resources that are created in AWS are identified by a unique identifier which is known as what option given below



- ☐ B. Amazon Resource Nametag
- ☐ C. Amazon Resource Name
- ☐ D. Amazon Resource Namespace

**Your answer is incorrect.**

Answer – C

Amazon Resource Names (ARNs) are used to uniquely identify AWS resources. For information on ARN's, refer to the link -

<http://docs.aws.amazon.com/general/latest/gr/aws-arns-and-namespaces.html>

The correct answer is: Amazon Resource Name

[Feedback about this question and answer](#)

**QUESTION 40**

**NOT ANSWERED**

[Submit Feedback](#)

When you are using Route53 for a web site hosted in S3 , what are the one of the rules that must be adhered to? Choose the correct answer from the options below

Please select :

- ☐ A. The S3 bucket name must be the same as the domain name
- ☐ B. The record set cannot use an alias
- ☐ C. The record set must be of type "MX"
- ☐ D. The S3 bucket must be in the same region as the hosted zone

**Your answer is incorrect.**

Answer – A

This is given in the aws documentation

### Prerequisites

Before you get started, you need the following:

- An S3 bucket that is configured to host a static website. For more information, see [Configure a Bucket for Website Hosting](#) in the *Amazon Simple Storage Service Developer Guide*.

#### Important

The bucket must have the same name as your domain or subdomain. For example, if you want to use the subdomain `acme.example.com`, the name of the bucket must be `acme.example.com`.

The correct answer is: The S3 bucket name must be the same as the domain name

[Feedback about this question and answer](#)

QUESTION 41

NOT ANSWERED

[Submit Feedback](#)

What are some of the benefits of using the Cloudformation service? Choose 2 answers from the options given below

Please select :

- ☐ A. Can automatically increase instance capacity
- ☐ B. A storage location for your applications code
- ☐ C. Version control your infrastructure
- ☐ D. A great disaster recovery option

**Your answer is incorrect.**

Answer – C and D

The justification for Infrastructure as code is given in the aws documentation

## Infrastructure as Code

A template can be used repeatedly to create identical copies of the same stack (or to use as a foundation to start a new stack). You can capture and control region-specific infrastructure variations such as Amazon EC2 AMIs, as well as Amazon EBS and Amazon RDS snapshot names. Templates are simple JSON or YAML formatted text files that can be placed under your normal source control mechanisms, stored in private or public locations such as Amazon S3, and exchanged via email. With AWS CloudFormation, you can "open the hood," to see exactly which AWS resources make up a stack. You retain full control and have the ability to modify any of the AWS resources created as part of a stack.

For the justification on disaster recovery, please visit the below link

<https://aws.amazon.com/blogs/aws/new-whitepaper-use-aws-for-disaster-recovery/>

For more information on Cloudformation, please visit the link

<https://aws.amazon.com/cloudformation/>

The correct answers are: Version control your infrastructure, A great disaster recovery option

[Feedback about this question and answer](#)

AWS thrives on the concept of high availability. Which of the below follows the concept of high availability. Choose the correct answer from the options below

Please select :

- ☐ A. Implementing security procedures
- ☐ B. Implementing multiple AWS services
- ☐ C. The ability of system to easily increase in size.
- ☐ D. A durable system that can operate for long periods of time without failure.

**Your answer is incorrect.**

Answer – D

High availability is a characteristic of a system, which aims to ensure an agreed level of operational performance, usually uptime, for a higher than normal period.

For more information on high availability, please refer to the following link

[https://en.wikipedia.org/wiki/High\\_availability](https://en.wikipedia.org/wiki/High_availability)

The correct answer is: A durable system that can operate for long periods of time without failure.

[Feedback about this question and answer](#)

QUESTION 43

NOT ANSWERED

[Submit Feedback](#)

What best describes the "Principal of Least Privilege"? Choose the correct answer from the options given below

Please select :

- ☐ A. All users should have the same baseline permissions granted to them to use basic AWS services.
- ☐ B. Users should be granted permission to access only resources they need to do their assigned job.
- ☐ C. Users should submit all access request in written so that there is a paper trail of who needs access to different AWS resources.
- ☐ D. Users should always have a little more access granted to them than they need, just in case they end up needing it in the future.

**Your answer is incorrect.**

Answer – B

For more information on principle of least privilege, please refer to the following link

[https://en.wikipedia.org/wiki/Principle\\_of\\_least\\_privilege](https://en.wikipedia.org/wiki/Principle_of_least_privilege)

The correct answer is: Users should be granted permission to access only resources they need to do their assigned job.

[Feedback about this question and answer](#)

QUESTION 44

NOT ANSWERED

[Submit Feedback](#)

Which of the following best describes the purpose of an Elastic Load Balancer. Choose an answer from the options given below. Choose the correct answer from the options given below

Please select :

- ☐ A. To scale more EC2 instances on demand
- ☐ B. To evenly distribute traffic among multiple EC2 instances located in single or different Availability Zones.
- ☐ C. To distribute traffic to a second instance once the first instance capacity has reached it's limit.
- ☐ D. To evenly distribute traffic among multiple EC2 instances in the same Availability Zone.

**Your answer is incorrect.**

Answer – B

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve fault tolerance in your applications, seamlessly providing the required amount of load balancing capacity needed to route application traffic.

**And the ELB is used to distribute traffic between instances in Multiple AZ's.**

For more information on Elastic Load Balancer, please refer to the following link

- <https://aws.amazon.com/elasticloadbalancing/>

Some more key points about ELB:

- Elastic Load Balancer (ELB) is used for rotating traffic to various EC2 instances located across the multiple Availability Zones(AZs).
- ELB can detect the healthy and unhealthy EC2 instances. It will not route traffic to the unhealthy EC2 instances.
- If all the instances in the same AZ is not healthy, it will route the traffic to other AZ EC2 instances.
- Achieve higher levels of fault tolerance for your applications by using Elastic Load Balancing to automatically route traffic across multiple instances and multiple Availability Zones. Elastic Load Balancing ensures that only healthy Amazon EC2 instances receive traffic by detecting unhealthy

The correct answer is: To evenly distribute traffic among multiple EC2 instances located in single or different Availability Zones.

[Feedback about this question and answer](#)

**QUESTION 45**

**NOT ANSWERED**

[Submit Feedback](#)

When you create a default VPC, what are the services you get by default in the VPC?  
Select 2 options.

Please select :

- ☐ A. An Elastic Load Balancer
- ☐ B. Default subnet in each Availability Zone
- ☐ C. An Internet Gateway attached to the default VPC

Your answer is incorrect.

Answer: B and C.

For the list of default services given for a default VPC, follow the link

<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/default-vpc.html>

to get more information on what comes as part of a default VPC.

The correct answers are: Default subnet in each Availability Zone, An Internet Gateway attached to the default VPC

[Feedback about this question and answer](#)

QUESTION 46

NOT ANSWERED

[Submit Feedback](#)

In an Autoscaling policy what best describes what the purpose of a scaling policy is. Choose an answer from the options below.

Please select :

- ☐ A. A set of CloudWatch metric thresholds that dictate when to add or remove instances from the Auto Scaling group.
- ☐ B. The IAM access policy granted to an Auto Scaling group.
- ☐ C. The percentage at which an ELB will send traffic to an instance before it sends traffic to a different instance.
- ☐ D. An SNS notification alert.

Your answer is incorrect.

Answer – A

You can create a scaling policy that uses CloudWatch alarms to determine when your Auto Scaling group should scale out or scale in. Each CloudWatch alarm watches a single metric and sends messages to Auto Scaling when the metric breaches a threshold that you specify in your policy. You can use alarms to monitor any of the metrics that the services in AWS that you're using send to CloudWatch, or you can create and monitor your own custom metrics.

For more information on Scaling policies, please refer to the following link

[http://docs.aws.amazon.com/autoscaling/latest/userguide/policy\\_creating.html](http://docs.aws.amazon.com/autoscaling/latest/userguide/policy_creating.html)

The correct answer is: A set of CloudWatch metric thresholds that dictate when to add or remove instances from the Auto Scaling group.

[Feedback about this question and answer](#)

A company has a solution hosted in AWS. This solution consists of a set of EC2 instances. They have been recently getting attacks as their IT security departments identified that attacks are from a set of IP addresses. Which of the following methods can be adopted to help in this situation.

Please select :

- ☐ A. Place the EC2 instances into private subnets, and set up an NAT gateway so employees can access them.
- ☐ B. Remove the IGW from the VPC so that no outside traffic can reach the EC2 instances.
- ☐ C. Lock down of NACL for the set to IP address.
- ☐ D. Place the EC2 instances into private subnets, and set up a bastion host so employees can access them.

**Your answer is incorrect.**

Answer – C

The NACL's can be modified to be most secure by only denying the traffic from the set of IP addresses.

For more information on NACL, please refer to the following link

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_ACLS.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLS.html)

The correct answer is: Lock down of NACL for the set to IP address.

[Feedback about this question and answer](#)

QUESTION 48

NOT ANSWERED

[Submit Feedback](#)

You have an ELB distributing traffic a fleet of EC2 instances inside your VPC, evenly spread across two Availability Zones. However, you realize that only half of our instances are actually receiving traffic. What is the most likely cause of this problem? Choose the correct answer from the options given below

Please select :

- ☐ A. The ELBs listener is not set to port 80.
- ☐ B. One or more security groups do not allow HTTP traffic.
- ☐ C. Cross-zone load balancing has not been enabled.
- ☐ D. The health check ping port is set to port 80, but should be set to port 22.

For environments where clients cache DNS lookups, incoming requests might favor one of the Availability Zones. Using cross-zone load balancing, this imbalance in the request load is spread across all available instances in the region, reducing the impact of misbehaving clients.

By default, your Classic Load Balancer distributes incoming requests evenly across its enabled Availability Zones. For example, if you have ten instances in Availability Zone us-west-2a and two instances in us-west-2b, the requests are distributed evenly between the two Availability Zones. As a result, the two instances in us-west-2b serve the same amount of traffic as the ten instances in us-west-2a. To ensure that your load balancer distributes incoming requests evenly across all instances in its enabled Availability Zones, enable cross-zone load balancing.

For more information on ELB Cross load balancer, please refer to the following link

<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/enable-disable-crosszone-lb.html>

The correct answer is: Cross-zone load balancing has not been enabled.

[Feedback about this question and answer](#)

QUESTION 49

NOT ANSWERED

[Submit Feedback](#)

You have an application currently running on five EC2 instances as part of an Auto Scaling group. For the past 30 minutes all five instances have been running at 100 CPU Utilization; however, the Auto Scaling group has not added any more instances to the group. What is the most likely cause? Choose 2 likely answers from the options given below

Please select :

- ☐ A. You already have 20 on-demand instances running.
- ☐ B. The Auto Scaling group's MAX size is set at five.
- ☐ C. The Auto Scaling group's scale down policy is too high.
- ☐ D. The Auto Scaling group's scale up policy has not yet been reached.

**Your answer is incorrect.**

Answer – A and B

Twenty instances limit is at the account level and you might have other applications running more EC2 instances across your account (may be in another region) which may cause a total number to exceed the limit.

This is provided in the aws documentation



- **Cause:** The Auto Scaling group has reached the limit set by the `DesiredCapacity` parameter.
- **Solution:**
  - Update your Auto Scaling group by providing a new value for the `--desired-capacity` parameter using the `update-auto-scaling-group` command.
  - If you've reached your limit for number of EC2 instances, you can request an increase. For more information, see [AWS Service Limits](#).

For more information on troubleshooting Autoscaling, please refer to the following link

<http://docs.aws.amazon.com/autoscaling/latest/userguide/ts-as-capacity.html>

The correct answers are: You already have 20 on-demand instances running., The Auto Scaling group's MAX size is set at five.

[Feedback about this question and answer](#)

QUESTION 50	NOT ANSWERED	<a href="#">Submit Feedback</a>
-------------	--------------	---------------------------------

If need to upload a file to S3 that is 600MB in size, which of the following is the best option to use? Choose an answer from the options below. Choose the correct answer from the options below

Please select :

- ☐ A. Single operation upload
- ☐ B. Snowball
- ☐ C. AWS Import/Export
- ☐ D. Multi-part upload

**Your answer is incorrect.**

Answer – D

The Multipart upload API enables you to upload large objects in parts. You can use this API to upload new large objects or make a copy of an existing object (see Operations on Objects).

Multipart uploading is a three-step process: You initiate the upload, you upload the object parts, and after you have uploaded all the parts, you complete the multipart upload. Upon receiving the complete multipart upload request, Amazon S3 constructs the object from the uploaded parts, and you can then access the object just as you would any other object in your bucket.

For more information on Multi-part file upload, please refer to the following link

<http://docs.aws.amazon.com/AmazonS3/latest/dev/mpuoverview.html>

**QUESTION 51****NOT ANSWERED**[Submit Feedback](#)

A company has the requirement to store files in S3. After a period of a month, these files can be archived. The archived files might be required after a period of 3-4 months. Which of the following suits the requirements

Please select :

- ☐ A. Use EC2 instances with EBS volumes, one for normal storage and the other for archived storage
- ☐ B. Use S3 for normal file storage and use lifecycle policies for moving the files to glacier.
- ☐ C. Use EC2 instances with EBS volumes and use lifecycle policies for moving the files to glacier.
- ☐ D. Use glacier for normal file storage and use lifecycle policies for moving the files to S3.

**Your answer is incorrect.**

Answer – B

Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects. These actions can be classified as follows:

- Transition actions – In which you define when objects transition to another storage class. For example, you may choose to transition objects to the STANDARD\_IA (IA, for infrequent access) storage class 30 days after creation, or archive objects to the GLACIER storage class one year after creation.

For more information on Lifecycle policies, please refer to the following link

<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>

The correct answer is: Use S3 for normal file storage and use lifecycle policies for moving the files to glacier.

[Feedback about this question and answer](#)

**QUESTION 52****NOT ANSWERED**[Submit Feedback](#)

Your team has an application hosted on Docker containers. You want to port that application in the easiest way possible onto AWS for your development community. Which of the following service can be used to fulfil this requirement

- ☒ B. AWS SNS
- ☐ C. AWS SQS
- ☐ D. AWS Elastic Beanstalk

**Your answer is incorrect.**

Answer – D

The AWS documentation mentions the following

Elastic Beanstalk supports the deployment of web applications from Docker containers. With Docker containers, you can define your own runtime environment. You can choose your own platform, programming language, and any application dependencies (such as package managers or tools), that aren't supported by other platforms. Docker containers are self-contained and include all the configuration information and software your web application requires to run.

For more information on Elastic beanstalk and docker , please visit the below URL:

- [http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/create\\_deploy\\_docker.html](http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/create_deploy_docker.html)

The correct answer is: AWS Elastic Beanstalk

[Feedback about this question and answer](#)

QUESTION 53

NOT ANSWERED

[Submit Feedback](#)

Which of the following tools for EC2 can be used to administer instances without the need to SSH or RDP into the instance.

Please select :

- ☐ A. AWS Config
- ☐ B. AWS CodePipeline
- ☐ C. Run Command
- ☐ D. EC2Config

**Your answer is incorrect.**

Answer - C

You can use Run Command from the Amazon EC2 console to configure instances without having to login to each instance

For more information on the Run Command , please visit the below URL:

- <http://docs.aws.amazon.com/systems-manager/latest/userguide/rc-console.html>

QUESTION 54

NOT ANSWERED

[Submit Feedback](#)

If you wanted to extend your on-premise infrastructure with AWS, which of the below options would help. Choose 2 answers from the options given below

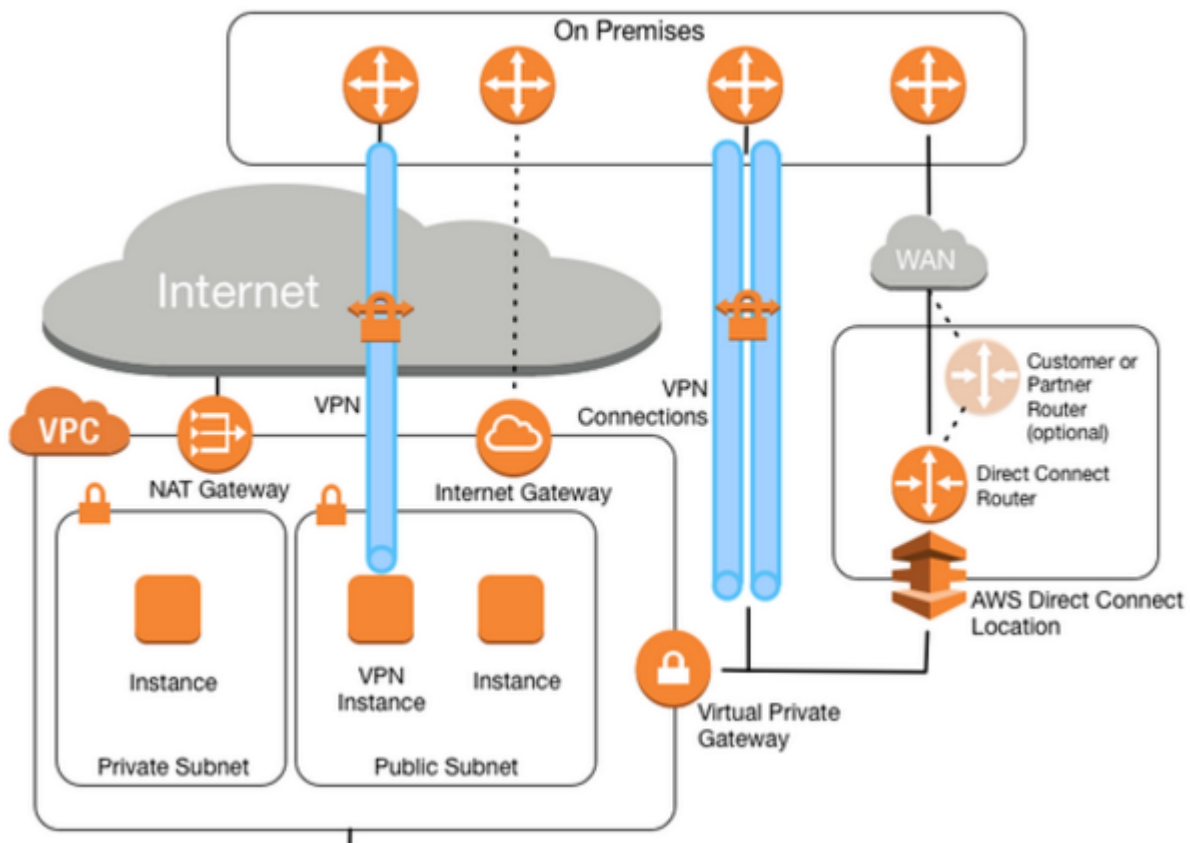
Please select :

- ☐ A. Virtual Private Network
- ☐ B. Customer gateway
- ☐ C. Direct Connect
- ☐ D. Primary Connection

**Your answer is incorrect.**

Answer – A and C

You can either build a VPN or have a direct connect connection



For more information on VPC to on-premise networks, please refer to the following link

[Feedback about this question and answer](#)

QUESTION 55

NOT ANSWERED

[Submit Feedback](#)

Why does stopping and starting an instance help in fixing a System Status Check error? Choose an answer from the options given below

Please select :

- ☐ A. Stopping and starting an instance causes the instance to change the AMI.
- ☐ B. Stopping and starting an instance causes the instance to be provisioned on different AWS hardware.
- ☐ C. Stopping and starting an instance reboots the operating system.
- ☐ D. None of the above

**Your answer is incorrect.**

Answer – B

Refer below steps published by AWS support: <https://aws.amazon.com/premiumsupport/knowledge-center/system-reachability-check/>

This is given in the aws documentation

When you stop a running instance, the following happens:

- The instance performs a normal shutdown and stops running; its status changes to *stopping* and then *stopped*.
- Any Amazon EBS volumes remain attached to the instance, and their data persists.
- Any data stored in the RAM of the host computer or the instance store volumes of the host computer is gone.
- In most cases, the instance is migrated to a new underlying host computer when it's started.

For more information on starting and stopping instances, please refer to the following link

[http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Stop\\_Start.html](http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Stop_Start.html)

The correct answer is: Stopping and starting an instance causes the instance to be provisioned on different AWS hardware.

[Feedback about this question and answer](#)

QUESTION 56

NOT ANSWERED

[Submit Feedback](#)

In consolidated billing what are the 2 different types of accounts.

- ☐ B. Parent account and Child account
- ☐ C. Main account and Sub account.
- ☐ D. Primary account and Secondary account.

**Your answer is incorrect.**

Answer: A – You can have a combination of Paying accounts and linked accounts. When you have consolidated billing you have the facility to reduce the costs for the paying account. This is one of the main advantages of consolidated billing.

The correct answer is: Paying account and Linked account

[Feedback about this question and answer](#)

**QUESTION 57**

**NOT ANSWERED**

[Submit Feedback](#)

What is the term often used to describe a DR scenario in which a minimal version of an environment is always running in the cloud. Choose an answer from the options given below

Please select :

- ☐ A. Backup and Restore
- ☐ B. Pilot Light
- ☐ C. Warm standby
- ☐ D. Multi Site

**Your answer is incorrect.**

Answer - B

This is given in a whitepaper published by AWS

For more information on disaster recovery, please refer to the below link

[https://media.amazonwebservices.com/AWS\\_Disaster\\_Recovery.pdf](https://media.amazonwebservices.com/AWS_Disaster_Recovery.pdf)

The correct answer is: Pilot Light

**QUESTION 58****NOT ANSWERED**[Submit Feedback](#)

Which of the following features ensures even distribution of traffic to Amazon EC2 instances in multiple Availability Zones registered with a load balancer?

Please select :

- ☐ A. Elastic Load Balancing request routing
- ☐ B. An Amazon Route 53 weighted routing policy
- ☐ C. Elastic Load Balancing cross-zone load balancing
- ☐ D. An Amazon Route 53 latency routing policy

**Your answer is incorrect.**

Answer: C – To ensure that traffic is evenly distributed, you need to ensure the “Enable Cross-Zone Load balancing option” is chosen. This option comes up when you are creating a classic load balancer in Step 5 of Add EC2 instances.

The correct answer is: Elastic Load Balancing cross-zone load balancing

[Feedback about this question and answer](#)

Currently you have a VPC with EC2 Security Group and several running EC2 instances. You change the Security Group rules to allow inbound traffic on a new port and protocol, and launch several new instances in the same Security Group. When will the Security Group changes be applied to the EC2 instances? Please choose the correct answer.

Please select :

- ☐ A. Immediately to all instances in the security group.
- ☐ B. Immediately to the new instances only.
- ☐ C. Immediately to the new instances, but old instances must be stopped and restarted before the new rules apply.
- ☐ D. To all instances, but it may take several minutes for old instances to see the changes.

**Your answer is incorrect.**

Answer: A – By default whatever changes you make to security rules will be applied in all instances which are part of that security group.

When you add or remove rules, they are automatically applied to all instances associated with the security group.

For more information, please refer the below URL:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_SecurityGroups.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html)

The correct answer is: Immediately to all instances in the security group.

[Feedback about this question and answer](#)

QUESTION 60

NOT ANSWERED

[Submit Feedback](#)

What is the minimum size of an object that can be uploaded to Amazon S3?

Please select :

- ☐ A. 1Megabyte
- ☐ B. 0Bytes
- ☐ C. 1Byte
- ☐ D. 5TB

**Your answer is incorrect.**



response to the question is as follows:

**Q: How much data can I store?**

The total volume of data and number of objects you can store are unlimited. Individual Amazon S3 objects can range in size from a minimum of 0 bytes to a maximum of 5 terabytes. The largest object that can be uploaded in a single PUT is 5 gigabytes. For objects larger than 100 megabytes, customers should consider using the [Multipart Upload](#) capability.

The correct answer is: 0Bytes

[Feedback about this question and answer](#)

There is no Incorrect answer(s)

---

## Company

[About Us](#)

[Discussions](#)

[Blog](#)

## Support

contact us @ [support@whizlabs.com](mailto:support@whizlabs.com)

## Follow Us

