


## Whizlabs Learning Center

[HOME](#) > [AWS CSAA PRACTICE TESTS](#) > [FULL TEST\(S\)](#) > [NEW DIAGNOSTIC TEST](#)**Started on** Wednesday, 5 July 2017, 3:13 AM**State** Finished**Completed on** Wednesday, 5 July 2017, 3:14 AM**Time taken** 33 secs**Grade** 0.00 out of 60.00 (0%)**Result** FAIL**Question 1**

Not answered

 Mark for Review

In AWS what is used for encrypting and decrypting login information to EC2 instances.

Please select :

- ☐ A. Templates
- ☐ B. AMI's
- ☐ C. Key pairs
- ☐ D. None of the above

Your answer is incorrect.

Answer - C

Amazon EC2 uses public-key cryptography to encrypt and decrypt login information. Public-key cryptography uses a public key to encrypt a piece of data, such as a password, then the recipient uses the private key to decrypt the data. The public and private keys are known as a key pair.

To log in to your instance, you must create a key pair, specify the name of the key pair when you launch the instance, and provide the private key when you connect to the instance. Linux instances have no password, and you use a key pair to log in using SSH. With Windows instances, you use a key pair to obtain the administrator password and then log in using RDP.

For more information on key pairs, please visit the below url


- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-key-pairs.html>

The correct answer is: Key pairs

**Feedback about this question and answer**

**Question 2**

Not answered

 Mark for Review

For DynamoDB, what are the scenario's in which you would want to enable Cross-region replication?

Please select :

- ☐ A. Live data migration
- ☐ B. Easier Traffic management
- ☐ C. Disaster Recovery
- ☐ D. All of the above

Your answer is incorrect.

Answer - D

From the AWS Documentation, it clearly states the reason for why you would want to enable Cross-Region Replication

**Q: When should I use cross-region replication?**

You can use cross-region replication for the following scenarios.

- **Efficient disaster recovery:** By replicating tables in multiple data centers, you can switch over to using DynamoDB tables from another region in case a data center failure occurs.
- **Faster reads:** If you have customers in multiple regions, you can deliver data faster by reading a DynamoDB table from the closest AWS data center.
- **Easier traffic management:** You can use replicas to distribute the read workload across tables and thereby consume less read capacity in the master table.
- **Easy regional migration:** By creating a read replica in a new region and then promoting the replica to be a master, you migrate your application to that region more easily.
- **Live data migration:** To move a DynamoDB table from one region to another, you can create a replica of the table from the source region in the destination region. When the tables are in sync, you can switch your application to write to the destination region.

For more information on DynamoDB, please visit the url

- <https://aws.amazon.com/dynamodb/faqs/>

The correct answer is: All of the above

**Feedback about this question and answer**

**Question 3**

Not answered

Mark for Review

In VPCs with private and public subnets, web servers should ideally be launched into:

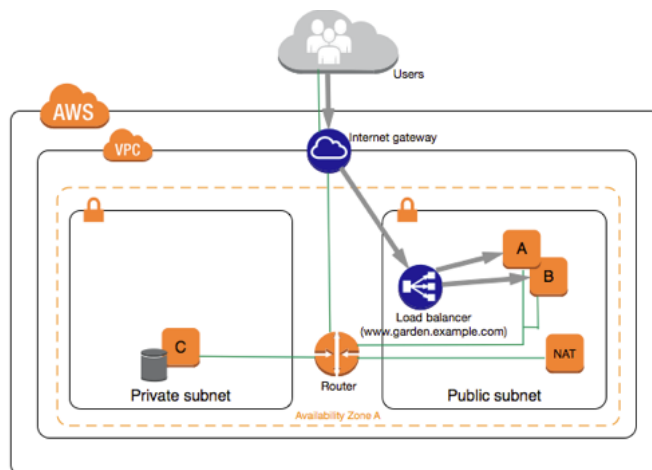
Please select :

- ☐ A. The public subnet
- ☐ B. The private subnet
- ☐ C. Either of them
- ☐ D. Not recommended, they should ideally be launched outside VPC

Your answer is incorrect.

Answer – A

Normally web servers are launched in the public subnet which has access to the internet.



For more information on private and public subnets, please visit the url

[http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Scenario2.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario2.html)

The correct answer is: The public subnet

**Feedback about this question and answer**

**Question 4**

Not answered

Mark for Review

True or False: A webserver running on an EC2 instance can be used as an origin server in CloudFront.

Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer – A

Currently Cloudfront supports the following types of distributions

- S3 buckets - When you use Amazon S3 as an origin for your distribution, you place any objects that you want CloudFront to deliver in an Amazon S3 bucket.
- Custom Origin - A custom origin is an HTTP server, for example, a web server. The HTTP server can be an Amazon EC2 instance or an HTTP server that you manage privately. When you use a custom origin, you specify the DNS name of the server, along with the HTTP and HTTPS ports and the protocol that you want CloudFront to use when fetching objects from your origin.

For more information on Cloudfront Distributions, please visit the url


<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html>

The correct answer is: True

**Feedback about this question and answer**

#### Question 5

Not answered

 Mark for Review

Which of the following is mandatory when defining a cloudformation template?

Please select :

- ☐ A. Resources
- ☐ B. Parameters
- ☐ C. Outputs
- ☐ D. Mappings

Your answer is incorrect.

Answer – A

This is clearly given in the aws documentation

#### **Resources (required)**

Specifies the stack resources and their properties, such as an Amazon Elastic Compute Cloud instance or an Amazon Simple Storage Service bucket. You can refer to resources in the **Resources** and **Outputs** sections of the template.

For more information on Cloudformation templates, please visit the url


<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/template-anatomy.html>

The correct answer is: Resources

**Feedback about this question and answer**

#### Question 6

Not answered

 Mark for Review

In IAM, what is the representation of a person or service ?

Please select :

- ☐ A. User
- ☐ B. Group
- ☐ C. Team
- ☐ D. Role

Your answer is incorrect.

Answer – A

An IAM *user* is an entity that you create in AWS. The IAM user represents the person or service who uses the IAM user to interact with AWS

An IAM *group* is a collection of IAM users. You can use groups to specify permissions for a collection of users, which can make those permissions easier to manage for those users

An IAM *role* is very similar to a user, in that it is an identity with permission policies that determine what the identity can and cannot do in AWS

For more information on IAM entities , please visit the url


<http://docs.aws.amazon.com/IAM/latest/UserGuide/id.html>

The correct answer is: User

**Feedback about this question and answer**

**Question 7**

Not answered

 Mark for Review

True or False: One of the advantages of a Multi-AZ RDS deployment is automatic failover.

Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer - A

Some of the advantages of Multi AZ rds deployments are given below

- If an Availability Zone failure or DB Instance failure occurs, your availability impact is limited to the time automatic failover takes to complete
- The availability benefits of Multi-AZ deployments also extend to planned maintenance and backups. In the case of system upgrades like OS patching or DB Instance scaling, these operations are applied first on the standby, prior to the automatic failover. As a result, your availability impact is, again, only the time required for automatic failover to complete.
- If a storage volume on your primary fails in a Multi-AZ deployment, Amazon RDS automatically initiates a failover to the up-to-date standby

For more information on Multi AZ rds deployments please visit the link


<https://aws.amazon.com/rds/details/multi-az/>

The correct answer is: True

**Feedback about this question and answer**

**Question 8**

Not answered

 Mark for Review

Which of the below instances is used normally for massive parallel computations?

Please select :

- ☐ A. Spot Instances
- ☐ B. On-Demand Instances
- ☐ C. Dedicated Instances
- ☐ D. This is not possible in AWS

Your answer is incorrect.

Answer – A

This is clearly given in the aws documentation

### Improve Application Throughput

You can run and scale applications such as stateless web services, image rendering, big data analytics and massively parallel computations on Spot instances. Since Spot instances typically cost 50-90% less, you can increase your compute capacity by 2-10x within the same budget.

For more information on Spot Instances, please visit the link –


<https://aws.amazon.com/ec2/spot/>

The correct answer is: Spot Instances

**Feedback about this question and answer**

**Question 9**

Not answered

 Mark for Review

Which of the below are incremental backups of your EBS volumes? Choose one answer from the options given below.


Please select :

- ☐ A. Volumes
- ☐ B. State Manager
- ☐ C. Placement Groups
- ☐ D. Snapshots

Your answer is incorrect.

Answer - D

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.



For more information on EBS snapshots, please visit the link -

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

The correct answer is: Snapshots

**Feedback about this question and answer**

#### Question 10

Not answered

Mark for Review

There is a requirement to host a NoSQL database with a need for low latency.

Which class of instances from the below list should they choose from ?

Please select :

- ☐ A. T2
- ☐ B. I2
- ☐ C. T1
- ☐ D. G2

Your answer is incorrect.

Answer - B

I2 instances are optimized to deliver tens of thousands of low-latency, random I/O operations per second (IOPS) to applications. They are well suited for the following scenarios:

- NoSQL databases (for example, Cassandra and MongoDB)
- Clustered databases
- Online transaction processing (OLTP) systems

For more information on I2 instances, please visit the link

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/i2-instances.html>

The correct answer is: I2

**Feedback about this question and answer**

#### Question 11

Not answered

Mark for Review

Which of the below elements can you manage in the IAM dashboard ?

Select 2 options.

Please select :

- ☐ A. Roles
- ☐ B. EC2
- ☐ C. Cost Allocation Reports
- ☐ D. Policies

Your answer is incorrect.

Answer – A and D

When you go to your IAM dashboard, below are the set of elements which can be configured.

**Dashboard**

Groups

Users

Roles

Policies

Identity providers

Account settings

Credential report

Encryption keys

For more information on IAM, please visit the link


<http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>

The correct answer is: Roles, Policies

**Feedback about this question and answer**

**Question 12**

Not answered

 Mark for Review

Which is the service provided by AWS for controlling encryption keys in AWS?

Please select :

- ☐ A. SNS
- ☐ B. SES
- ☐ C. KMS
- ☐ D. None of the above

Your answer is incorrect.

Answer - C

AWS Key Management Service (KMS) is a managed service that makes it easy for you to create and control the encryption keys used to encrypt your data, and uses Hardware Security Modules (HSMs) to protect the security of your keys. AWS Key Management Service (KMS) is integrated with several other AWS services to help you protect the data you store with these services. AWS Key Management Service is also integrated with AWS CloudTrail to provide you with logs of all key usage to help meet your regulatory and compliance needs.

Option A is wrong because it is the Notification service from aws.

Option B is wrong because it is the emailing service from aws.

For more information on AWS Key Management Service, please visit the below url


- [https://aws.amazon.com/kms/?nc2=h\\_m1](https://aws.amazon.com/kms/?nc2=h_m1)

The correct answer is: KMS

**Feedback about this question and answer**

**Question 13**

Not answered

 Mark for Review

Which of the below elements can you manage in the Billing dashboard ?

Select 2 options.

Please select :

- ☐ A. Budgets
- ☐ B. Policies
- ☐ C. Credential Report
- ☐ D. Cost Explorer

Your answer is incorrect.

Answer – A and D

When you go to your Billing dashboard, below are the set of elements which can be configured.

**Dashboard**

Bills  
Cost Explorer  
Budgets  
Reports  
Cost Allocation Tags  
Payment Methods  
Payment History  
Consolidated Billing  
Preferences  
Credits  
Tax Settings  
DevPay

For more information on AWS cloud billing and pricing, please visit the link


<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/billing-getting-started.html>

The correct answer is: Budgets, Cost Explorer

**Feedback about this question and answer**

**Question 14**

Not answered

 Mark for Review

What is the name of VPC that is automatically created for your AWS account for the first time ?

Please select :

- ☐ A. Primary VPC
- ☐ B. First VPC
- ☐ C. Default VPC
- ☐ D. Initial VPC

Your answer is incorrect.

Answer – C

A default VPC is a logically isolated virtual network in the AWS cloud that is automatically created for your AWS account the first time you provision Amazon EC2 resources. When you launch an instance without specifying a subnet-ID, your instance will be launched in your default VPC.

## Default VPCs

### Q. What is a default VPC?

A default VPC is a logically isolated virtual network in the AWS cloud that is automatically created for your AWS account the first time you provision Amazon EC2 resources. When you launch an instance without specifying a subnet-ID, your instance will be launched in your default VPC.

### Q. What are the benefits of a default VPC?

When you launch resources in a default VPC, you can benefit from the advanced networking functionalities of Amazon VPC (EC2-VPC) with the ease of use of Amazon EC2 (EC2-Classical). You can enjoy features such as changing security group membership on the fly, security group egress filtering, multiple IP addresses, and multiple network interfaces without having to explicitly create a VPC and launch instances in the VPC.

For more information on VPC, please refer to the link


- <https://aws.amazon.com/vpc/faqs/>

The correct answer is: Default VPC

### Feedback about this question and answer

#### Question 15

Not answered

 Mark for Review

Which of the following databases that supports read replica feature?

Select 3 options.

Please select :

- ☐ A. MySQL
- ☐ B. MariaDB
- ☐ C. PostgreSQL
- ☐ D. Oracle

Your answer is incorrect.

Answer – A, B and C

Read replicas are available in Amazon RDS for MySQL, MariaDB, and PostgreSQL. When you create a read replica, you specify an existing DB Instance as the source. Amazon RDS takes a snapshot of the source instance and creates a read-only instance from the snapshot. For MySQL, MariaDB and PostgreSQL, Amazon RDS uses those engines' native asynchronous replication to update the read replica whenever there is a change to the source DB instance.

The read replica operates as a DB instance that allows only read-only connections; applications can connect to a read replica just as they would to any DB instance. Amazon RDS replicates all databases in the source DB instance.

For more information on rds read replicas, please refer to the link


- <https://aws.amazon.com/rds/details/read-replicas/>

The correct answer is: MySQL, MariaDB, PostgreSQL

### Feedback about this question and answer

#### Question 16

Not answered

 Mark for Review

What can be used from AWS to import existing virtual machines into AWS? Choose one answer from the options given below.

Please select :

- ☐ A. VM Import/Export
- ☐ B. AWS Import/Export
- ☐ C. AWS Storage Gateway
- ☐ D. This is not possible in AWS

Your answer is incorrect.

Answer – A

VM Import/Export enables customers to import Virtual Machine (VM) images in order to create Amazon EC2 instances. Customers can also export previously imported EC2 instances to create VMs. Customers can use VM Import/Export to leverage their previous investments in building VMs by migrating their VMs to Amazon EC2.

For more information on AWS EC2, please visit


- <https://aws.amazon.com/ec2/faqs/>

The correct answer is: VM Import/Export

### Feedback about this question and answer

#### Question 17

Not answered

 Mark for Review

What is the service used by AWS to segregate control over the various AWS services ?

Please select :

- ☐ A. AWS RDS
- ☐ B. AWS Integrity Management
- ☐ C. AWS Identity and Access Management
- ☐ D. Amazon EMR



Your answer is incorrect.

Answer – C

AWS Identity and Access Management (IAM) is a web service that helps you securely control access to AWS resources for your users. You use IAM to control who can use your AWS resources (authentication) and what resources they can use and in what ways (authorization).

For more information on IAM, please visit:


- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>

The correct answer is: AWS Identity and Access Management

**Feedback about this question and answer**

**Question 18**

Not answered

 Mark for Review

Which of the following programming languages have an officially supported AWS SDK?

Select 2 options.

Please select :

- ☐ A. PHP
- ☐ B. Pascal
- ☐ C. Java
- ☐ D. SQL
- ☐ E. Perl

Your answer is incorrect.

Answer - A and C

This is as per the AWS documentation

**Java**

[Install »](#)

[Documentation »](#)

[Learn more »](#)

**.NET**

[Install »](#)

[Documentation »](#)

[Learn more »](#)

**Node.js**

[Install »](#)

[Documentation »](#)

[Learn more »](#)

**PHP**

[Install »](#)

[Documentation »](#)

[Learn more »](#)

**Python**

[Install »](#)

[Documentation »](#)

[Learn more »](#)

**Ruby**

[Install »](#)

[Documentation »](#)

[Learn more »](#)

**Go**

[Install »](#)

[Documentation »](#)

**C++**

[Install »](#)

[Documentation »](#)

**AWS Mobile SDK**

[Install »](#)

[Documentation »](#)

For more information on aws toolkits available, please refer to the below url:


- <https://aws.amazon.com/tools/>

The correct answer is: PHP, Java

**Feedback about this question and answer**

**Question 19**

Not answered

 Mark for Review

How long can messages live in a SQS queue ?

Please select :

- ☐ A. 12 hours
- ☐ B. 10 days
- ☐ C. 14 days
- ☐ D. 1 year

Your answer is incorrect.

Answer – C

This is clearly given in the AWS documentation

**Q: How long can I keep my messages in Amazon SQS message queues?**

Longer message retention provides greater flexibility to allow for longer intervals between message production and consumption.

You can configure the Amazon SQS message retention period to a value from 1 minute to 14 days. The default is 4 days. Once the message retention limit is reached, your messages are automatically deleted.

For more information on SQS , please visit the following url


- <https://aws.amazon.com/sqs/faqs/>

The correct answer is: 14 days

**Feedback about this question and answer**

**Question 20**

Not answered

 Mark for Review

True or False: By default in AWS, the storage type for a read replica is the same as the storage type for the source DB.

Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer – A

This is clearly provided in the AWS documentation

By default, a Read Replica is created with the same storage type as the source DB instance. However, you can create a Read Replica that has a different storage type from the source DB instance based on the options listed in the following table.

Source DB Instance Storage Type	Source DB Instance Storage Allocation	Read Replica Storage Type Options
PIOPS	100 GB - 3 TB	PIOPS   GP2   Standard
GP2	100 GB - 3 TB	PIOPS   GP2   Standard
GP2	Less than 100 GB	GP2   Standard
Standard	100 GB - 3 TB	PIOPS   GP2   Standard
Standard	Less than 100 GB	GP2   Standard

For more information on read replica please visit the below link


- [http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER\\_ReadRepl.html](http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html)

The correct answer is: True

**Feedback about this question and answer**

**Question 21**

Not answered

 Mark for Review

What are some of the common causes why you cannot connect to a DB instance on AWS ?

Select 3 options.

Please select :

- ☐ A. There is a read replica being created, hence you cannot connect
- ☐ B. The DB is still being created
- ☐ C. The local firewall is stopping the communication traffic
- ☐ D. The security groups for the DB are not properly configured.

Your answer is incorrect.

Answer – B, C and D

There are some steps clearly given in the AWS documentation:

When you cannot connect to a DB instance, the following are common causes:

- The access rules enforced by your local firewall and the ingress IP addresses that you authorized to access your DB instance in the instance's security group are not in sync. The problem is most likely the ingress rules in your security group. By default, DB instances do not allow access; access is granted through a security group. To grant access, you must create your own security group with specific ingress and egress rules for your situation. For more information about setting up a security group, see [Provide Access to the DB Instance in the VPC by Creating a Security Group](#).
- The port you specified when you created the DB instance cannot be used to send or receive communications due to your local firewall restrictions. In this case, check with your network administrator to determine if your network allows the specified port to be used for inbound and outbound communication.
- Your DB instance is still being created and is not yet available. Depending on the size of your DB instance, it can take up to 20 minutes before an instance is available.

For more information on rds troubleshooting please visit the below link


- [http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP\\_Troubleshooting.html](http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Troubleshooting.html)

The correct answer is: The DB is still being created, The local firewall is stopping the communication traffic, The security groups for the DB are not properly configured.

**Feedback about this question and answer**

**Question 22**

Not answered

 Mark for Review

For which of the following databases does Amazon RDS provides high availability and failover support using Amazon's failover technology for DB instances using Multi-AZ deployments.

Select 3 options.

Please select :

- ☐ A. SQL Server
- ☐ B. MySQL
- ☐ C. Oracle
- ☐ D. MariaDB

Your answer is incorrect.

Answer – B, C and D

This is clearly provided in the aws documentation.

Amazon RDS provides high availability and failover support for DB instances using Multi-AZ deployments. Multi-AZ deployments for Oracle, PostgreSQL, MySQL, and MariaDB DB instances use Amazon technology, while SQL Server DB instances use SQL Server Mirroring.

For more information on MultiAZ please visit the below link


<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>

The correct answer is: MySQL, Oracle, MariaDB

**Feedback about this question and answer**

**Question 23**

Not answered

 Mark for Review

Which is the service provided by AWS for providing a petabyte-scale data warehouse ?

Please select :

- ☐ A. Amazon DynamoDB
- ☐ B. Amazon Redshift
- ☐ C. Amazon Kinesis
- ☐ D. Amazon Simple Queue Service

Your answer is incorrect.

Answer – B

Amazon Redshift is a fast, fully managed, petabyte-scale data warehouse that makes it simple and cost-effective to analyze all your data using your existing business intelligence tools. Start small for \$0.25 per hour with no commitments and scale to petabytes for \$1,000 per terabyte per year, less than a tenth the cost of traditional solutions.

Option A is wrong because it is used as a NOSQL solution.

Option C is wrong because it is used for processing streams and not for storage.

Option D is wrong because it is a de-coupling solution.

For more information on Redshift, please visit the below url


- [https://aws.amazon.com/redshift/?nc2=h\\_m1](https://aws.amazon.com/redshift/?nc2=h_m1)

The correct answer is: Amazon Redshift

**Feedback about this question and answer**

**Question 24**

Not answered

 Mark for Review

An image named photo.jpg has been uploaded to a bucket named examplebucket in the us-east-1 region. Which of the below is the right URL to access the image, if it were made public ? Consider that S3 is used as a static website.

Please select :

- ☐ A. <http://examplebucket.s3-website-us-east-1.amazonaws.com/photo.jpg>
- ☐ B. <http://examplebucket.website-us-east-1.amazonaws.com/photo.jpg>
- ☐ C. <http://examplebucket.s3-us-east-1.amazonaws.com/photo.jpg>
- ☐ D. <http://examplebucket.amazonaws.s3-website-us-east-1/photo.jpg>

Your answer is incorrect.

Answer - A

The URL for a S3 web site is as shown below

<bucket-name>.s3-website-<AWS-region>.amazonaws.com

Hence the right option is option A

For more information on the bucket and the URL format for S3 buckets , please visit the below url:


- <http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteEndpoints.html>
- <http://docs.aws.amazon.com/AmazonS3/latest/dev/HostingWebsiteOnS3Setup.html>

The correct answer is: <http://examplebucket.s3-website-us-east-1.amazonaws.com/photo.jpg>

**Feedback about this question and answer**

**Question 25**

Not answered

 Mark for Review

A company has an EC2 instance that is hosting a web solution which is mostly used for read-only purposes. The CPU utilization is constantly 100% on the EC2 instance. Which of the below solutions can help alleviate and provide a quick resolution to the problem.

Please select :

- ☐ A. Use Cloudfront and place the EC2 instance as the origin
- ☐ B. Let the EC2 instance continue to run at 100%, since the AWS environment can handle the load.
- ☐ C. Use SNS to notify the IT admin when it reaches 100% so that they can disconnect some sessions to help alleviate the load
- ☐ D. Use SES to notify the IT admin when it reaches 100% so that they can disconnect some sessions to help alleviate the load

Your answer is incorrect.

Answer – A

Cloudfront can be used alleviate the load on web based solutions by caching the recent reads in its edge locations and reduce the burden on the EC2 instance.

Amazon CloudFront is a global content delivery network (CDN) service that accelerates delivery of your websites, APIs, video content or other web assets

For more information on AWS Cloudfront please visit the below url


<https://aws.amazon.com/cloudfront/>

The correct answer is: Use Cloudfront and place the EC2 instance as the origin

**Feedback about this question and answer**

**Question 26**

Not answered

 Mark for Review

Which of the below mentioned AWS services uses the concept of shards that is uniquely identified group of data records in a stream.

Please select :

- ☐ A. Cloudfront

- ☐ B. SQS
- ☐ C. Kinesis
- ☐ D. SES

Your answer is incorrect.

Answer – C

In Amazon Kinesis, a shard is a uniquely identified group of data records in a stream. A stream is composed of one or more shards, each of which provides a fixed unit of capacity. Each shard can support up to 5 transactions per second for reads, up to a maximum total data read rate of 2 MB per second and up to 1,000 records per second for writes, up to a maximum total data write rate of 1 MB per second (including partition keys).

For more information on AWS Kinesis please visit the below url

- <http://docs.aws.amazon.com/streams/latest/dev/key-concepts.html>

The correct answer is: Kinesis

**Feedback about this question and answer**

Question **27**

Not answered

Mark for Review

Which of the below mentioned services are the building blocks for creating a basic high availability architecture in AWS.

Select 2 options.

Please select :

- ☐ A. EC2
- ☐ B. SQS
- ☐ C. Elastic Load Balancer
- ☐ D. Cloudwatch

Your answer is incorrect.

Answer – A and C

Having EC2 instances hosting your applications in multiple subnets, hence multiple AZ's and placing them behind an ELB is the basic building block of a high availability architecture in AWS.



For more information on High availability and Fault tolerance please visit the below url

[https://media.amazonwebservices.com/architecturecenter/AWS\\_ac\\_ra\\_ftha\\_04.pdf](https://media.amazonwebservices.com/architecturecenter/AWS_ac_ra_ftha_04.pdf)

The correct answer is: EC2, Elastic Load Balancer

**Feedback about this question and answer**

Question **28**

Not answered

Mark for Review

If you want to change the AMI in Autoscaling, in which of the following sections would you make the change ?

Please select :

- ☐ A. Desired Capacity
- ☐ B. Launch Configuration
- ☐ C. Groups
- ☐ D. Scaling Plans

Your answer is incorrect.

Answer – B

Your group uses a launch configuration as a template for its EC2 instances. When you create a launch configuration, you can specify information such as the AMI ID, instance type, key pair, security groups, and block device mapping for your instances

For more information on Autoscaling please visit the below url


<http://docs.aws.amazon.com/autoscaling/latest/userguide/WhatIsAutoScaling.html>

The correct answer is: Launch Configuration

**Feedback about this question and answer**

Question **29**

Not answered

 Mark for Review

Which of the following services provides an object store which can also be used to store files ?

Please select :

- ☐ A. S3
- ☐ B. SQS
- ☐ C. SNS
- ☐ D. EC2

Your answer is incorrect.

Answer – A

Amazon Simple Storage Service is storage for the Internet. It is designed to make web-scale computing easier for developers.

Amazon S3 has a simple web services interface that you can use to store and retrieve any amount of data, at any time, from anywhere on the web. It gives any developer access to the same highly scalable, reliable, fast, inexpensive data storage infrastructure that Amazon uses to run its own global network of web sites.

For more information on S3 please visit the below url


<http://docs.aws.amazon.com/AmazonS3/latest/dev/Welcome.html>

The correct answer is: S3

**Feedback about this question and answer**

Question **30**

Not answered

 Mark for Review

What are the different types of scale out options available in the AutoScaling service provided by AWS?

Select 3 options.

Please select :

- ☐ A. Scheduled Scaling
- ☐ B. Dynamic Scaling
- ☐ C. Manual Scaling
- ☐ D. Static Scaling

Your answer is incorrect.

Answer – A, B and C

This is clearly given in the AWS documentation:

#### Scale Out

The following scale out events direct the Auto Scaling group to launch EC2 instances and attach them to the group:

- You manually increase the size of the group. For more information, see [Manual Scaling](#).
- You create a scaling policy to automatically increase the size of the group based on a specified increase in demand. For more information, see [Dynamic Scaling](#).
- You set up scaling by schedule to increase the size of the group at a specific time. For more information, see [Scheduled Scaling](#).

For more information on Autoscaling please visit the below URL:


- <http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html>

The correct answer is: Scheduled Scaling, Dynamic Scaling, Manual Scaling

**Feedback about this question and answer**

**Question 31**

Not answered

 Mark for Review

Which of the following services provides edge locations that can be used to cache frequently accessed pages of a web application ?

Please select :

- ☐ A. SQS
- ☐ B. Cloudfront
- ☐ C. Subnets
- ☐ D. EC2

Your answer is incorrect.

Answer - B

Amazon CloudFront is a global content delivery network (CDN) service that accelerates delivery of your websites, APIs, video content or other web assets

For more information on AWS Cloudfront please visit the below url


<https://aws.amazon.com/cloudfront/>

The correct answer is: Cloudfront

**Feedback about this question and answer**

**Question 32**

Not answered

 Mark for Review

True or False : A subnet spans across Multiple Availability Zones.

Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer – B

This is clearly given in the AWS documentation:

When you create a VPC, it spans all the Availability Zones in the region. After creating a VPC, you can add one or more subnets in each Availability Zone. When you create a subnet, you specify the CIDR block for the subnet, which is a subset of the VPC CIDR block. Each subnet must reside entirely within one Availability Zone and cannot span zones. Availability Zones are distinct locations that are engineered to be isolated from failures in other Availability Zones. By launching instances in separate Availability Zones, you can protect your applications from the failure of a single location. We assign a unique ID to each subnet.

**Q. Can a subnet span Availability Zones?**

**No. A subnet must reside within a single Availability Zone.**

For more information on VPC and Subnets please visit the below URL:


- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)
- <https://aws.amazon.com/vpc/faqs/>

The correct answer is: False

**Feedback about this question and answer**

**Question 33**

Not answered

 Mark for Review

True or False: All subnets within a VPC can route traffic between each other by default whether they are public or private?

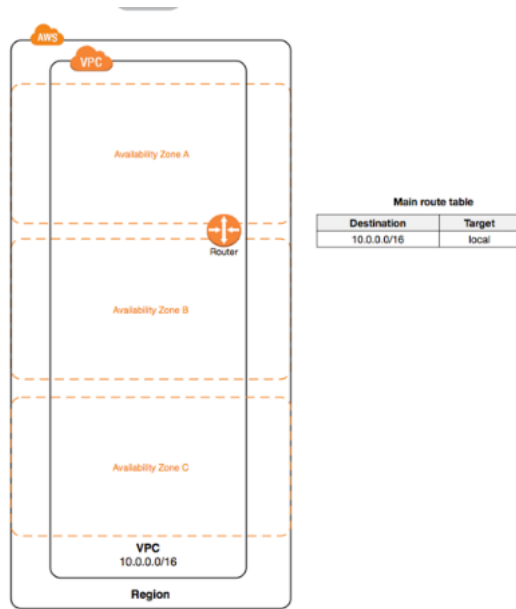
Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer – A

Yes , by default all subnets created within a VPC will have the Main route table attached to it which allows for communication on all subnets within the VPC.



For more information on VPC and Subnets please visit the below URL:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)

The correct answer is: True

**Feedback about this question and answer**

#### Question 34

Not answered

Mark for Review

Which is the service provided by AWS for collecting and processing large streams of data in real time?

Please select :

- ☐ A. Amazon Kinesis
- ☐ B. AWS Data Pipeline
- ☐ C. Amazon AppStream
- ☐ D. Amazon Simple Queue Service

Your answer is incorrect.

Answer – A

Use Amazon Kinesis Streams to collect and process large streams of data records in real time.

You'll create data-processing applications, known as *Amazon Kinesis Streams applications*. A typical Amazon Kinesis Streams application reads data from an *Amazon Kinesis stream* as data records. These applications can use the Amazon Kinesis Client Library, and they can run on Amazon EC2 instances. The processed records can be sent to dashboards, used to generate alerts, dynamically change pricing and advertising strategies, or send data to a variety of other AWS services

For more information on Amazon Kinesis, please visit this URL:

- <http://docs.aws.amazon.com/streams/latest/dev/introduction.html>

The correct answer is: Amazon Kinesis

**Feedback about this question and answer**

#### Question 35

Not answered

Mark for Review

Which of the following criteria are some of which must be met when attaching an EC2 instance to an existing AutoScaling Group ?

Select 3 options.

Please select :

- ☐ A. The instance is in the running state.
- ☐ B. The AMI used to launch the instance must still exist.
- ☐ C. The instance is not a member of another Auto Scaling group.
- ☐ D. They should have the same private key



Your answer is incorrect.

Answer – A, B and C

This is clearly mentioned in the AWS documentation:

### Attach EC2 Instances to Your Auto Scaling Group

Auto Scaling provides you with an option to enable Auto Scaling for one or more EC2 instances by attaching them to your existing Auto Scaling group. After the instances are attached, they become a part of the Auto Scaling group.

The instance that you want to attach must meet the following criteria:

- The instance is in the `running` state.
- The AMI used to launch the instance must still exist.
- The instance is not a member of another Auto Scaling group.
- The instance is in the same Availability Zone as the Auto Scaling group.
- If the Auto Scaling group has an attached load balancer, the instance and the load balancer must both be in EC2-Classic or the same VPC. If the Auto Scaling group has an attached target group, the instance and the Application Load Balancer must both be in the same VPC.

For more information on the criteria for attaching an EC2 instance to an existing AutoScaling group please visit the below URL:


- <http://docs.aws.amazon.com/autoscaling/latest/userguide/attach-instance-asg.html>

The correct answer is: The instance is in the running state., The AMI used to launch the instance must still exist., The instance is not a member of another Auto Scaling group.

**Feedback about this question and answer**

Question **36**

Not answered

 Mark for Review

A company wants to make use of serverless code. Which service in AWS provides such a facility?

Please select :

- ☐ A. SQS
- ☐ B. Cloudfront
- ☐ C. EC2
- ☐ D. Lambda

Your answer is incorrect.

Answer – D

AWS Lambda is a compute service that lets you run code without provisioning or managing servers. AWS Lambda executes your code only when needed and scales automatically, from a few requests per day to thousands per second. You pay only for the compute time you consume - there is no charge when your code is not running.

For more information on AWS Lambda please visit the below URL:


- <http://docs.aws.amazon.com/lambda/latest/dg/welcome.html>

The correct answer is: Lambda

**Feedback about this question and answer**

Question **37**

Not answered

 Mark for Review

Which of the following AWS services provides real-time data analytics ?

Please select :

- ☐ A. Kinesis
- ☐ B. Cloudfront
- ☐ C. Cloudwatch

Your answer is incorrect.

Answer - A

This is clearly mentioned in the AWS documentation:

#### Real-time data analytics

This combines the power of parallel processing with the value of real-time data. For example, process website clickstreams in real time, and then analyze site usability engagement using multiple different Streams applications running in parallel.

Amazon Kinesis Analytics is the easiest way to process streaming data in real time with standard SQL without having to learn new programming languages or processing frameworks. Amazon Kinesis Analytics enables you to query streaming data or build entire streaming applications using SQL, so that you can gain actionable insights and respond to your business and customer needs promptly.

Amazon Kinesis Analytics takes care of everything required to run your queries continuously and scales automatically to match the volume and throughput rate of your incoming data. With Amazon Kinesis Analytics, you only pay for the resources your queries consume. There is no minimum fee or setup cost.

For more information on AWS Kinesis please visit the below URL:


- <http://docs.aws.amazon.com/streams/latest/dev/introduction.html>
- <https://aws.amazon.com/kinesis/analytics/>

The correct answer is: Kinesis

**Feedback about this question and answer**

#### Question 38

Not answered

 Mark for Review

True or False: NACLs are defined for subnets stateful in nature

Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer – B

Network ACLs are stateless; responses to allowed inbound traffic are subject to the rules for outbound traffic (and vice versa).

For more information on AWS NACL please visit the below url


- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_ACLs.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html)

The correct answer is: False

**Feedback about this question and answer**

#### Question 39

Not answered

 Mark for Review

You have some EC2 instances hosted in your AWS environment. You have a concern that not all of the EC2 instances are being utilized.

Which of the below mentioned services can help you find underutilized resources in AWS ?

Select 2 options.

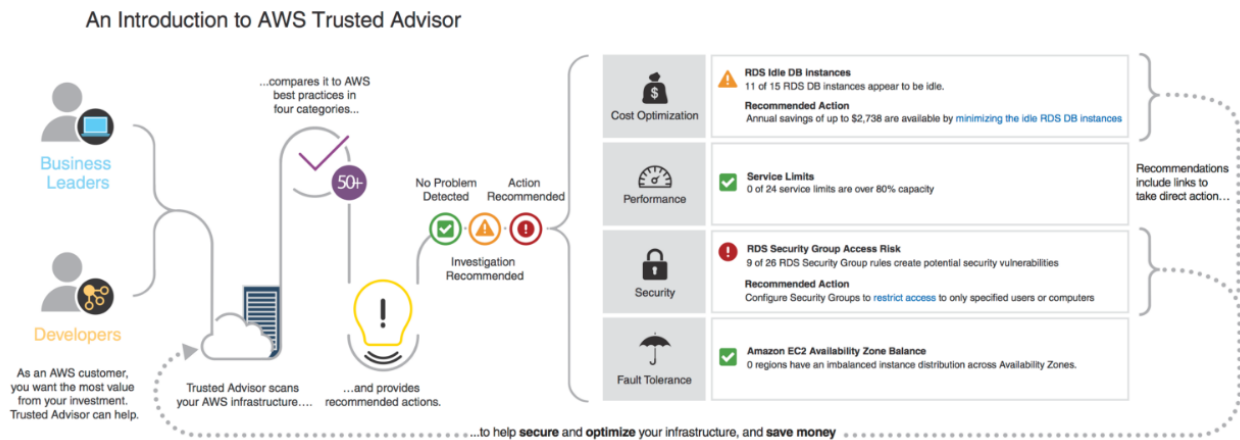
Please select :

- ☐ A. AWS Cloudwatch
- ☐ B. SNS
- ☐ C. AWS Trusted Advisor
- ☐ D. Cloudtrail

Your answer is incorrect.

Answer – A and C

The AWS Trusted Advisor can help you identify underutilized resources in AWS.



For more information on AWS trusted advisor please visit the below URL:

- <https://aws.amazon.com/premiumsupport/trustedadvisor/>

If You look at the Cloudwatch graphs, the CPU utilization of your resources and you can see the trend over time in the graphs.

For more information on AWS Cloudwatch please visit the below URL:

- <https://aws.amazon.com/cloudwatch/>

The correct answer is: AWS Cloudwatch, AWS Trusted Advisor

**Feedback about this question and answer**

Question 40

Not answered

Mark for Review

Which of the following features allows you to capture information about the IP traffic outgoing and incoming to network interfaces in your VPC ?

Please select :

- ☐ A. Cloudwatch
- ☐ B. EC2
- ☐ C. SQS
- ☐ D. Flow Logs

Your answer is incorrect.

Answer - D

VPC Flow Logs is a feature that enables you to capture information about the IP traffic going to and from network interfaces in your VPC. Flow log data is stored using Amazon CloudWatch Logs. After you've created a flow log, you can view and retrieve its data in Amazon CloudWatch Logs.

For more information on VPC Flowlogs please visit the below URL:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/flow-logs.html>

The correct answer is: Flow Logs

**Feedback about this question and answer**

Question 41

Not answered

Mark for Review

What are the main benefits of AWS regions?

Select 2 options.

Please select :

- ☐ A. Regions allow you to design applications to conform to specific laws and regulations for specific parts of the world.
- ☐ B. All regions offer the same service at the same prices.
- ☐ C. Regions allow you to choose a location in any country in the world.
- ☐ D. Regions allow you to place AWS resources in the area of the world closest to your customers who access those resources.

Your answer is incorrect.

Answer – A and D

AWS developer data centers across the world to help develop solutions that are close to the customer as possible. They also have center in core countries to help tie up with the specific laws and regulations for specific parts of the world.

AWS does not have centers in every location of the world, hence option C is invalid.

Services and prices are specific to every region, hence option B is invalid.

For more information on Regions please visit the below URL:


- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

The correct answer is: Regions allow you to design applications to conform to specific laws and regulations for specific parts of the world., Regions allow you to place AWS resources in the area of the world closest to your customers who access those resources.

**Feedback about this question and answer**

**Question 42**

Not answered

 Mark for Review

Which of the following are ways that users can interface with AWS?

Select 2 options.

Please select :

- ☐ A. AWS Cloudfront
- ☐ B. AWS CLI
- ☐ C. AWS Console
- ☐ D. AWS Cloudwatch

Your answer is incorrect.

Answer – B and C

**AWS Management Console** - The console is a browser-based interface to manage IAM and AWS resources.

**AWS Command Line Tools** - You can use the AWS command line tools to issue commands at your system's command line to perform IAM and AWS tasks; this can be faster and more convenient than using the console.

For more information on interfacing with AWS please visit the below URL:


- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>

The correct answer is: AWS CLI, AWS Console

**Feedback about this question and answer**

**Question 43**

Not answered

 Mark for Review

What are the two layers of security provided by AWS in the VPC?

Please select :

- ☐ A. Security Groups and NACLs
- ☐ B. NACLs and DHCP Options
- ☐ C. Route Tables and Internet gateway
- ☐ D. None of the above

Your answer is incorrect.

Answer – A

This is clearly given in the AWS documentation

## Security

Amazon VPC provides features that you can use to increase and monitor the security for your VPC:

- Security groups — Act as a firewall for associated Amazon EC2 instances, controlling both inbound and outbound traffic at the instance level
- Network access control lists (ACLs) — Act as a firewall for associated subnets, controlling both inbound and outbound traffic at the subnet level
- Flow logs — Capture information about the IP traffic going to and from network interfaces in your VPC

For more information on VPC Security please visit the below URL:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Security.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Security.html)

The correct answer is: Security Groups and NACLs

**Feedback about this question and answer**

### Question 44

Not answered

Mark for Review

A company wants to have a 50 Mbps dedicated connection to its AWS resources. Which of the below services can help fulfil this requirement ?

Please select :

- ☐ A. Virtual Private Gateway
- ☐ B. Virtual Private Connection (VPN)
- ☐ C. Direct Connect
- ☐ D. Internet Gateway

Your answer is incorrect.

Answer - C

AWS Direct Connect makes it easy to establish a dedicated network connection from your premises to AWS. Using AWS Direct Connect, you can establish private connectivity between AWS and your datacenter, office, or colocation environment, which in many cases can reduce your network costs, increase bandwidth throughput, and provide a more consistent network experience than Internet-based connections

For more information on AWS Direct Connect, please visit the below URL:

- <https://aws.amazon.com/directconnect/>

The correct answer is: Direct Connect

**Feedback about this question and answer**

### Question 45

Not answered

Mark for Review

What is the service name in AWS that can display costs in a chart format?

Please select :

- ☐ A. Cost Explorer
- ☐ B. Cost Allocation Tags
- ☐ C. AWS Budgets
- ☐ D. Payment History

Your answer is incorrect.

Answer – A

Cost Explorer is a free tool that you can use to view charts of your costs (also known as spend data) for up to the last 13 months, and forecast how much you are likely to spend for the next three months. You can use Cost Explorer to see patterns in how much you spend on AWS resources over time, identify areas that need further inquiry, and see trends that you can use to understand your costs. You can also specify time ranges for the data you want to see, and you can view time data by day or by month.

For more information on Cost Explorer, please visit the URL:

- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-explorer-what-is.html>

The correct answer is: Cost Explorer

**Feedback about this question and answer**

### Question 46

In the shared responsibility model, what is the customer not responsible for?

Not answered

Mark for Review

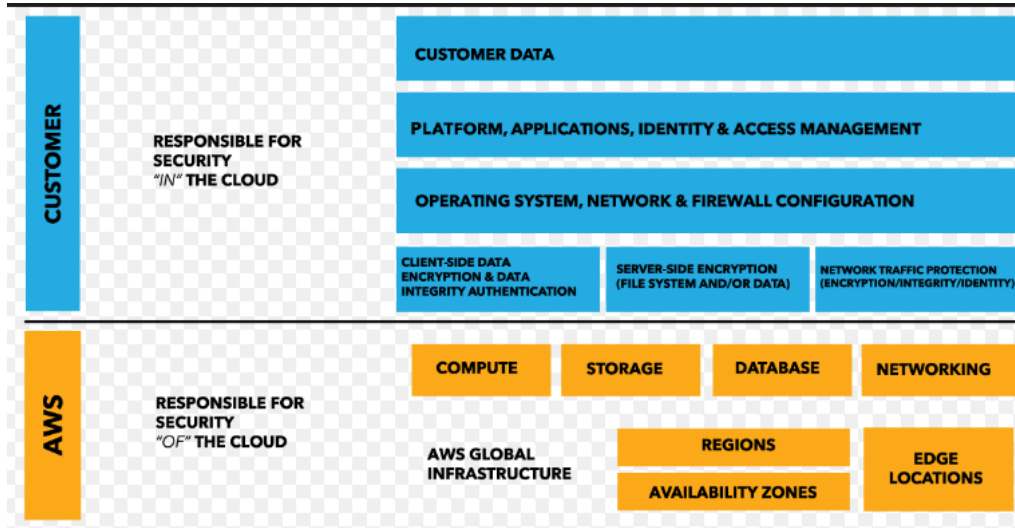
Please select :

- ☐ A. Edge locations
- ☐ B. Installation of custom firewall software
- ☐ C. Security Groups
- ☐ D. Applying an SSL Certificate to an ELB

Your answer is incorrect.

Answer - A

AWS has published the Shared Responsibility Model. And the Physical networking comes as part of the responsibility of AWS.



For more information on the Shared Responsibility Model, please refer to the below URL:

- <https://aws.amazon.com/compliance/shared-responsibility-model/>

The correct answer is: Edge locations

Feedback about this question and answer

Question 47

Not answered

Mark for Review

Which of the below options best describes how EBS snapshots work?

Please select :

- ☐ A. Each snapshot stores the entire volume in S3
- ☐ B. Snapshots are not possible for EBS volumes
- ☐ C. Snapshots are incremental in nature and are stored in S3
- ☐ D. Snapshots are stored in DynamoDB

Your answer is incorrect.

Answer – C

You can back up the data on your EBS volumes to Amazon S3 by taking point-in-time snapshots. Snapshots are incremental backups, which means that only the blocks on the device that have changed after your most recent snapshot are saved. This minimizes the time required to create the snapshot and saves on storage costs. When you delete a snapshot, only the data unique to that snapshot is removed.

For more information on EBS Snapshots, please refer to the below link

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

The correct answer is: Snapshots are incremental in nature and are stored in S3

Feedback about this question and answer

Question 48

Not answered

Mark for Review

True or False: You have database instances hosted in a private subnet. It Is possible to create an Elastic Load Balancer for private subnets.

Please select :

- ☐ A. True

☐ B. False

Your answer is incorrect.

Answer – A

When you create a load balancer in a VPC, you must choose whether to make it an internal load balancer or an Internet-facing load balancer.

The nodes of an Internet-facing load balancer have public IP addresses. The DNS name of an Internet-facing load balancer is publicly resolvable to the public IP addresses of the nodes. Therefore, Internet-facing load balancers can route requests from clients over the Internet. For more information, see Internet-Facing Classic Load Balancers.

The nodes of an internal load balancer have only private IP addresses. The DNS name of an internal load balancer is publicly resolvable to the private IP addresses of the nodes. Therefore, internal load balancers can only route requests from clients with access to the VPC for the load balancer.

For more information on Internal load balancers, please refer to the below URL:


- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-internal-load-balancers.html>

The correct answer is: True

**Feedback about this question and answer**

Question **49**

Not answered

 Mark for Review

Which of the following best describes the main feature of an Elastic Load Balancer (ELB) in AWS?

Please select :

- ☐ A. To evenly distribute traffic among multiple EC2 instances in separate Availability Zones.
- ☐ B. To evenly distribute traffic among multiple EC2 instances in a single Availability Zones.
- ☐ C. To evenly distribute traffic among multiple EC2 instances in a multiple regions.
- ☐ D. To evenly distribute traffic among multiple EC2 instances in a multiple counties.

Your answer is incorrect.

Answer – A

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve fault tolerance in your applications, seamlessly providing the required amount of load balancing capacity needed to route application traffic. The ELB is best used for EC2 instances multiple AZ's.

You cannot use ELB for distributing traffic across regions.

For more information on AWS ELB, please refer to the below link:


- <https://aws.amazon.com/elasticloadbalancing/>

The correct answer is: To evenly distribute traffic among multiple EC2 instances in separate Availability Zones.

**Feedback about this question and answer**

Question **50**

Not answered

 Mark for Review

What are the 2 main components of AutoScaling?

Select 2 options.

Please select :

- ☐ A. Launch Configuration
- ☐ B. Cloudtrail
- ☐ C. Cloudwatch
- ☐ D. AutoScaling Groups

Your answer is incorrect.

Answer – A and D

Groups - Your EC2 instances are organized into groups so that they can be treated as a logical unit for the purposes of scaling and management. When you create a group, you can specify its minimum, maximum, and, desired number of EC2 instances.

Launch configurations - Your group uses a launch configuration as a template for its EC2 instances. When you create a launch configuration, you can specify information such as the AMI ID, instance type, key pair, security groups, and block device mapping for your instances.

For more information on AWS Autoscaling, please refer to the below link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/WhatIsAutoScaling.html>

The correct answer is: Launch Configuration, AutoScaling Groups

**Feedback about this question and answer**

**Question 51**

Not answered

Mark for Review

What are two primary requirements of a NAT Instance? Choose the correct answer from the options below:

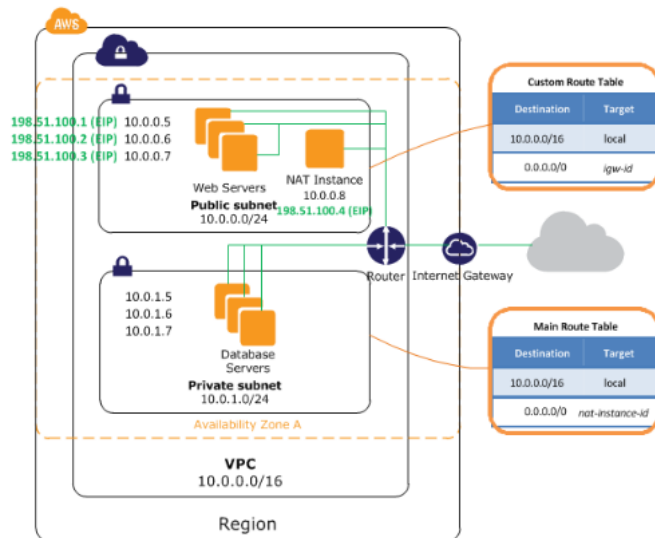
Please select :

- ☐ A. A NAT instance must be provisioned into a private subnet, and it must part of the private subnet's route table.
- ☐ B. A NAT instance must be provisioned into a public subnet, and it must part of the private subnet's route table.
- ☐ C. A NAT instance must be provisioned into a private subnet, and does not require a public IP address.
- ☐ D. A NAT instance must be provisioned into a public subnet, and must be combined with a bastion host.

Your answer is incorrect.

Answer – B

The below snapshot from the AWS documentation shows how the NAT instance is setup.



It needs to be placed in the public subnet and the private subnet should have a route to it.

For more information on AWS NAT Instance, please refer to the below link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_NAT\\_Instance.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_NAT_Instance.html)

The correct answer is: A NAT instance must be provisioned into a public subnet, and it must part of the private subnet's route table.

**Feedback about this question and answer**

**Question 52**

Not answered

Mark for Review

You have just provisioned a fleet of EC2 instances and realized that none of them have a public IP address. What settings would need to be changed for the next fleet of instances to be created with public IP addresses?

Please select :

- ☐ A. Modify the auto-assign public IP setting on the subnet.
- ☐ B. Modify the auto-assign public IP setting on the instance type.
- ☐ C. Modify the auto-assign public IP setting on the route table.
- ☐ D. Modify the auto-assign public IP setting on the VPC.

Your answer is incorrect.

Answer – A

This setting is done at the subnet level and if marked as true, all instances launched in that subnet will get a public IP address by default.



## Modify Auto-Assign Public IP



Enable auto-assign public IP to automatically request a public IP address for instances launched into this subnet.

☒ Enable auto-assign Public IP

Note: You can override the auto-assign public IP setting for each individual instance at launch time. Regardless of how you've configured the auto-assign public IP feature, you can assign a public IP address to an instance that has a single, new network interface with a device index of eth0.

Cancel

Save

For more information on AWS IP Addressing, please refer to the below link:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-ip-addressing.html>

The correct answer is: Modify the auto-assign public IP setting on the subnet.

**Feedback about this question and answer**

### Question 53

Not answered



Mark for Review

You keep getting an error while you try to attach an IGW to a VPC. What is the most likely cause of the error?

Please select :

- ☐ A. You need to have a customer gateway first before attaching an internet gateway
- ☐ B. You need to have a public subnet defined first before attaching an internet gateway
- ☐ C. You need to have a private subnet defined first before attaching an internet gateway
- ☐ D. An IGW is already attached to the VPC.

Your answer is incorrect.

Answer – D

You can only have one internet gateway attached to your VPC at one time, hence the error must be coming because there is already an internet gateway attached.

For more information on Internet gateways, please refer to the below link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Internet\\_Gateway.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html)

The correct answer is: An IGW is already attached to the VPC.

**Feedback about this question and answer**

### Question 54

Not answered



Mark for Review

What are the process to move objects from the standard storage class to Glacier?

Please select :

- ☐ A. S3 Events
- ☐ B. Object Versioning
- ☐ C. Selecting a different storage class
- ☐ D. Lifecycle Policies

Your answer is incorrect.

Answer – D

Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects. These actions can be classified as follows:

- Transition actions – In which you define when objects transition to another storage class. For example, you may choose to transition objects to the STANDARD\_IA (IA, for infrequent access) storage class 30 days after creation, or archive objects to the GLACIER storage class one year after creation.
- Expiration actions – In which you specify when the objects expire. Then Amazon S3 deletes the expired objects on your behalf.

For more information on Lifecycle policies, please refer to the below link:


- <http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>

The correct answer is: Lifecycle Policies

**Feedback about this question and answer**

Question **55**

Not answered

 Mark for Review

You work for a company that stores records for a minimum of 10 years. Most of these records will never be accessed but must be made available upon request (within a few hours). What is the most cost-effective storage option?

Please select :

- ☐ A. S3-IA
- ☐ B. Reduced Redundancy Storage (RRS)
- ☐ C. Glacier
- ☐ D. AWS Import/Export

Your answer is incorrect.

Answer – C

Amazon Glacier is a secure, durable, and extremely low-cost cloud storage service for data archiving and long-term backup. Customers can reliably store large or small amounts of data for as little as \$0.004 per gigabyte per month, a significant savings compared to on-premises solutions. To keep costs low yet suitable for varying retrieval needs, Amazon Glacier provides three options for access to archives, from a few minutes to several hours.

For more information on Amazon Glacier, please refer to the below link:


- <https://aws.amazon.com/glacier/>

The correct answer is: Glacier

**Feedback about this question and answer**

Question **56**

Not answered

 Mark for Review

What are the services from AWS helps to migrate databases to AWS easily ?

Please select :

- ☐ A. AWS Snowball
- ☐ B. AWS Direct Connect
- ☐ C. AWS Database Migration Service (DMS)
- ☐ D. None of the above

Your answer is incorrect.

Answer – C

AWS Database Migration Service helps you migrate databases to AWS easily and securely. The source database remains fully operational during the migration, minimizing downtime to applications that rely on the database. The AWS Database Migration Service can migrate your data to and from most widely used commercial and open-source databases.

For more information on aws database migration service, please visit the URL:

- <https://aws.amazon.com/dms>

The correct answer is: AWS Database Migration Service (DMS)

**Feedback about this question and answer**

Question **57**

Not answered

 Mark for Review

Your company has petabytes of data that it wants to move from their on-premise network to AWS. What AWS solution should you use?

Please select :

- ☐ A. AWS Direct Connect
- ☐ B. AWS Doesn't Support This Migration
- ☐ C. AWS Import/Export
- ☐ D. AWS Snowball

Your answer is incorrect.

Answer – D

Snowball is a petabyte-scale data transport solution that uses secure appliances to transfer large amounts of data into and out of the AWS cloud. Using Snowball addresses common challenges with large-scale data transfers including high network costs, long transfer times, and security concerns. Transferring data with Snowball is simple, fast, secure, and can be as little as one-fifth the cost of high-speed Internet.

For more information, please refer the below link:

- <http://docs.aws.amazon.com/snowball/latest/ug/whatissnowball.html>

The correct answer is: AWS Snowball

**Feedback about this question and answer**

Question **58**

Not answered

Mark for Review

There are multiple issues reported from an EC2 instance hence it is required to analyze the logs files. What can be used in AWS to store and analyze the log files?

Please select :

- ☐ A. SQS
- ☐ B. S3
- ☐ C. Cloudtrail
- ☐ D. Cloudwatch Logs

Your answer is incorrect.

Answer – D

You can use Amazon CloudWatch Logs to monitor, store, and access your log files from Amazon Elastic Compute Cloud (Amazon EC2) instances, AWS CloudTrail, and other sources. You can then retrieve the associated log data from CloudWatch Log. To enable Cloudwatch logs , follow the below steps

Step 1) Go to the Cloudwatch section and click on Logs

#### CloudWatch

Dashboards

Alarms

ALARM

INSUFFICIENT

OK

Billing

Events

Rules

Logs

Metrics **NEW**

Step 2) Once you create a log group, you can then configure your EC2 instance to send logs to this Log Group.

#### Create log group



Log Group

Name:

Demo

Cancel

Create log group

For more information on Cloudwatch logs please visit the link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/WhatIsCloudWatchLogs.html>

The correct answer is: Cloudwatch Logs

**Feedback about this question and answer**

Question **59**

Not answered

Mark for Review

You are not able to connect to an EC2 instance via SSH, and you have already verified that the instance has a public IP and the Internet gateway and route tables are in place, what should you check next?

Please select :

- ☐ A. Adjust the security group to allow traffic from port 22

- ☐ B. Adjust the security group to allow traffic from port 3389
- ☐ C. Restart the instance since there might be some issue with the instance
- ☐ D. Create a new instance since there might be some issue with the instance

Your answer is incorrect.

Answer – A

The reason why you cannot connect to the instance is because maybe the SSH protocol has not been enabled in the security group.

Go to your EC2 Security groups, click on the required security groups to make the changes. Go to the Inbound Tab. Ensure that the Inbound rules has a rule for SSH protocol.

Name	Group ID	Group Name	VPC ID
<input checked="" type="checkbox"/>	sg-042ebf79	launch-wizard-5	vpc-6dcc550a

**Security Group: sg-042ebf79**

Description Inbound Outbound Tags

Edit

Type	Protocol	Port Range
SSH	TCP	22

The correct answer is: Adjust the security group to allow traffic from port 22

**Feedback about this question and answer**

#### Question 60

Not answered

Mark for Review

Your company VPC has a need to communicate with another company VPC. What can be used from AWS to interface between the two VPC?

Please select :

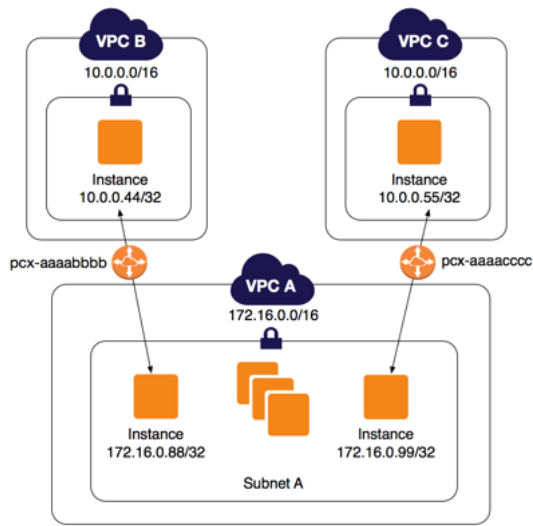
- ☐ A. VPC Connection
- ☐ B. VPN Connection
- ☐ C. Direct Connect
- ☐ D. VPC Peering

Your answer is incorrect.

Answer - D

A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IP addresses. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, or with a VPC in another AWS account within a single region

The below diagram shows an example of VPC peering. Now please note that VPC B cannot communicate to VPC C because there is no peering between them.



For more information on VPC peering, please visit the URL:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-peering.html>

The correct answer is: VPC Peering

**Feedback about this question and answer**

Finish review

#### QUIZ NAVIGATION

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15
16	17	18	19	20
21	22	23	24	25
26	27	28	29	30
31	32	33	34	35
36	37	38	39	40
41	42	43	44	45
46	47	48	49	50
51	52	53	54	55
56	57	58	59	60

Show all incorrect

Show all correct

Show one page at a time

Finish review