

2 MILLIONS+ USERS, 14+ YEARS, 100+ COUNTRIES

Visit our Website

Whizlabs Learning Center

HOME > AWS CSAA PRACTICE TESTS > FULL TEST(S) > NEW PRACTICE TEST II

Started on Wednesday, 5 July 2017, 3:27 AM

State Finished

Completed on Wednesday, 5 July 2017, 3:27 AM

Time taken 19 secs

Grade 0.00 out of 60.00 (**0**%)

Result FAIL

Ouestion 1

Not answered



Amazon EC2 provides a repository of public data sets that can be seamlessly integrated into AWS cloud-based applications. What is the monthly charge for using the public data sets?

Please select:

- A. 1 time charge of 1\$ for all the datasets.
- B. 1\$ per dataset per month
- C. 10 \$ per month for all datasets
- D. There is no charge for using public data sets

Your answer is incorrect.

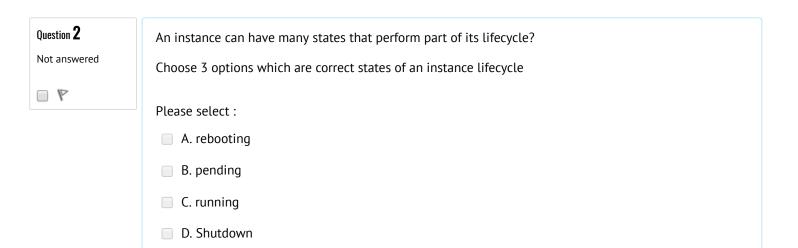
Answer - D

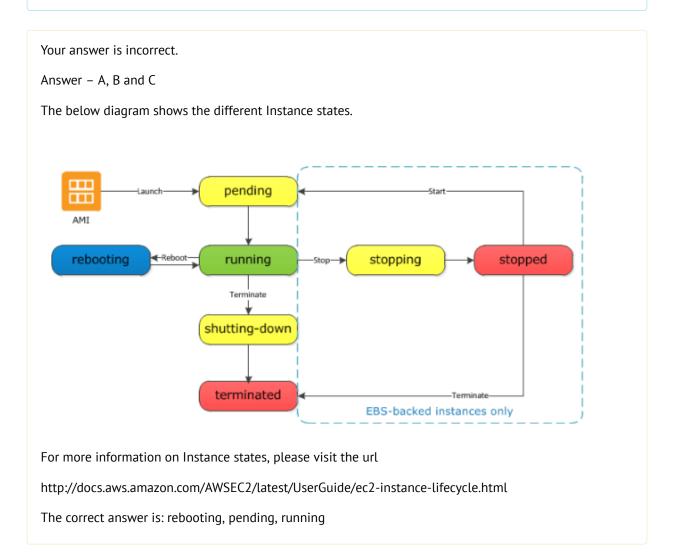
AWS hosts a variety of public datasets that anyone can access for free. Previously, large datasets such as the mapping of the Human Genome required hours or days to locate, download, customize, and analyze. Now, anyone can access these datasets via the AWS centralized data repository and analyze those using Amazon EC2 instances or Amazon EMR (Hosted Hadoop) clusters. By hosting this important data where it can be quickly and easily processed with elastic computing resources, AWS hopes to enable more innovation, more quickly.

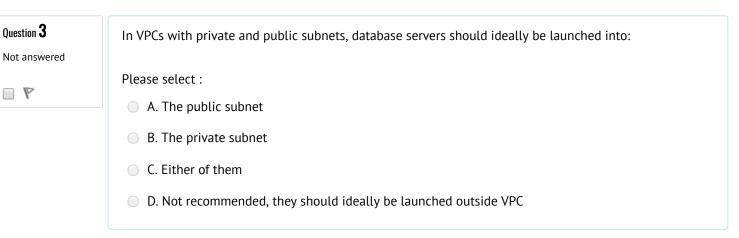
For more information on datasets please visit the below link

https://aws.amazon.com/public-datasets/

The correct answer is: There is no charge for using public data sets



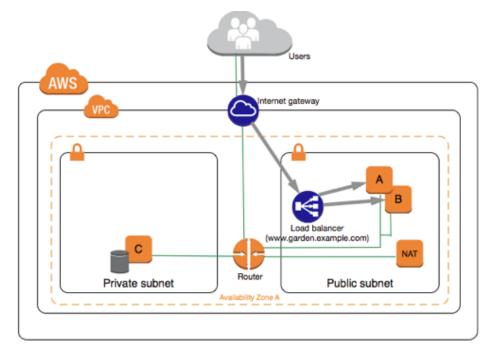




Answer - B

Normally database servers should not be exposed to the internet and should reside in private subnets.

The web servers will be part of the public subnet and exposed to the end users.



The correct answer is: The private subnet

Ouestion 4

Not answered



Which of the following can be used as an origin server in CloudFront?

Choose 3 answers from the options given below.

Please select:

- A. A webserver running on EC2
- B. A webserver running in your own datacenter
- C. A RDS instance
- D. An Amazon S3 bucket

Your answer is incorrect.

Answer - A, B and D

Currently Cloudfront supports the following types of distributions

- S3 buckets When you use Amazon S3 as an origin for your distribution, you place any objects that you want CloudFront to deliver in an Amazon S3 bucket.
- Custom Origin A custom origin is an HTTP server, for example, a web server. The HTTP server can be an Amazon EC2 instance or an HTTP server that you manage privately. When you use a

custom origin, you specify the DNS name of the server, along with the HTTP and HTTPS ports and the protocol that you want CloudFront to use when fetching objects from your origin.

For more information on Cloudfront Distributions, please visit the url

http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html

The correct answer is: A webserver running on EC2, A webserver running in your own datacenter, An Amazon S3 bucket

Ouestion **5**

Not answered



You have written a CloudFormation template that creates 1 elastic load balancer fronting 2 EC2 instances. Which section of the template should you edit so that the DNS of the load balancer is returned upon creation of the stack?

Please select:

- A. Resources
- B. Parameters
- C. Outputs
- D. Mappings

"Value":

Your answer is incorrect. Answer - C The below example shows a simple CloudFormation template. It creates an EC2 instance based on the AMI - ami-d6f32ab5. When the instance is created, it will output the AZ in which it is created. { "Resources": { "MyEC2Instance": { "Type": "AWS::EC2::Instance", "Properties": { "ImageId": "ami-d6f32ab5" } } }, "Outputs": { "Availability": { "Description": "The Instance ID",

```
{ "Fn::GetAtt" : [ "MyEC2Instance", "AvailabilityZone" ]}
}
}
}
```

To understand more on CloudFormation, please visit the url

https://aws.amazon.com/cloudformation/

The correct answer is: Outputs

Question **6**

Not answered



Is it true that EBS can always tolerate an Availability Zone failure?

Please select:

- A. No, all EBS volume is stored in a single Availability Zone
- B. Yes, EBS volume has multiple copies so it should be fine
- C. Depends on how it is setup
- D. Depends on the Region where EBS volume is initiated

Your answer is incorrect.

Answer - A

The EBS volumes are always associated with one AZ. So if the AZ fails, the EBS volume will also fail.

Option B is wrong because EBS volumes do not have multiple copies

Option C is wrong because no matter how it is setup, it always stays in one AZ

Option C is wrong because no matter in which region it is setup in, it always stays in one AZ

For more information on EBS volumes please visit the link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumes.html

The correct answer is: No, all EBS volume is stored in a single Availability Zone

Ouestion **7**

Not answered



Which of the following Auto scaling CANNOT do?

Choose 2 answers from the options given below

Please select:

- A. Start up EC2 instances when CPU utilization is above threshold
- B. Release EC2 instances when CPU utilization is below threshold
- C. Increase the instance size when utilization is above threshold

D. Decrease the instance size when utilization is below threshold

Your answer is incorrect.

Answer - C and D

You can only scale horizontally and not vertically. So, you can not increase or decrease the size of instances when the utilization is above threshold. As per the AWS documentation, below is what can be done with Auto Scaling?

Types of Scaling

- Vertical Scaling
 - · Changing instance size
 - Increasing EBS Capacity
- Horizontal Scaling
 - Adding / removing instances
 - ELB
- Scale out Amazon EC2 instances seamlessly and automatically when demand increases.
- Shed unneeded Amazon EC2 instances automatically and save money when demand subsides.
- Scale dynamically based on your Amazon CloudWatch metrics, or predictably according to a schedule that you define.
- Replace unhealthy or unreachable instances to maintain higher availability of your applications.
- Receive notifications via Amazon Simple Notification Service (Amazon SNS) to be alerted when
 you use Amazon CloudWatch alarms to initiate Auto Scaling actions, or when Auto Scaling
 completes an action.
- Run On-Demand or Spot Instances, including those inside your virtual private cloud (VPC) or high performance computing (HPC) clusters.
- If you're signed up for the Amazon EC2 service, you're already registered to use Auto Scaling and can begin using the feature via the API or command line interface.

For more information on Auto scaling please visit the link

https://aws.amazon.com/autoscaling/

The correct answer is: Increase the instance size when utilization is above threshold, Decrease the instance size when utilization is below threshold

Ouestion 8

Not answered



Which of the following benefits does adding Multi-AZ deployment in RDS provide?

Choose 2 answers from the options given below

Please select:

A. MultiAZ deployed database can tolerate an Availability Zone failure

B. Decrease latencies if app servers accessing database are in multiple Availability zones					
C. Make database access times faster for all app servers					
D. Make database more available during maintenance tasks					
Your answer is incorrect.					
Answer - A and D					
Some of the advantages of Multi AZ rds deployments are given below					
 If an Availability Zone failure or DB Instance failure occurs, your availability impact is limited to the time automatic failover takes to complete The availability benefits of Multi-AZ deployments also extend to planned maintenance and backups. In the case of system upgrades like OS patching or DB Instance scaling, these operations are applied first on the standby, prior to the automatic failover. As a result, your availability impact is, again, only the time required for automatic failover to complete. If a storage volume on your primary fails in a Multi-AZ deployment, Amazon RDS automatically initiates a failover to the up-to-date standby For more information on Multi AZ rds deployments please visit the link https://aws.amazon.com/rds/details/multi-az/ The correct answer is: MultiAZ deployed database can tolerate an Availability Zone failure, Make database more available during maintenance tasks 					
What happens to data when an EC2 instance terminates ?					
Select 3 options.					
Please select :					
 A. For EBS backed AMI, the EBS volume with operation system on it is preserved 					
■ B. For EBS backed AMI, any volume attached other than the OS volume is preserved					
C. All the snapshots of the EBS volume with operating system is preserved					

Answer - B, C and D

Question **9**Not answered

Option B is correct because when an instance is terminated, the volume will remain, unless you specifically delete the volume. When you create an instance, you have the root volume that does get deleted on deletion of the instance. But when you add a new volume, by default the "Delete on termination flag" is unchecked. So unless you don't check this, the volume will remain.

D. For S3 backed AMI, all the data in the local (ephemeral) hard drive is deleted

Option C is correct because this is the whole idea of snapshots to remain even if the volume or instance is deleted.

Option D is correct because ephemeral storage is temporary storage by default and gets deleted when the system is terminated.

For more information on EBS volumes, please visit the link http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumes.html

The correct answer is: For EBS backed AMI, any volume attached other than the OS volume is preserved, All the snapshots of the EBS volume with operating system is preserved, For S3 backed AMI, all the data in the local (ephemeral) hard drive is deleted

Question 10

Not answered



A company has the following EC2 instance configuration. They are trying to connect to the instance from the internet. They have verified the existence of the Internet gateway and the route tables are iin place. What could be the issue?

Please select:

- A. It's launched in the wrong Availability Zone
- B. The AMI used to launch the instance cannot be accessed from the internet
- C. The private IP is wrongly assigned
- D. There is no Elastic IP Assigned

Your answer is incorrect.

Answer - D

An instance must either have a public or Elastic IP in order to be accessible from the internet.

A public IP address is reachable from the Internet. You can use public IP addresses for communication between your instances and the Internet.

An Elastic IP address is a static IP address designed for dynamic cloud computing. An Elastic IP address is associated with your AWS account. With an Elastic IP address, you can mask the failure of an instance or software by rapidly remapping the address to another instance in your account.

An Elastic IP address is a public IP address, which is reachable from the Internet. If your instance does not have a public IP address, you can associate an Elastic IP address with your instance to enable communication with the Internet; for example, to connect to your instance from your local computer.

For more information on Elastic IP's, please visit the link http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/elastic-ip-addresses-eip.html

The correct answer is: There is no Elastic IP Assigned

Question 11	You have built an AMI. Another AWS account holder wants to use your AMI, but is not able to access
Not answered	it. What could be the issue? Below is the settings of the AMI in the AWS Console.
	Please select :
	A. It is not possible to share AMI's
	B. The AMI needs to made Public
	C. The Owner of AMI needs to be changed
	D. The AMI ID does not allow sharing
	Your answer is incorrect.
	Answer - B
	An Amazon Machine Image (AMI) provides the information required to launch an instance, which is a virtual server in the cloud.
	After you create an AMI, you can keep it private so that only you can use it, or you can share it with a
	specified list of AWS accounts. You can also make your custom AMI public so that the community can
	use it. To make the AMI public, carry out the following steps
	Step 1) Choose the AMI and then choose the menu option of Moidy Image Permissions
	Step 2) Make the Image as Public
	For more information on sharing an AMI to public, please visit the below URL:
	http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/sharingamis-intro.html
	The correct answer is: The AMI needs to made Public

Question **12**

Not answered



What is the basic requirement for logging into an EC2 instance on the AWS cloud?

Please select :

- A. Volumes
- B. AMI's
- C. Key Pairs
- D. S3

Answer - C

Amazon EC2 uses public-key cryptography to encrypt and decrypt login information. Public-key cryptography uses a public key to encrypt a piece of data, such as a password, then the recipient uses the private key to decrypt the data. The public and private keys are known as a key pair.

To log in to your instance, you must create a key pair, specify the name of the key pair when you launch the instance, and provide the private key when you connect to the instance. Linux instances have no password, and you use a key pair to log in using SSH. With Windows instances, you use a key pair to obtain the administrator password and then log in using RDP.

When you launch an EC2 instance, you will either be asked to create a new key pair or an existing key pair. This is .pem file which can then use to log into your instance.

For more information on Key Pairs please visit the below link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-key-pairs.html

The correct answer is: Key Pairs

Question 13

Not answered



Which of the below features allows you to take backups of your EBS volumes? Choose one answer from the options given below.

Please select:

- A. Volumes
- B. State Manager
- C. Placement Groups
- D. Snapshots

Your answer is incorrect.

Answer - D

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.

For more information on EBS snapshots, please visit the link http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html

The correct answer is: Snapshots

Question 14 Not answered	A company wants to host a selection of MongoDB instances. They are expecting a high load and want to have as low latency as possible. Which class of instances from the below list should they choose from.					
	Please select : A. T2					
	B. 12					
	C. T1					
	D. G2					
	D. G2					
	Your answer is incorrect.					
	Answer -B					
	I2 instances are optimized to deliver tens of thousands of low-latency, random I/O operations per second (IOPS) to applications. They are well suited for the following scenarios:					
	NoSQL databases (for example, Cassandra and MongoDB)					
	 Clustered databases Online transaction processing (OLTP) systems 					
	For more information on I2 instances, please visit the link					
	http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/i2-instances.html					
	The correct answer is: I2					
4=						
Question 15 Not answered	Which of the below elements can you manage in the IAM dashboard?					
	Choose 3 answers from the options given below					
	Please select :					
	A. Users					
	■ B. Encryption Keys					
	C. Cost Allocation Reports					
	D. Policies					
	Your answer is incorrect.					
	Answer – A, B and D					
	When you go to your IAM dashboard, below are the set of elements which can be configured.					

The correct answer is: Users, Encryption Keys, Policies Ouestion 16 What are the languages currently supported by AWS Lamda? Not answered Choose 3 answers from the options given below Please select: A. Node.js B. Angular JS C. Java D. Python Your answer is incorrect. Answer - A, C and D AWS Lambda supports code written in Node.js (JavaScript), Python, Java (Java 8 compatible), and C# (using the .NET Core runtime). For more information on Amazon Lambda, please visit https://aws.amazon.com/lambda/?nc2=h_m1 The correct answer is: Node.js, Java, Python Question 17 What is the number of Read Replica's that can be created for a particular Amazon RDS for MySQL Not answered Instance? Please select: A. 2 B. 5 C. 10 D. 20 Your answer is incorrect. Answer - B Here is the snapshot of the AWS documentation. For more information on AWS RDS, please visit

https://aws.amazon.com/rds/faqs/

The correct answer is: 5

Ouestion 18

Not answered



A company wants to pan out the steps visually on the automated steps that can be used to release their software. Which of the below AWS services makes this possible.

Please select:

- A. AWS CodeCommit
- B. AWS CodePipeline
- C. AWS Workflow
- D. AWS S3

Your answer is incorrect.

Answer - B

AWS CodePipeline is a continuous delivery service that enables you to model, visualize, and automate the steps required to release your software. With AWS CodePipeline, you model the full release process for building your code, deploying to pre-production environments, testing your application and releasing it to production. AWS CodePipeline then builds, tests, and deploys your application according to the defined workflow every time there is a code change. You can integrate partner tools and your own custom tools into any stage of the release process to form an end-to-end continuous delivery solution.

For more information on AWS CodePipeline, please visit

https://aws.amazon.com/codepipeline/faqs/

The correct answer is: AWS CodePipeline

Question 19

Not answered



In Cloudtrail, where does it store all of the logs that it creates? Choose one answer from the options given below.

Please select:

- A. A separate EC2 instance with EBS storage
- B. A RDS instance
- C. A DynamoDB instance
- D. Amazon S3

Your answer is incorrect.

Answer - D

When you enable Cloudtrail, you need to provide an S3 bucket where all the logs can be written to. For more information on AWS Cloudtrail, please visit https://aws.amazon.com/cloudtrail/ The correct answer is: Amazon S3 Ouestion 20 A customer wants to import their existing virtual machines to the cloud. Which service can they use Not answered for this? Choose one answer from the options given below. Please select: A. VM Import/Export B. AWS Import/Export C. AWS Storage Gateway D. This is not possible in AWS Your answer is incorrect. Answer - A VM Import/Export enables customers to import Virtual Machine (VM) images in order to create Amazon EC2 instances. Customers can also export previously imported EC2 instances to create VMs. Customers can use VM Import/Export to leverage their previous investments in building VMs by migrating their VMs to Amazon EC2. For more information on AWS EC2, please visit https://aws.amazon.com/ec2/faqs/ The correct answer is: VM Import/Export Ouestion 21 In the event of an unplanned outage of your primary DB, AWS RDS automatically switches over to the Not answered secondary. In such a case which record in Route 53 is changed? Select one answer from the options given below Please select: A. DNAME B. CNAME C. TXT

Answer - B

The AWS documentation clearly highlights what happens in the event of an automatic failover for an AWS RDS instance.

For more information on AWS RDS, please visit

https://aws.amazon.com/rds/faqs/

The correct answer is: CNAME

Ouestion 22

Not answered



Which of the below resources cannot be tagged in AWS

Please select:

- A. Images
- B. EBS Volumes
- C. Elastic IP
- D. VPC

Your answer is incorrect.

Answer - C

Tags enable you to categorize your AWS resources in different ways, for example, by purpose, owner, or environment. This is useful when you have many resources of the same type — you can quickly identify a specific resource based on the tags you've assigned to it. Each tag consists of a key and an optional value, both of which you define.

But you cannot tag an Elastic IP

For more information on AWS Resourcing Tagging, please visit

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using_Tags.html

The correct answer is: Elastic IP

Ouestion 23

Not answered



What is the key aspect for the below AMI image. Choose one answer from the options below

 A. Since it's a EBS volume AMI, it is special in nature 						
B. Since it's a Linux based AMI, it is special in nature						
C. Since it's a HVM based AMI , it is special in nature						
D. Since it's a NAT based AMI , it is special in nature						
Your answer is incorrect.						
Answer – D						
Amazon provides Amazon Linux AMIs that are configured to run as NAT instances. These AMIs						
include the string amzn-ami-vpc-nat in their names, so you can search for them in the Amazon EC2						
console. When you launch an instance from a NAT AMI, the following configuration occurs on the instance:						
 IPv4 forwarding is enabled and ICMP redirects are disabled in /etc/sysctl.d/10-nat-settings.conf A script located at /usr/sbin/configure-pat.sh runs at startup and configures iptables IP masquerading 						
For more information on NAT instances please visit the below link						
http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_NAT_Instance.html						
The correct answer is: Since it's a NAT based AMI, it is special in nature						
What type of monitoring for EBS volumes is available automatically in 5 minute periods at no						
charge?						
Please select :						
A. Basic						
■ B. Primary						
○ C. Detailed						
D. Local						

Answer – A

Question **24**

Not answered

Please select:

As per the AWS documentation, below is the types of monitoring data.

For more information on Volume monitoring, please visit

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/monitoring-volume-status.html

The correct answer is: Basic

Question **25**

Not answered



There is a company website that is going to be launched in the coming weeks. There is a probability that the traffic will be quite high in the first couple of weeks. In the event of a load failure, how can you set up DNS failover to a static website? Choose the correct answer from the options given below.

Please select:

- A. Duplicate the exact application architecture in another region and configure DNS weightbased routing
- B. Enable failover to an on-premise data center to the application hosted there.
- C. Use Route 53 with the failover option to failover to a static S3 website bucket or CloudFront distribution.
- D. Add more servers in case the application fails.

Your answer is incorrect.

Answer - C

Amazon Route 53 health checks monitor the health and performance of your web applications, web servers, and other resources.

If you have multiple resources that perform the same function, you can configure DNS failover so that Amazon Route 53 will route your traffic from an unhealthy resource to a healthy resource. For example, if you have two web servers and one web server becomes unhealthy, Amazon Route 53 can route traffic to the other web server. So you can route traffic to a website hosted on S3 or to a cloudfront distribution.

For more information on DNS failover using Route53, please refer to the below link

http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover.html

The correct answer is: Use Route 53 with the failover option to failover to a static S3 website bucket or CloudFront distribution.

Ouestion 26

Not answered



ot unswered

What is one of the major advantages of having a VPC in AWS?

Please select:

- A. You don't have to worry about security, this is managed by AWS.
- B. You can connect your cloud resources to on-premise data centers using VPN connections
- C. You can provision unlimited number of S3 resources.
- D. None of the above

Answer - B

One of the major advantages is that you can combine your on-premise data center to AWS via a VPN connection.

You can create an IPsec, hardware VPN connection between your VPC and your remote network. On the AWS side of the VPN connection, a *virtual private gateway* provides two VPN endpoints for automatic failover. You configure your *customer gateway*, which is the physical device or software application on the remote side of the VPN connection.

For more information on VPN connections, please refer to the below link

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html

The correct answer is: You can connect your cloud resources to on-premise data centers using VPN connections

Ouestion 27

Not answered



One of your instances is reporting an unhealthy system status check. However, this is not something you should have to monitor and repair on your own. How might you automate the repair of the system status check failure in an AWS environment? Choose the correct answer from the options given below

Please select:

- A. Create CloudWatch metrics that stop and start the instance based off of status check alarms
- B. Write a script that queries the EC2 API for each instance status check
- C. Write a script that periodically shuts down and starts instances based on certain stats.
- D. Implement a third party monitoring tool.

Your answer is incorrect.

Answer - A

Using Amazon CloudWatch alarm actions, you can create alarms that automatically stop, terminate, reboot, or recover your EC2 instances. You can use the stop or terminate actions to help you save money when you no longer need an instance to be running. You can use the reboot and recover actions to automatically reboot those instances or recover them onto new hardware if a system impairment occurs.

For more information on using alarm actions, please refer to the below link

http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/UsingAlarmActions.html

The correct answer is: Create CloudWatch metrics that stop and start the instance based off of status check alarms

Question 28

Not answered



A company is running three production web server reserved EC2 instances with EBS-backed root volumes. These instances have a consistent CPU load of 80%. Traffic is being distributed to these instances by an Elastic Load Balancer. They also have production and development Multi-AZ RDS MySQL databases. What recommendation would you make to reduce cost in this environment without affecting availability of mission-critical systems? Choose the correct answer from the options given below

Please select:

- A. Consider using on-demand instances instead of reserved EC2 instances
- B. Consider not using a Multi-AZ RDS deployment for the development database
- C. Consider using spot instances instead of reserved EC2 instances
- D. Consider removing the Elastic Load Balancer

Your answer is incorrect.

Answer - B

Multi-AZ databases is better for production environments rather than for development environments, so you can reduce costs by not using this for development environments

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure, and is engineered to be highly reliable. In case of an infrastructure failure, Amazon RDS performs an automatic failover to the standby (or to a read replica in the case of Amazon Aurora), so that you can resume database operations as soon as the failover is complete. Since the endpoint for your DB Instance remains the same after a failover, your application can resume database operation without the need for manual administrative intervention

For more information on Multi-AZ RDS, please refer to the below link

https://aws.amazon.com/rds/details/multi-az/

The correct answer is: Consider not using a Multi-AZ RDS deployment for the development database

Question 29

Not answered



A company has assigned two web server instances in a VPC subnet to an Elastic Load Balancer (ELB). However, the instances and the ELB are not reachable via URL to the Elastic Load Balancer (ELB). How can you resolve the issue so that your web server instances can start serving the web app data to the public Internet? Choose the correct answer from the options given below

Please select: A. Attach an Internet gateway to the VPC and route it to the subnet B. Add an elastic IP address to the instance C. Use Amazon Elastic Load Balancer to serve requests to your instances located in the internal subnet D. None of the above

Your answer is incorrect.

Answer - A

If the Internet gateway is not attached to the VPC, which is a pre-requisite for the instances to be accessed from the internet then the instances will not be reachable.

For more information on Internet gateways, please refer to the below link:

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html

The correct answer is: Attach an Internet gateway to the VPC and route it to the subnet

Question **30**Not answered



A company has EC2 instances running in AWS. The EC2 instances are running via an Autoscaling solution. There is a lot of requests being lost because of the load on the servers. The Autoscaling solution is launching new instances to take the load but there are still some requests which are being lost. Which of the following is likely to provide the most cost-effective solution to avoid losing recently submitted requests? Choose the correct answer from the options given below

Please select:

- A. Use an SQS queue to decouple the application components
- B. Keep one extra EC2 instance always powered on in case a spike occurs
- C. Use larger instances for your application
- D. Pre-warm your Elastic Load Balancer

Your answer is incorrect.

Answer - A

Amazon Simple Queue Service (SQS) is a fully-managed message queuing service for reliably communicating among distributed software components and microservices - at any scale. Building applications from individual components that each perform a discrete function improves scalability and reliability, and is best practice design for modern applications

For more information on SQS, please refer to the below link

https://aws.amazon.com/sqs/

The correct answer is: Use an SQS queue to decouple the application components

Question 31

Not answered



After migrating an application architecture from on-premise to AWS you will not be responsible for the ongoing maintenance of packages for the following AWS services that your application uses.

Choose the 2 correct answers from the options below

Please select:

- A. Elastic Beanstalk
- B. RDS
- C. DynamoDB
- D. EC2

Your answer is incorrect.

Answer - B and C

Both RDS and DynamoDB are managed solutions provided by AWS.

Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, freeing you up to focus on your applications and business.

For more information on RDS, please refer to the below link

https://aws.amazon.com/rds/

Amazon DynamoDB is a fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale. It is a fully managed cloud database and supports both document and key-value store models.

For more information on DynamoDB, please refer to the below link

https://aws.amazon.com/dynamodb/

The correct answer is: RDS, DynamoDB

Question 32

Not answered



What is the difference between an availability zone and an edge location? Choose the correct answer from the options below

Please select:

- A. Edge locations are used as control stations for AWS resources
- B. An edge location is used as a link when building load balancing between regions

C. An availability zone is an Amazon resource within an AWS region; an edge location will deliver cached content to the closest location to reduce latency
 D. An availability zone is a grouping of AWS resources in a specific region; an edge location is a specific resource within the AWS region
Your answer is incorrect.
Answer – C
Edge locations
Using a network of edge locations around the world, Amazon CloudFront caches copies of your static content close to viewers, lowering latency when they download your objects and giving you the high, sustained data transfer rates needed to deliver large popular objects to end users at scale.
For more information on Cloudfront and edge locations, please refer to the below link
https://aws.amazon.com/cloudfront/
Availability Zones
Each region is completely independent. Each Availability Zone is isolated, but the Availability Zones in a region are connected through low-latency links. The following diagram illustrates the relationship between regions and Availability Zones.
For more information on AZ, please refer to the below link
http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html
The correct answer is: An availability zone is an Amazon resource within an AWS region; an edge location will deliver cached content to the closest location to reduce latency
When storing sensitive data on the cloud which of the below options should be carried out on AWS?
Choose 3 answers from the options given below.
Please select :
A. With AWS you do not need to worry about encryption
B. Enable EBS Encryption
C. Encrypt the file system on an EBS volume using Linux tools
D. Enable S3 Encryption
Your answer is incorrect.
Answer – B,C and D

Question **33**Not answered

Amazon EBS encryption offers you a simple encryption solution for your EBS volumes without the need for you to build, maintain, and secure your own key management infrastructure. When you create an encrypted EBS volume and attach it to a supported instance type, the following types of data are encrypted:

- Data at rest inside the volume
- All data moving between the volume and the instance
- All snapshots created from the volume

For more information on EBS Encryption, please refer to the below link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html

Data protection refers to protecting data while in-transit (as it travels to and from Amazon S3) and at rest (while it is stored on disks in Amazon S3 data centers). You can protect data in transit by using SSL or by using client-side encryption.

For more information on S3 Encryption, please refer to the below link

http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingEncryption.html

The correct answer is: Enable EBS Encryption, Encrypt the file system on an EBS volume using Linux tools, Enable S3 Encryption

Question **34**

Not answered



What is an AWS service which can help protect web applications from common security threats from the outside world? Choose one answer from the options below

Please select:

- A. NAT
- B. WAF
- C. SQS
- D. SES

Your answer is incorrect.

Answer - B

Option A is wrong because this is used to relay information from private subnets to the internet.

Option C is wrong because this is used as a queuing service in aws.

Option D is wrong because this is used as an emailing service in aws.

AWS WAF is a web application firewall that helps protect your web applications from common web exploits that could affect application availability, compromise security, or consume excessive resources. AWS WAF gives you control over which traffic to allow or block to your web applications

by defining customizable web security rules. You can use AWS WAF to create custom rules that block common attack patterns, such as SQL injection or cross-site scripting, and rules that are designed for your specific application. New rules can be deployed within minutes, letting you respond quickly to changing traffic patterns. Also, AWS WAF includes a full-featured API that you can use to automate the creation, deployment, and maintenance of web security rules.

In WAF, you can create a set of Conditions and Rules to protect your network against attacks from outside.

For more information on AWS WAF please visit the below link

https://aws.amazon.com/waf/

The correct answer is: WAF

Question **35**

Not answered



Your supervisor asks you to create a decoupled application whose process includes dependencies on EC2 instances and servers located in your company's on-premises data center. Which of these are you least likely to recommend as part of that process? Choose the correct answer from the options below:

Please select:

- A. SQS polling from an EC2 instance deployed with an IAM role
- B. An SWF workflow
- C. SQS polling from an EC2 instance using IAM user credentials
- D. SQS polling from an on-premises server using IAM user credentials

Your answer is incorrect.

Answer - C

Note that the question asks you for the least likely recommended option.

You should use IAM roles for secure communication between EC2 instances and resources on AWS.

Your most likely scenario will actually be SQS polling from an EC2 instance deployed with an IAM role because when your polling SQS from EC2 you should use IAM roles. What you should never do is use IAM user api keys for authentication to poll sqs messages.

An IAM *role* is similar to a user, in that it is an AWS identity with permission policies that determine what the identity can and cannot do in AWS. However, instead of being uniquely associated with one person, a role is intended to be assumable by anyone who needs it. Also, a role does not have any credentials (password or access keys) associated with it. Instead, if a user is assigned to a role, access keys are created dynamically and provided to the user.

For more information on IAM Roles, please refer to the below link:

• http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles.html

The correct answer is: SQS polling from an EC2 instance using IAM user credentials

Ouestion 36

Not answered



An EC2 instance retrieves a message from an SQS queue, begins processing the message, then crashes. What happens to the message? Choose the correct answer from the options below:

Please select:

- A. Before the message visibility timeout expires, the message becomes available for processing by other EC2 instances
- B. It will remain in the queue and again assigned to the same EC2 instance when the instance comes online.
- C. The message is deleted and becomes duplicated when the EC2 instance comes online.

Your answer is incorrect.

Answer - B

When a consumer receives and processes a message from a queue, the message remains in the queue. Amazon SQS doesn't automatically delete the message: Because it's a distributed system, there is no guarantee that the component will actually receive the message (the connection can break or a component can fail to receive the message). Thus, the consumer must delete the message from the queue after receiving and processing it.

Q: How does Amazon SQS allow multiple readers to access the same message queue without losing messages or processing them multiple times?

Every Amazon SQS queue has a configurable visibility timeout. A message is not visible to any other reader for a designated amount of time when it is read from a message queue. As long as the amount of time it takes to process the message is less than the visibility timeout, every message is processed and deleted.

If the component processing of the message fails or becomes unavailable, the message again becomes visible to any component reading the message queue once the visibility timeout ends. This allows multiple components to read messages from the same message queue, each one working to process different messages.

For more information on SQS Visibility timeout, please refer to the below link

- http://sqs-publicimages.s3.amazonaws.com/Building_Scalabale_EC2_applications_with_SQS2.pdf (this document explains in detail how EC2 and SQS works together in all scenarios. There is also explanation what happens if the EC2 instance crashes before it deletes a message from Queue)
- http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-visibilitytimeout.html

The correct answer is: It will remain in the queue and again assigned to the same EC2 instance when the instance comes online.

Question **37**

Not answered



You are running an instance store based instance. You shutdown and then start the instance. You then notice that the data which you have saved earlier is no longer available. What might be the cause of this? Choose the correct answer from the options below

Please select:

- A. The volume was not big enough to handle all of the processing data
- B. The EC2 instance was using EBS backed root volumes, which are ephemeral and only live for the life of the instance
- C. The EC2 instance was using instance store volumes, which are ephemeral and only live for the life of the instance
- D. The instance might have been compromised

Your answer is incorrect.

Answer - C

The data in an instance store persists only during the lifetime of its associated instance. If an instance reboots (intentionally or unintentionally), data in the instance store persists. However, data in the instance store is lost under the following circumstances:

- The underlying disk drive fails
- The instance stops
- The instance terminates

For more information on Instance store, please refer to the below link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html

The correct answer is: The EC2 instance was using instance store volumes, which are ephemeral and only live for the life of the instance

Ouestion 38

Not answered



You have been told that you need to set up a bastion host by your manager in the cheapest, most secure way, and that you should be the only person that can access it via SSH. Which of the following setups would satisfy your manager's request? Choose the correct answer from the options below

Please select:

A. A small EC2 instance and a security group which only allows access on port 22 via your IP address

 B. A large EC2 instance and a security group which only allows access on port 22 via your IP address
C. A large EC2 instance and a security group which only allows access on port 22
D. A small EC2 instance and a security group which only allows access on port 22
Your answer is incorrect.
Answer – A
The bastion host should only have a security group from a particular IP address for maximum
security. Since the request is to have a cheapest infrastructure, then you should use a small instance.
The correct answer is: A small EC2 instance and a security group which only allows access on port 22
via your IP address
W. I. S.I. S.II. I. W. I
Which of the following are Invalid VPC peering configurations?
Choose 3 answers from the options below
Please select :
A. Overlapping CIDR blocks
■ B. Transitive Peering
C. Edge to Edge routing via a gateway
D. One to one relationship between 2 VPC's
Your answer is incorrect.
Answer – A,B and C
Allower App and C

This is given in the aws documentation

Question **39**Not answered

For more information on VPC Peering configurations, please refer to the below link http://docs.aws.amazon.com/AmazonVPC/latest/PeeringGuide/invalid-peering-configurations.html The correct answer is: Overlapping CIDR blocks, Transitive Peering, Edge to Edge routing via a gateway

Question 40

Not answered



You've been tasked with building out a duplicate environment in another region for disaster recovery purposes. Part of your environment relies on EC2 instances with preconfigured software. What steps would you take to configure the instances in another region? Choose the correct answer from the options below

Please select:

- A. Create an AMI of the EC2 instance
- B. Create an AMI of the EC2 instance and copy the AMI to the desired region
- C. Make the EC2 instance shareable among other regions through IAM permissions
- D. None of the above

Your answer is incorrect.

Answer - B

You can copy an Amazon Machine Image (AMI) within or across an AWS region using the AWS Management Console, the AWS command line tools or SDKs, or the Amazon EC2 API, all of which support the CopyImage action. You can copy both Amazon EBS-backed AMIs and instance store-backed AMIs. You can copy AMIs with encrypted snapshots and encrypted AMIs.

For more information on copying AMI's, please refer to the below link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html

The correct answer is: Create an AMI of the EC2 instance and copy the AMI to the desired region

Question 41

Not answered



In order to establish a successful site-to-site VPN connection from your on-premise network to the VPC (Virtual Private Cloud), which of the following needs to be configured outside of the VPC? Choose the correct answer from the options below

Please select:

- A. The main route table to route traffic through a NAT instance
- B. A public IP address on the customer gateway for the on-premise network
- C. A dedicated NAT instance in a public subnet
- D. An Elastic IP address to the Virtual Private Gateway

Your answer is incorrect.

Answer - B

On the customer side gateway you need to have a public IP address which can be addressed by the VPN connection.

For more information on VPN connections, please refer to the below link:

• http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html

The correct answer is: A public IP address on the customer gateway for the on-premise network

Question 42

Not answered



You have 5 CloudFormation templates. Each template has been defined for a specific purpose. What determines the cost of using the CloudFormation templates? Choose the correct answer from the options below

Please select:

- A. \$1.10 per template per month
- B. The length of time it takes to build the architecture with CloudFormation
- C. It depends on the region the template is created in
- D. CloudFormation does not have a cost but you are charged for the underlying resources it builds

Your answer is incorrect.

Answer - D

This is given in the aws documentation

For more information on Cloudformation pricing, please refer to the below link

https://aws.amazon.com/cloudformation/pricing/

The correct answer is: CloudFormation does not have a cost but you are charged for the underlying resources it builds

Question 43

Not answered



Does S3 provide read-after-write consistency for new objects? Choose the correct answer from the options below

Please select:

- A. Yes, for all regions
- B. No, not for any region
- C. Yes, but only for certain regions and for new objects
- D. Yes, but only for certain regions, not the us-standard region

Answer - A

This is given in the aws documentation

For more information on S3, please refer to the below link

https://aws.amazon.com/s3/faqs/

The correct answer is: Yes, for all regions

Question 44

Not answered



Your organization has been using a HSM (Hardware Security Module) for secure key storage. It is only used for generating keys for your EC2 instances. Unfortunately, the HSM has been zeroized after someone attempted to log in as the administrator three times using an invalid password. This means that the encryption keys on it have been wiped. You did not have a copy of the keys stored anywhere else. How can you obtain a new copy of the keys that you had stored on HSM? Choose the correct answer from the options below

Please select:

- A. You cannot; the keys are lost if you did not have a copy.
- B. Contact AWS Support; your incident will be routed to the team that supports AWS CloudHSM and a copy of the keys will be sent to you after verification
- C. Restore a snapshot of the HSM
- D. You can still connect via CLI; use the command 'get-client-configuration' and you can get a copy of the keys

Your answer is incorrect.

Answer - A

This is given in the aws documentation

For more information on CloudHSM, please refer to the below link

https://aws.amazon.com/cloudhsm/faqs/

The correct answer is: You cannot; the keys are lost if you did not have a copy.

Ouestion 45

What service from AWS can help manage the budgets for all resources in AWS? Choose one answer from the options below

Please select:

- A. Cost Explorer
- B. Cost Allocation Tags
- C. AWS Budgets
- D. Payment History

Your answer is incorrect.

Answer - C

A budget is a way to plan your usage and your costs (also known as spend data), and to track how close your usage and costs are to exceeding your budgeted amount. Budgets use data from Cost Explorer to provide you with a quick way to see your usage-to-date and current estimated charges from AWS, and to see how much your predicted usage accrues in charges by the end of the month. Budgets also compare the current estimated usage and charges to the amount that you indicated that you want to use or spend, and lets you see how much of your budget has been used. AWS updates your budget status several times a day. Budgets track your unblended costs, subscriptions, and refunds.

You can create budgets for different types of usage and different types of cost. For example, you can create a budget to see how many EC2 hours you have used, or how many GB you have stored in an S3 bucket. You can also create a budget to see how much you are spending on a particular service, or how often you call a particular API operation. Budgets use the same data filters as Cost Explorer.

To create your budget, you can perform the below steps

Step 1) Go to your billing section, go to Budgets and create a new Budget

Step 2) In the next screen, you can then mention the budget amount and what services to link the budget to.

For more information on AWS Budgets please visit the below link

http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/budgets-managing-costs.html

The correct answer is: AWS Budgets

Question 46

Not answered



You are testing an application that uses EC2 instances to poll an SQS queue. At this stage of testing, you have verified that the EC2 instances can retrieve messages from the queue, but your coworkers are complaining about not being able to manually retrieve any messages from the queue from their on-premises workstations. What is the most likely source of this problem? Choose the correct answer from the options below

Please select: A. Your coworkers may not have permission to the SQS queue B. Manual polling of SQS queues is not possible C. SQS queues only accept traffic from within AWS D. They are using short polling

Your answer is incorrect.

Answer - A

If the user has not been granted access to use SQS , then they will not be able to retrieve messages from SQS.

For more information on SQS Access permissions, please refer to the below link

http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-authentication-and-access-control.html

The correct answer is: Your coworkers may not have permission to the SQS queue

Question 47

Not answered



Currently you're helping design and architect a highly available application. After building the initial environment, you've found that part of your application does not work correctly until port 443 is added to the security group. After adding port 443 to the appropriate security group, how much time will it take before the changes are applied and the application begins working correctly? Choose the correct answer from the options below

Please select:

- A. Generally, it takes 2-5 minutes in order for the rules to propagate
- B. Immediately after a reboot of the EC2 instances belong to that security group
- C. Changes apply instantly to the security group, and the application should be able to respond to 443 requests
- D. It will take 60 seconds for the rules to apply to all availability zones within the region

Your answer is incorrect.

Answer - C

This is given in the aws documentation

For more information on Security Groups, please refer to the below link

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC SecurityGroups.html

The correct answer is: Changes apply instantly to the security group, and the application should be able to respond to 443 requests

Ouestion 48

Not answered



Which of the following services allow the administrator access to the underlying operating system?

Choose the 2 correct answers from the options below

Please select:

- A. Amazon RDS
- B. Amazon EMR
- C. Amazon EC2
- D. DynamoDB

Your answer is incorrect.

Answer - B and C

Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides secure, resizable compute capacity in the cloud. It is designed to make web-scale cloud computing easier for developers.

For more information on EC2, please refer to the below link

https://aws.amazon.com/ec2/

Amazon EMR provides a managed Hadoop framework that makes it easy, fast, and cost-effective to process vast amounts of data across dynamically scalable Amazon EC2 instances. You can also run other popular distributed frameworks such as Apache Spark, HBase, Presto, and Flink in Amazon EMR, and interact with data in other AWS data stores such as Amazon S3 and Amazon DynamoDB.

For more information on EMR, please refer to the below link

https://aws.amazon.com/emr/

The correct answer is: Amazon EMR, Amazon EC2

Question **49**

Not answered



Besides regions and their included availability zones, which of the following is another "regional" data center location used for content distribution? Choose the correct answer from the options below

Please select:

- A. Edge Location
- B. Front Location
- C. Backend Location
- D. Cloud Location

Answer - A

Using a network of edge locations around the world, Amazon CloudFront caches copies of your static content close to viewers, lowering latency when they download your objects and giving you the high, sustained data transfer rates needed to deliver large popular objects to end users at scale.

For more information on Cloudfront and edge locations, please refer to the below link

https://aws.amazon.com/cloudfront/

The correct answer is: Edge Location

Question **50**

Not answered



What are the main benefits of IAM groups?

Choose 2 answers from the options below

Please select:

- A. Ability to create custom permission policies.
- B. Allow for EC2 instances to gain access to S3.
- C. Easier user/policy management.
- D. Assign IAM permission policies to more than one user at a time.

Your answer is incorrect.

Answer - C and D

An IAM *group* is a collection of IAM users. Groups let you specify permissions for multiple users, which can make it easier to manage the permissions for those users. For example, you could have a group called *Admins* and give that group the types of permissions that administrators typically need. Any user in that group automatically has the permissions that are assigned to the group. If a new user joins your organization and needs administrator privileges, you can assign the appropriate permissions by adding the user to that group. Similarly, if a person changes jobs in your organization, instead of editing that user's permissions, you can remove him or her from the old groups and add him or her to the appropriate new groups.

For more information on IAM Groups, please refer to the below link

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_groups.html

The correct answer is: Easier user/policy management., Assign IAM permission policies to more than one user at a time.

Ouestion 51

Not answered

API Access Keys are required to make programmatic call to AWS from which of the following?

Choose the 3 correct answers from the options below

Please select :
A. Windows PowerShell
■ B. Managing AWS resources through the AWS console
C. Direct HTTP call using the API
D. AWS CLI

Answer - A, C and D

By default, when you create an access key, its status is Active, which means the user can use the access key for AWS CLI, Tools for Windows PowerShell, and API calls. Each user can have two active access keys, which is useful when you must rotate the user's access keys. You can disable a user's access key, which means it can't be used for API calls. You might do this while you're rotating keys or to revoke API access for a user

For more information on API Access keys, please refer to the below link

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_access-keys.html

The correct answer is: Windows PowerShell, Direct HTTP call using the API, AWS CLI

Question **52**

Not answered



In the default VPC, all subnets have a route to the Internet. Choose the correct answer from the options below

Please select:

A. True

B. False

Your answer is incorrect.

Answer - A

When a default VPC is created, the following elements are created

- Create a VPC with a size /16 IPv4 CIDR block.
- Create a default subnet in each Availability Zone.
- Create an Internet gateway and connect it to your default VPC.
- Create a main route table for your default VPC with a rule that sends all IPv4 traffic destined for the Internet to the Internet gateway.
- Create a default security group and associate it with your default VPC.
- Create a default network access control list (ACL) and associate it with your default VPC.
- Associate the default DHCP options set for your AWS account with your default VPC.

For more information on default VPC, please refer to the below link

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/default-vpc.html The correct answer is: True How are Network access rules evaluated? Choose the correct answer from the options below Please select: A. Rules are evaluated by rule number, from highest to lowest, and executed immediately when a matching allow/deny rule is found. B. All rules are evaluated before any traffic is allowed or denied. C. Rules are evaluated by rule number, from lowest to highest, and executed immediately when a matching allow/deny rule is found. D. Rules are evaluated by rule number, from lowest to highest, and executed after all rules are checked for conflicting allow/deny rules. Your answer is incorrect. Answer - C This is given in the aws documentation For more information on NACL, please refer to the below link http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html The correct answer is: Rules are evaluated by rule number, from lowest to highest, and executed immediately when a matching allow/deny rule is found. A VPC can have up to 5 Internet Gateways attached at any given time. Please select: A. True B. False Your answer is incorrect. Answer - B This is given in the aws documentation

For more information on VPC limits, please refer to the below link

Ouestion **53**

Not answered

Question **54**

Not answered

http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Appendix_Limits.html The correct answer is: False What are three attributes of DynamoDB? Choose the 3 correct answers from the options below Please select: A. Used for data warehousing B. A NoSQL database platform C. Uses key-value store D. Fully-managed Your answer is incorrect. Answer - B,C and D Amazon DynamoDB is a fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale. It is a fully managed cloud database and supports both document and key-value store models. Its flexible data model and reliable performance make it a great fit for mobile, web, gaming, ad tech, IoT, and many other applications. AWS Redshift can be used for data warehousing For more information on DynamoDB, please refer to the below link https://aws.amazon.com/dynamodb/ The correct answer is: A NoSQL database platform, Uses key-value store, Fully-managed If you cannot connect to your Ec2 instance via remote desktop, and you have already verified the instance has a public IP and the Internet gateway and route tables are in place, what should you check next? Choose one answer from the options given below Please select: A. Adjust the security group to allow traffic from port 22 B. Adjust the security group to allow traffic from port 3389

C. Restart the instance since there might be some issue with the instance

D. Create a new instance since there might be some issue with the instance

Your answer is incorrect.

Ouestion **55**

Not answered

Question **56**

Not answered

Answer - B

The reason why you cannot connect to the instance is because by default RDP protocol will not be enabled on the Security Group.

Option A is wrong because this is for the SSH protocol and here we want to RDP into the instance. Option C and D are wrong because there is no mention of anything wrong with the instance.

Step 1) Go to your EC2 Security groups, click on the required security groups to make the changes. Go to the Inbound Tab.

Step 2) Make sure to add a rule for the RDP protocol for the instance and then click the Save button.

The correct answer is: Adjust the security group to allow traffic from port 3389

Question **57**

Not answered



What database service should you choose if you need petabyte-scale data warehousing? Choose the correct answer from the options below

Please select:

- A. DynamoDB
- B. ElastiCache
- C. RDS
- D. Redshift

Your answer is incorrect.

Answer - D

Amazon Redshift is a fast, fully managed data warehouse that makes it simple and cost-effective to analyze all your data using standard SQL and your existing Business Intelligence (BI) tools. It allows you to run complex analytic queries against petabytes of structured data, using sophisticated query optimization, columnar storage on high-performance local disks, and massively parallel query execution

For more information on Redshift, please refer to the below link

http://docs.aws.amazon.com/redshift/latest/mgmt/welcome.html

The correct answer is: Redshift

Question 58 Not answered	Which feature in AWS allows 2 VPC's to talk to each other? Choose one answer from the options given below						
	Please select: A. VPC Connection B. VPN Connection C. Direct Connect D. VPC Peering						
	Your answer is incorrect. Answer - D A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IP addresses. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, or with a VPC in another AWS account within a single region The below diagram shows an example of VPC peering. Now please note that VPC B cannot communicate to VPC C because there is no peering between them.						
	For more information on VPC peering, please visit the url http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-peering.html The correct answer is: VPC Peering						
Question 59 Not answered	In AWS Security Groups what are the 2 types of rules you can define? Select 2 options. Please select: A. Inbound B. Transitional C. Bi-Directional D. Outbound						
	Your answer is incorrect. Answer – A and D						

A security group acts as a virtual firewall that controls the traffic for one or more instances. When you launch an instance, you associate one or more security groups with the instance. You add rules to each security group that allow traffic to or from its associated instances. You can modify the rules for a security group at any time; the new rules are automatically applied to all instances that are associated with the security group.

The below diagram's show that rules can be defined for Inbound and Outbound

For more information on Security Groups, please visit the url

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html

The correct answer is: Inbound, Outbound

Question **60**

Not answered



How much data from Amazon Glacier can you download for free? Choose the right answer from the options given below

Please select:

- A. 5GB
- B. 10GB
- C. 15GB
- D. There is no limit

Your answer is incorrect.

Answer - B

The documentation for aws clearly mentions the retrieval policy for Amazon Glacier

For more information on S3 and Glacier, please visit the url

https://aws.amazon.com/s3/faqs/

The correct answer is: 10GB

QUIZ **NAVIGATION** 20 21 23 24 26 27 38 39 41 42 44 45 47 48 50 51 53 54 56 57 58 59 60 Show one page at a time Finish review