


Whizlabs Learning Center

HOME > AWS CSAA PRACTICE TESTS > FULL TEST(S) > NEW PRACTICE TEST III

Started on Wednesday, 5 July 2017, 3:28 AM**State** Finished**Completed on** Wednesday, 5 July 2017, 3:28 AM**Time taken** 14 secs**Grade** 0.00 out of 60.00 (0%)**Result** FAIL**Question 1**

Not answered

 Mark for Review

A VPC has been setup with public subnet and an internet gateway. You setup an EC2 instance with a public IP. But you are still not able to connect to it via the Internet. You can see that the right Security groups are in place. What should you do to ensure you can connect to the EC2 instance from the internet?

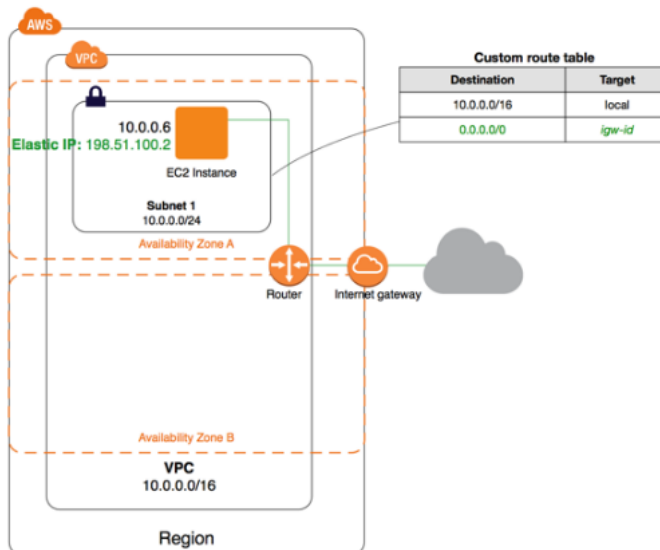
Please select :

- ☐ A. Set an Elastic IP Address to the EC2 instance
- ☐ B. Set a Secondary Private IP Address to the EC2 instance
- ☐ C. Ensure the right route entry is there in the Route table
- ☐ D. There must be some issue in the EC2 instance. Check the system logs.

Your answer is incorrect.

Answer – C

You have to ensure that the Route table has an entry to the Internet gateway because this is required for instances to communicate over the internet. The diagram shows the configuration of the public subnet in a VPC.



Option A is wrong because you already have a public IP Assigned to the instance, so this should be enough to connect to the Internet.

Option B is wrong because private IP's cannot be accessed from the internet

Option D is wrong because the Route table is what is causing the issue and not the system

For more information on AWS public subnet, please visit the link


http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html

The correct answer is: Ensure the right route entry is there in the Route table

Feedback about this question and answer

Question 2

Not answered

 Mark for Review

Which of the following approaches provides the lowest cost for Amazon Elastic Block Store snapshots while giving you the ability to fully restore data?

Please select :

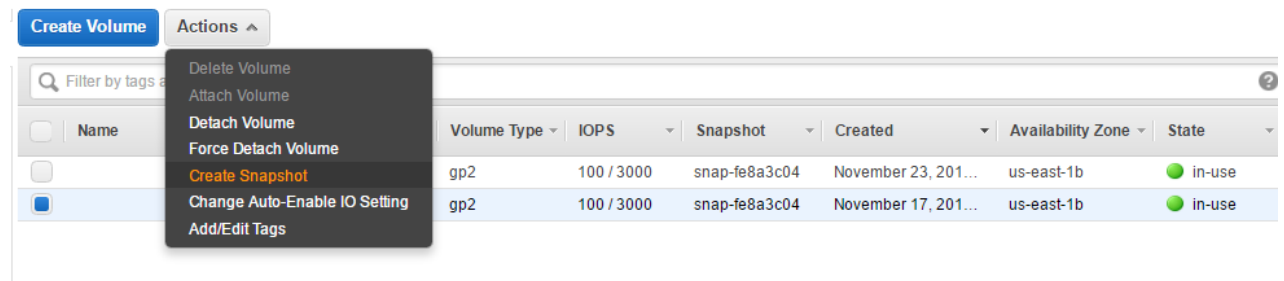
- ☐ A. Maintain two snapshots: the original snapshot and the latest incremental snapshot.
- ☐ B. Maintain a volume snapshot; subsequent snapshots will overwrite one another
- ☐ C. Maintain a single snapshot the latest snapshot is both Incremental and complete.
- ☐ D. Maintain the most current snapshot, archive the original and incremental to Amazon Glacier.

Your answer is incorrect.

Answer – C

EBS snapshots are incremental and complete, so you don't need to maintain multiple snapshots if you are looking on reducing costs.

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.




For more information on EBS snapshots, please visit the link - <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

The correct answer is: Maintain a single snapshot the latest snapshot is both Incremental and complete.

Feedback about this question and answer

Question 3

Not answered

 Mark for Review

An existing application stores sensitive information on a non-boot Amazon EBS data volume attached to an Amazon Elastic Compute Cloud instance. Which of the following approaches would protect the sensitive data on an Amazon EBS volume?

Please select :

- ☐ A. Upload your customer keys to AWS CloudHSM. Associate the Amazon EBS volume with AWS CloudHSM. Remount the Amazon EBS volume.
- ☐ B. Create and mount a new, encrypted Amazon EBS volume. Move the data to the new volume. Delete the old Amazon EBS volume.
- ☐ C. Unmount the EBS volume. Toggle the encryption attribute to True. Re-mount the Amazon EBS volume.
- ☐ D. Snapshot the current Amazon EBS volume. Restore the snapshot to a new, encrypted Amazon EBS volume. Mount the Amazon EBS volume

Your answer is incorrect.

Answer – B

Here the only option available is to create a new mount volume

Option A is wrong because you cannot encrypt a volume once it is created. You would need to use some local encrypting algorithm if you want to encrypt the data on the volume.

Option C is wrong because even if you unmounts the volume, you cannot encrypt the volume. Encryption has to be done during volume creation.


Option D is wrong because even if the volume is not encrypted, the snapshot will also not be encrypted.

The correct answer is: Create and mount a new, encrypted Amazon EBS volume. Move the data to the new volume. Delete the old Amazon EBS volume.

Feedback about this question and answer

Question 4

Not answered

 Mark for Review

Which of the following are use cases for Amazon DynamoDB? Choose 3 answers

Please select :

- ☐ A. Storing BLOB data.
- ☐ B. Managing web sessions.

- ☐ C. Storing JSON documents.
- ☐ D. Storing metadata for Amazon S3 objects.
- ☐ E. Running relational joins and complex updates.
- ☐ F. Storing large amounts of infrequently accessed data

Your answer is incorrect.

Answer - B, C and D

Amazon DynamoDB stores structured data, indexed by primary key, and allows low latency read and write access to items ranging from 1 byte up to 400KB.

Amazon S3 stores unstructured blobs and suited for storing large objects up to 5 TB.

DynamoDB IS a good choice to store the metadata for a BLOB, such as name, date created, owner, etc... The Binary Large Object itself would be stored in S3.

Q: When should I use Amazon DynamoDB vs Amazon S3?

Amazon DynamoDB stores structured data, indexed by primary key, and allows low latency read and write access to items ranging from 1 byte up to 400KB. Amazon S3 stores unstructured blobs and suited for storing large objects up to 5 TB. In order to optimize your costs across AWS services, large objects or infrequently accessed data sets should be stored in Amazon S3, while smaller data elements or file pointers (possibly to Amazon S3 objects) are best saved in Amazon DynamoDB.

For more information on Amazon Dynamo DB, please visit


- <https://aws.amazon.com/dynamodb/faqs/>

The correct answer is: Managing web sessions., Storing JSON documents., Storing metadata for Amazon S3 objects.

Feedback about this question and answer

Question 5

Not answered

 Mark for Review

In Amazon CloudWatch what is the retention period for a one minute datapoint. Choose the right answer from the options given below

Please select :

- ☐ a. 10 days
- ☐ b. 15 days
- ☐ c. 1 month
- ☐ d. 1 year

Your answer is incorrect.

Answer - B

Amazon CloudWatch is a monitoring service for AWS cloud resources and the applications you run on AWS. You can use Amazon CloudWatch to collect and track metrics, collect and monitor log files, set alarms, and automatically react to changes in your AWS resources. Amazon CloudWatch can monitor AWS resources such as Amazon EC2 instances, Amazon DynamoDB tables, and Amazon RDS DB instances, as well as custom metrics generated by your applications and services, and any log files your applications generate. You can use Amazon CloudWatch to gain system-wide visibility into resource utilization, application performance, and operational health.

Below is the retention period for the various data points

CloudWatch Metrics now supports the following three retention schedules:

- 1 minute datapoints are available for 15 days
- 5 minute datapoints are available for 63 days
- 1 hour datapoints are available for 455 days

For more information on Amazon Cloudwatch, please visit


<https://aws.amazon.com/cloudwatch/>

The correct answer is: 15 days

Feedback about this question and answer

Question 6

Not answered

 Mark for Review

What is the minimum time Interval for the data that Amazon CloudWatch receives and aggregates for EC2 in basic monitoring?

Please select :

- ☐ A. One second
- ☐ B. Five seconds
- ☐ C. One minute
- ☐ D. Three minutes
- ☐ E. Five minutes

Your answer is incorrect.

Answer – E

In Amazon CloudWatch for basic monitoring of EC2 instances, the following important metrics are collected at five minute intervals and stored for two weeks.

- CPU load
- disk I/O
- network I/O

For more information on Amazon Cloudwatch EC2 basic monitoring, please visit


<https://aws.amazon.com/blogs/aws/amazon-cloudwatch-basic-monitoring-for-ec2-at-no-charge/>

The correct answer is: Five minutes

Feedback about this question and answer

Question 7

Not answered

 Mark for Review

What log formats does Cloudwatch logs support? Choose 2 answers from the options given below

Please select :

- ☐ A. XML
- ☐ B. JSON Formatted
- ☐ C. SOAP
- ☐ D. Any text based common log

Your answer is incorrect.

Answer – B and D

As per the AWS Documentation for Cloudwatch it is clearly mentioned what log files are supported by Cloudwatch.

Q: What log formats does CloudWatch Logs support?

CloudWatch Logs can ingest, aggregate and monitor any text based common log data or JSON-formatted logs.

For more information on Amazon Cloudwatch FAQ, please visit


<https://aws.amazon.com/cloudwatch/faqs/>

The correct answer is: JSON Formatted, Any text based common log

Feedback about this question and answer

Question 8

Not answered

 Mark for Review

A customer wants to apply a group of database specific settings to their rds instances in their aws account. Which of the below options can be used to apply the settings in one go to all rds instances

Please select :

- ☐ A. Security Groups
- ☐ B. NACL Groups
- ☐ C. Parameter Groups
- ☐ D. IAM Roles.

Your answer is incorrect.

Answer – C

DB Parameter Groups are used to assign specific settings which can be applied to a set of RDS instances in aws.

In your RDS, when you go to Parameter Groups, you can create a new parameter group. In the parameter group itself, you have a lot of database related settings that can be assigned to the database.

RDS Dashboard

Instances

Clusters

Reserved Purchases

Snapshots

Security Groups

Parameter Groups

External Licenses

Option Groups

Subnet Groups

Events

Event Subscriptions

Notifications

Parameter Groups > demo

Parameters

Recent Events

Tags

Filter:

Cancel Editing

Preview Changes

Save Changes

Name	Edit Values	Allowed Values
allow-suspicious-udfs	<input type="text"/>	0, 1
auto_increment_increment	<input type="text"/>	1-65535
auto_increment_offset	<input type="text"/>	1-65535
autocommit	<engine-default>	
automatic_sp_privileges	<engine-default>	
back_log	<input type="text"/>	1-65535
basedir	/rdsdbbin/mysql	
binlog_cache_size	32768	4096-18446744073709547520

Option A, B and D are wrong because this is specific to what resources have access to the database.

For more information on EB parameter groups, please visit

http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithParamGroups.html

The correct answer is: Parameter Groups

Feedback about this question and answer

Question 9

Not answered

Mark for Review

When creating a MySQL instance on AWS, what are the different storage options available?

Select 3 options.

Please select :

- ☐ A. General Purpose (SSD)
- ☐ B. Provisioned IOPS
- ☐ C. RAID 5
- ☐ D. Magnetic

Your answer is incorrect.

Answer – A, B and D

When you create an RDS in AWS, you will see the below Storage options which has the options for General Purpose (SSD), Provisioned IOPS & Magnetic.

Instance Specifications

DB Engine

mysql

License Model

general-public-license

DB Engine Version

5.6.27

Review the Known Issues/Limitations to learn about potential compatibility issues with specific database versions.

DB Instance Class

db.t2.micro — 1 vCPU, 1 GiB RAM

Multi-AZ Deployment

- Select One -

Storage Type

General Purpose (SSD)

Allocated Storage*

- Select One -

General Purpose (SSD)

Provisioned IOPS (SSD)

Magnetic

For more information on AWS RDS, please visit

<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Welcome.html>

The correct answer is: General Purpose (SSD), Provisioned IOPS, Magnetic

Feedback about this question and answer

Question 10

Not answered

Mark for Review

Before I delete an EBS volume, what can I do if I want to recreate the volume later?

Please select :

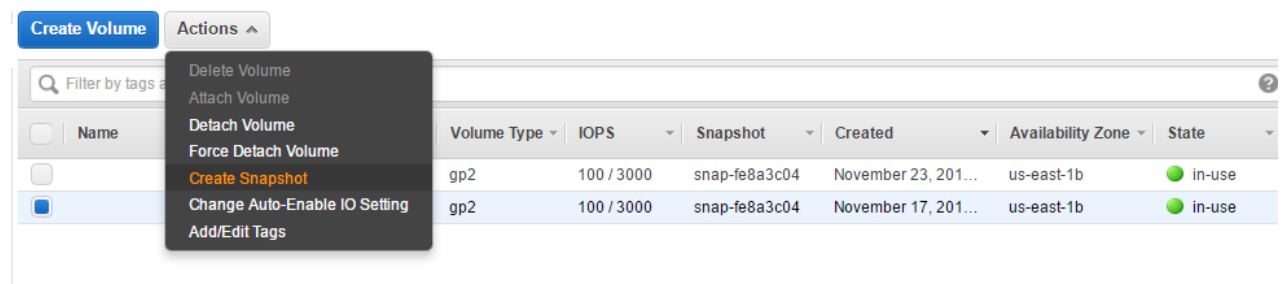
- ☐ A. Create a copy of the EBS volume (not a snapshot)
- ☐ B. Store a snapshot of the volume
- ☐ C. Download the content to an EC2 instance
- ☐ D. Back up the data in to a physical disk

Your answer is incorrect.

Answer – B

Snapshots occur asynchronously; the point-in-time snapshot is created immediately, but the status of the snapshot is pending until the snapshot is complete (when all of the modified blocks have been transferred to Amazon S3), which can take several hours for large initial snapshots or subsequent snapshots where many blocks have changed. While it is completing, an in-progress snapshot is not affected by ongoing reads and writes to the volume.

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.



For more information on EBS snapshots, please visit the link - <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

The correct answer is: Store a snapshot of the volume

Feedback about this question and answer

Question 11

Not answered

Mark for Review

All Amazon EC2 instances are assigned two IP addresses at launch, out of which one can only be reached from within the Amazon EC2 network?

Please select :

- ☐ A. Multiple IP address
- ☐ B. Public IP address
- ☐ C. Private IP address
- ☐ D. Elastic IP Address

Your answer is incorrect.

Answer – C

A private IP address is an IP address that's not reachable over the Internet. You can use private IP addresses for communication between instances in the same network (EC2-Classic or a VPC).

When an instance is launched a private IP address is allocated for the instance using DHCP. Each instance is also given an internal DNS hostname that resolves to the private IP address of the instance; for example, ip-10-251-50-12.ec2.internal. You can use the internal DNS hostname for communication between instances in the same network, but we can't resolve the DNS hostname outside the network that the instance is in.

For more information on IP Addressing, please visit the link –


<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-instance-addressing.html>

The correct answer is: Private IP address

Feedback about this question and answer

Question 12

Not answered

 Mark for Review

The below snapshot shows 2 instances in the Ec2 dashboard. What is the column name for the instances as highlighted in red below?

<input type="radio"/>	i-023260ed8c2050...	t2.micro	us-east-1b
<input checked="" type="radio"/>	i-09f54a79fc36966e4	t2.micro	us-east-1b

Please select :

- ☐ A. Instance ID
- ☐ B. VPC ID
- ☐ C. Subnet ID
- ☐ D. Public IP

Your answer is incorrect.

Answer - A

When resources are created, an assignment of each resource with a unique resource ID is done. You can use resource IDs to find your resources in the Amazon EC2 console. A resource ID takes the form of a resource identifier (such as snap for a snapshot) followed by a hyphen and a unique combination of letters and numbers

<input type="radio"/>	Name	Instance ID	Instance Type	Availability Zone	Instance State
<input type="radio"/>		i-023260ed8c2050...	t2.micro	us-east-1b	running
<input checked="" type="radio"/>		i-09f54a79fc36966e4	t2.micro	us-east-1b	running

For more information on Instance ID, please visit the link


<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/resource-ids.html>

The correct answer is: Instance ID

Feedback about this question and answer

Question 13

Not answered

 Mark for Review

While creating an Amazon RDS DB, what is the first task you should carry out to ensure that the right controls are in place to ensure the proper IP addresses or EC2 instances have access to your DB Instance.

Please select :

- ☐ A. Security Pool
- ☐ B. Secure Zone
- ☐ C. Security Token Pool
- ☐ D. Security Group

Your answer is incorrect.

Answer – D

A DB security group controls network access to a DB instance that is not inside a VPC. By default, network access is turned off to a DB instance. You can specify rules in a security group that allows access from an IP address range, port, or EC2 security group. Once ingress rules are configured, the same rules apply to all DB instances that are associated with that security group. You can specify up to 20 rules in a security group

When you go to the Configure Advanced Settings tab for the DB creation, you will see the option to specify the security group.

Step 1: [Select Engine](#)
Step 2: [Production?](#)
Step 3: [Specify DB Details](#)
Step 4: Configure Advanced Settings

Configure Advanced Settings

Network & Security

VPC* Default VPC (vpc-6dcc550a)

Subnet Group default

Publicly Accessible Yes

Availability Zone No Preference

VPC Security Group(s)
Create new Security Group
AutoScaling-Security-Group-1 (VPC)
Demo (VPC)
default (VPC)

For more information on DB security groups visit the below link

http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithSecurityGroups.html

The correct answer is: Security Group

Feedback about this question and answer

Question 14

Not answered

Mark for Review

Your AWS Engineer is trying to create a database in aws and is trying to have it encrypted. But when he reaches Database options screen, the Enable Encryption option is greyed out.

Database Options

Database Name

Note: if no database name is specified then no initial MySQL database will be created on the DB Instance.

Database Port 3306

DB Parameter Group default.mysql5.6

Option Group default.mysql-5-6

Copy Tags To Snapshots ☐

Enable Encryption No

As a solution Architect what is the solution you would give the engineer?

Please select :

- ☐ A. Tell that it's not possible , because AWS RDS does not support Encryption
- ☐ B. Since the storage option was chosen as SSD in the previous screen, that's why the option is not enabled.
- ☐ C. Since the instance option was not chosen as db.t2.large in the previous screen, that's why the option is not enabled. This would have been selected as db.t2.medium which causes encryption option is greyed out.
- ☐ D. If the selected database type is Oracle, only then Encryption can be enabled.

Your answer is incorrect.

Answer – C

The DB instance class has to be at least db.t2.large for Encryption to be enabled. Your aws engineer has to select the db.t2.large instance class as shown below:

DB Instance Class db.t2.large — 2 vCPU, 8 GiB RAM

Multi-AZ Deployment - Select One -

Storage Type General Purpose (SSD)

Allocated Storage* 5 GB

Amazon RDS encrypted instances use the industry standard AES-256 encryption algorithm to encrypt your data on the server that hosts your Amazon RDS instance. Once your data is encrypted, Amazon RDS handles authentication of access and decryption of your data transparently with a minimal impact on performance. You don't need to modify your database client applications to use encryption.

Availability of Amazon RDS Encrypted Instances

Amazon RDS encrypted instances are currently available for all database engines and storage types. Amazon RDS encryption is not currently available in the China (Beijing) region.

Amazon RDS encryption is available for all DB instance classes for Amazon Aurora, and for the following DB instance classes for all other Amazon RDS database engines:

Instance Type	Instance Class
General Purpose (M4)—Current Generation	db.m4.large
	db.m4.xlarge
	db.m4.2xlarge
	db.m4.4xlarge
	db.m4.10xlarge
Memory Optimized (R3)—Current Generation	db.r3.large
	db.r3.xlarge
	db.r3.2xlarge
	db.r3.4xlarge
	db.r3.8xlarge
Burst Capable (T2)—Current Generation	db.t2.large

For more information on Database Encryption visit the below link


<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Overview.Encryption.html>

The correct answer is: Since the instance option was not chosen as db.t2.large in the previous screen, that's why the option is not enabled. This would have been selected as db.t2.medium which causes encryption option is greyed out.

Feedback about this question and answer

Question 15

Not answered

 Mark for Review

Where does AWS beanstalk store the application files and server log files? Choose one answer from the options given below

Please select :

- ☐ A. On the local server within Elastic beanstalk
- ☐ B. AWS S3
- ☐ C. AWS Cloudtrail
- ☐ D. AWS DynamoDB

Your answer is incorrect.

Answer - B

AWS Elastic Beanstalk stores your application files and, optionally, server log files in Amazon S3. If you are using the AWS Management Console, the AWS Toolkit for Visual Studio, or AWS Toolkit for Eclipse, an Amazon S3 bucket will be created in your account for you and the files you upload will be automatically copied from your local client to Amazon S3. Optionally, you may configure Elastic Beanstalk to copy your server log files every hour to Amazon S3. You do this by editing the environment configuration settings

For more information on Elastic Beanstalk visit the below link


<https://aws.amazon.com/elasticbeanstalk/faqs/>

The correct answer is: AWS S3

Feedback about this question and answer

Question 16

Not answered

 Mark for Review

What does the AWS Storage Gateway provide?

Please select :

- ☐ A. It allows to integrate on-premises IT environments with Cloud Storage.

- ☐ B. A direct encrypted connection to Amazon S3.
- ☐ C. It's a backup solution that provides an on-premises Cloud storage.
- ☐ D. It provides an encrypted SSL endpoint for backups in the Cloud.

Your answer is incorrect.

Answer – A

Option B is wrong because it is not an encrypted solution to S3

Option C is wrong because you can use S3 as a backup solution

Option D is wrong because the SSL endpoint can be achieved via S3

The AWS Storage Gateway's software appliance is available for download as a virtual machine (VM) image that you install on a host in your datacenter. Once you've installed your gateway and associated it with your AWS Account through our activation process, you can use the AWS Management Console to create either gateway-cached volumes, gateway-stored volumes, or a gateway-virtual tape library (VTL), which can be mounted as iSCSI devices by your on-premises applications.

You have primarily 2 types of volumes

- 1) Gateway-cached volumes allow you to utilize Amazon S3 for your primary data, while retaining some portion of it locally in a cache for frequently accessed data.
- 2) Gateway-stored volumes store your primary data locally, while asynchronously backing up that data to AWS

For more information on AWS Storage gateways visit the below link


<https://aws.amazon.com/storagegateway/details/>

The correct answer is: It allows to integrate on-premises IT environments with Cloud Storage.

Feedback about this question and answer

Question 17

Not answered

 Mark for Review

What is the base URI for all requests for instance metadata? Choose one answer from the options given below

Please select :

- ☐ A. <http://254.169.169.254/latest/>
- ☐ B. <http://169.169.254.254/latest/>
- ☐ C. <http://127.0.0.1/latest/>
- ☐ D. <http://169.254.169.254/latest/>

Your answer is incorrect.

Answer – D

Instance metadata is data about your instance that you can use to configure or manage the running instance

Because your instance metadata is available from your running instance, you do not need to use the Amazon EC2 console or the AWS CLI. This can be helpful when you're writing scripts to run from your instance. For example, you can access the local IP address of your instance from instance metadata to manage a connection to an external application.

<http://169.254.169.254/latest/meta-data/>

For more information on Instance Metadata visit the below link


<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html>

The correct answer is: <http://169.254.169.254/latest/>

Feedback about this question and answer

Question 18

Not answered

 Mark for Review

When you disable automated backups for aws rds, what are you compromising on? Choose on answer from the options given below

Please select :

- ☐ A. Nothing, you are actually saving resources on aws
- ☐ B. You are kind of disable the point-in-time recovery.
- ☐ C. Nothing really, you can take automated backups.
- ☐ D. You cannot disable automated backups in RDS.

Your answer is incorrect.

Answer – B

Amazon RDS creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases. You can set the backup retention period when you create a DB instance. If you don't set the backup retention period, Amazon RDS uses a default period retention period of one day. You can modify the backup retention period; valid values are 0 (for no backup retention) to a maximum of 35 days

You will also specifically see AWS mentioning the risk of not allowing automated backups.

Important

We highly discourage disabling automated backups because it disables point-in-time recovery.

If you disable and then re-enable automated backups, you are only able to restore starting from the time you re-enabled automated backups.

For more information on Automated backups, please visit


http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html

The correct answer is: You are kind of disable the point-in-time recovery.

Feedback about this question and answer

Question 19

Not answered

 Mark for Review

For what databases are read replica's possible in AWS?

Select 3 options.

Please select :

- ☐ A. Oracle
- ☐ B. MySQL
- ☐ C. MariaDB
- ☐ D. Aurora

Your answer is incorrect.

Answer – B, C and D

Amazon RDS Read Replicas provide enhanced performance and durability for Database (DB) Instances. This replication feature makes it easy to elastically scale out beyond the capacity constraints of a single DB Instance for read-heavy database workloads. You can create one or more replicas of a given source DB Instance and serve high-volume application read traffic from multiple copies of your data, thereby increasing aggregate read throughput. Read replicas can also be promoted, so that they become standalone DB Instances.

For more information on read replicas, please visit


<https://aws.amazon.com/rds/details/read-replicas/>

The correct answer is: MySQL, MariaDB, Aurora

Feedback about this question and answer

Question 20

Not answered

 Mark for Review

In what events would cause Amazon RDS to initiate a failover to the standby replica?

Select 3 options.

Please select :

- ☐ A. Loss of availability in primary Availability Zone
- ☐ B. Loss of network connectivity to primary
- ☐ C. Storage failure on secondary
- ☐ D. Storage failure on primary

Your answer is incorrect.

Answer - A, B and D

Amazon RDS detects and automatically recovers from the most common failure scenarios for Multi-AZ deployments so that you can resume database operations as quickly as possible without administrative intervention. Amazon RDS automatically performs a failover in the event of any of the following:

- Loss of availability in primary Availability Zone
- Loss of network connectivity to primary
- Compute unit failure on primary

- Storage failure on primary

Note: When operations such as DB Instance scaling or system upgrades like OS patching are initiated for Multi-AZ deployments, for enhanced availability, they are applied first on the standby prior to an automatic failover. As a result, your availability impact is limited only to the time required for automatic failover to complete. Note that Amazon RDS Multi-AZ deployments do not failover automatically in response to database operations such as long running queries, deadlocks or database corruption errors.

For more information on read replicas, please visit

<https://aws.amazon.com/rds/details/read-replicas/>

The correct answer is: Loss of availability in primary Availability Zone, Loss of network connectivity to primary, Storage failure on primary

Feedback about this question and answer

Question **21**

Not answered

Mark for Review

What does the following command do with respect to the Amazon EC2 security groups?

revoke-security-group-ingress

Please select :

- ☐ A. Removes one or more security groups from a rule.
- ☐ B. Removes one or more security groups from an Amazon EC2 instance.
- ☐ C. Removes one or more rules from a security group.

Your answer is incorrect.

Answer – C

Removes one or more ingress rules from a security group. The values that you specify in the revoke request (for example, ports) must match the existing rule's values for the rule to be removed.

Each rule consists of the protocol and the CIDR range or source security group. For the TCP and UDP protocols, you must also specify the destination port or range of ports. For the ICMP protocol, you must also specify the ICMP type and code.

For more information on revoke-security-group-ingress CLI command, please visit

<http://docs.aws.amazon.com/cli/latest/reference/ec2/revoke-security-group-ingress.html>

The correct answer is: Removes one or more rules from a security group.

Feedback about this question and answer

Question **22**

Not answered

Mark for Review

What is the durability of S3 RRS?

Please select :

- ☐ A. 99.99%
- ☐ B. 99.95%
- ☐ C. 99.995%
- ☐ D. 99.99999999%

Your answer is incorrect.

Answer – A

RRS only has 99.99% durability and there is a chance that data can be lost. So you need to ensure you have the right steps in place to replace lost objects.

	Standard	Standard - Infrequent Access	Reduced Redundancy Storage
Durability	99.999999999%	99.999999999%	99.99%

For more information on RRS, visit the link

<https://aws.amazon.com/s3/reduced-redundancy/>

The correct answer is: 99.99%

Feedback about this question and answer

Question 23

Not answered

Mark for Review

Which aws service is used to as a global content delivery network (CDN) service in aws?

Please select :

- ☐ A. Amazon SES
- ☐ B. Amazon Cloudtrail
- ☐ C. Amazon CloudFront
- ☐ D. Amazon S3

Your answer is incorrect.

Answer – C

Amazon CloudFront is a web service that gives businesses and web application developers an easy and cost effective way to distribute content with low latency and high data transfer speeds. Like other AWS services, Amazon CloudFront is a self-service, pay-per-use offering, requiring no long term commitments or minimum fees. With CloudFront, your files are delivered to end-users using a global network of edge locations.

For more information on CloudFront, please visit the link

<https://aws.amazon.com/cloudfront/>

The correct answer is: Amazon CloudFront

Feedback about this question and answer

Question 24

Not answered

Mark for Review

What features in aws acts as a firewall that controls the traffic allowed to reach one or more instances.

Please select :

- ☐ A. Security group
- ☐ B. ACL
- ☐ C. IAM
- ☐ D. Private IP Addresses

Your answer is incorrect.

Answer - A

A security group acts as a virtual firewall that controls the traffic for one or more instances. When you launch an instance, you associate one or more security groups with the instance. You add rules to each security group that allow traffic to or from its associated instances. Below is an example of a security group for EC2 instances that allows inbound rules and ensure there is a rule for TCP on port 22.

EC2 Dashboard

- Events
- Tags
- Reports
- Limits
- INSTANCES
 - Instances
 - Spot Requests
 - Reserved Instances
 - Scheduled Instances
 - Dedicated Hosts
- IMAGES
 - AMIs
 - Bundle Tasks
- ELASTIC BLOCK STORE
 - Volumes
 - Snapshots
- NETWORK & SECURITY
 - Security Groups**

Create Security Group Actions

search: Demo Add filter

Name	Group ID	Group Name	VPC ID	Description
	sg-2831a055	Demo	vpc-6dcc550a	Demo

Security Group: sg-2831a055

Description Inbound Outbound Tags

Edit

Type	Protocol	Port Range	Source
SSH	TCP	22	0.0.0.0/0


For more information on EC2 Security groups, please visit the url - <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html>

The correct answer is: Security group

Feedback about this question and answer

Question **25**

Not answered

 Mark for Review

How many types of block devices does Amazon EC2 support? Choose one answer from the options below

Please select :

- ☐ A. 2
- ☐ B. 3
- ☐ C. 4
- ☐ D. 1

Your answer is incorrect.

Answer – A

A block device is a storage device that moves data in sequences of bytes or bits (blocks). These devices support random access and generally use buffered I/O. Examples include hard disks, CD-ROM drives, and flash drives. A block device can be physically attached to a computer or accessed remotely as if it were physically attached to the computer. Amazon EC2 supports two types of block devices:

- Instance store volumes (virtual devices whose underlying hardware is physically attached to the host computer for the instance)


EBS volumes (remote storage devices)

The correct answer is: 2

Feedback about this question and answer

Question **26**

Not answered

 Mark for Review

By default, when an EBS volume is attached to a Windows instance, it may show up as any drive letter on the instance. For which services can you use to change the settings of the drive letters of the EBS volumes per your specifications?

Please select :

- ☐ A. EBSConfig Service
- ☐ B. AMIConfig Service
- ☐ C. Ec2Config Service
- ☐ D. Ec2-AMIConfig Service

Your answer is incorrect.

Answer – C

Windows AMIs include an optional service called the EC2Config service (EC2Config.exe). EC2Config starts when the instance boots and performs tasks during startup and each time you stop or start the instance. EC2Config can also perform tasks on demand. Some of these tasks are automatically enabled, while others must be enabled manually. Although optional, this service provides access to advanced features that aren't otherwise available. This service runs in the LocalSystem account.

For more information on Ec2 Config service, please visit the link


http://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/UsingConfig_WinAMI.html

The correct answer is: Ec2Config Service

Feedback about this question and answer

Question **27**

Not answered

 Mark for Review

When running my DB Instance as a Multi-AZ deployment, can I use the standby for read and write operations?

Please select :

- ☐ A. Yes
- ☐ B. Only with MSSQL based RDS
- ☐ C. Only for Oracle RDS instances
- ☐ D. No

Your answer is incorrect.

Answer – D

This is clearly mentioned in the aws documentation that you cannot use the secondary DB instances for writing purposes.

Q: When running my DB Instance as a Multi-AZ deployment, can I use the standby for read or write operations?

No, the standby replica cannot serve read requests. Multi-AZ deployments are designed to provide enhanced database availability and durability, rather than read scaling benefits. As such, the feature uses synchronous replication between primary and standby. Our implementation makes sure the primary and the standby are constantly in sync, but precludes using the standby for read or write operations. If you are interested in a read scaling solution, please see the FAQs on [Read Replicas](#).

Here is the overview of Multi-AZ RDS Deployments:

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure, and is engineered to be highly reliable. In case of an infrastructure failure, Amazon RDS performs an automatic failover to the standby (or to a read replica in the case of Amazon Aurora), so that you can resume database operations as soon as the failover is complete. Since the endpoint for your DB Instance remains the same after a failover, your application can resume database operation without the need for manual administrative intervention.

For more information on Multi AZ RDS, please visit the link


<https://aws.amazon.com/rds/details/multi-az/>

The correct answer is: No

Feedback about this question and answer

Question 28

Not answered

 Mark for Review

Which Amazon service can I use to define a virtual network that closely resembles a traditional data center?

Please select :

- ☐ A. Amazon VPC
- ☐ B. Amazon ServiceBus
- ☐ C. Amazon EMR
- ☐ D. Amazon RDS

Your answer is incorrect.

Answer – A

Amazon Virtual Private Cloud (Amazon VPC) lets you provision a logically isolated section of the Amazon Web Services (AWS) cloud where you can launch AWS resources in a virtual network that you define. You have complete control over your virtual networking environment, including selection of your own IP address range, creation of subnets, and configuration of route tables and network gateways.

You can easily customize the network configuration for your Amazon Virtual Private Cloud. For example, you can create a public-facing subnet for your web servers that has access to the Internet, and place your backend systems such as databases or application servers in a private-facing subnet with no Internet access. You can leverage multiple layers of security, including security groups and network access control lists, to help control access to Amazon EC2 instances in each subnet

For more information on Amazon VPC, please visit the link


<https://aws.amazon.com/vpc/>

The correct answer is: Amazon VPC

Feedback about this question and answer

Question 29

Not answered

 Mark for Review

The common use for IAM is to manage what?

Select 3 options.

Please select :

- ☐ A. Security Groups
- ☐ B. API Keys
- ☐ C. Multi-Factor Authentication
- ☐ D. Roles

Your answer is incorrect.

Answer – B,C and D

If you go on the IAM console, you will see the options on the left hand side. The Security groups are managed as part of the EC2 dashboard and not the IAM console.

+	Password
+	Multi-Factor Authentication (MFA)
+	Access Keys (Access Key ID and Secret Access Key)
+	CloudFront Key Pairs
+	X.509 Certificates
+	Account Identifiers

For more information on IAM, please refer to the below link

<https://aws.amazon.com/iam/>

The correct answer is: API Keys, Multi-Factor Authentication, Roles

Feedback about this question and answer

Question **30**

Not answered

Mark for Review

You have instances running on your VPC. You have both production and development based instances running in the VPC. You want to ensure that people who are responsible for the development instances don't have the access to work on the production instances to ensure better security. Using policies, which of the following would be the best way to accomplish this? Choose the correct answer from the options given below

Please select :

- ☐ A. Launch the test and production instances in separate VPC's and use VPC peering
- ☐ B. Create an IAM policy with a condition which allows access to only instances that are used for production or development
- ☐ C. Launch the test and production instances in different Availability Zones and use Multi Factor Authentication
- ☐ D. Define the tags on the test and production servers and add a condition to the IAM policy which allows access to specific tags

Your answer is incorrect.

Answer – D

You can easily add tags which define which instances are production and which are development instances and then ensure these tags are used when controlling access via an IAM policy.

For more information on tagging your resources, please refer to the below link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using_Tags.html

The correct answer is: Define the tags on the test and production servers and add a condition to the IAM policy which allows access to specific tags

Feedback about this question and answer

Question **31**

Not answered

Mark for Review

Your company is concerned with EBS volume backup on Amazon EC2 and wants to ensure they have proper backups and that the data is durable. What solution would you implement and why? Choose the correct answer from the options below

Please select :

- ☐ A. Configure Amazon Storage Gateway with EBS volumes as the data source and store the backups on premise through the storage gateway
- ☐ B. Write a cronjob on the server that compresses the data that needs to be backed up using gzip compression, then use AWS CLI to copy the data into an S3 bucket for durability
- ☐ C. Use a lifecycle policy to back up EBS volumes stored on Amazon S3 for durability
- ☐ D. Write a cronjob that uses the AWS CLI to take a snapshot of production EBS volumes. The data is durable because EBS snapshots are stored on the Amazon S3 standard storage class

Your answer is incorrect.

Answer – D

You can take snapshots of EBS volumes and to automate the process you can use the CLI. The snapshots are automatically stored on S3 for durability.

For more information on EBS snapshots, please refer to the below link

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

The correct answer is: Write a cronjob that uses the AWS CLI to take a snapshot of production EBS volumes. The data is durable because EBS snapshots are stored on the Amazon S3 standard storage class

Feedback about this question and answer

Question **32**

Not answered

Mark for Review

You are a consultant tasked with migrating an on-premise application architecture to AWS. During your design process you have to give consideration to current on-premise security and determine which security attributes you are responsible for on AWS. Which of the following does AWS provide for you as part of the shared responsibility model? Choose the correct answer from the options given below

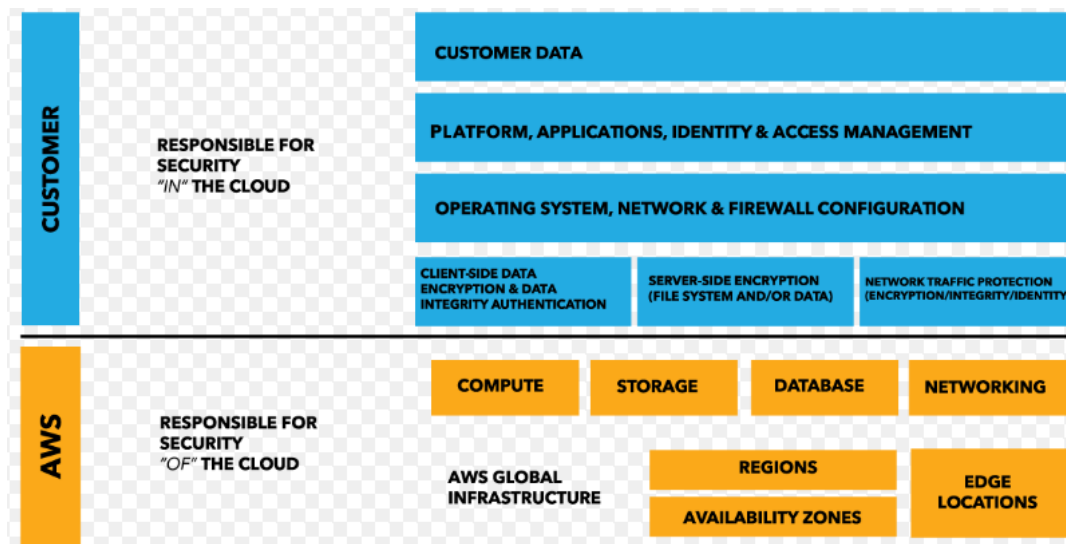
Please select :

- ☐ A. Customer Data
- ☐ B. Physical network infrastructure
- ☐ C. Instance security
- ☐ D. User access to the AWS environment

Your answer is incorrect.

Answer – B

As per the Shared responsibility model, the Physical network infrastructure is taken care by AWS. The below diagram clearly shows what has to be managed by customer and what is managed by AWS.



For more information on the Shared Responsibility model, please refer to the below link

<https://aws.amazon.com/compliance/shared-responsibility-model/>

The correct answer is: Physical network infrastructure

Feedback about this question and answer

Question **33**

Not answered

Mark for Review

Which of the following will occur when an EC2 instance in a VPC with an associated Elastic IP is stopped and started?

Select 2 options.

Please select :

- ☐ A. The underlying host for the instance can be changed
- ☐ B. The ENI (Elastic Network Interface) is detached
- ☐ C. All data on instance-store devices will be lost

- ☐ D. The Elastic IP will be dissociated from the instance

Your answer is incorrect.

Answer – A and C

If you have an EBS backed instance store , then the underlying host is changed when the instance is stopped and started.

And if you have instance store volumes, the data on the instance store devices will be lost.

For more information on the AMI types, please refer to the below link


<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>

The correct answer is: The underlying host for the instance can be changed, All data on instance-store devices will be lost

Feedback about this question and answer

Question **34**

Not answered

 Mark for Review

A customer is hosting their company website on a cluster of web servers that are behind a public-facing load balancer. The customer also uses Amazon Route 53 to manage their public DNS. How should the customer configure the DNS zone apex record to point to the load balancer?

Please select :

- ☐ A. Create an A record pointing to the IP address of the load balancer
- ☐ B. Create a CNAME record pointing to the load balancer DNS name.
- ☐ C. Create an alias for CNAME record to the load balancer DNS name.
- ☐ D. Create an alias for A record to the load balancer DNS name

Your answer is incorrect.

Answer – D

Alias resource record sets are virtual records that work like CNAME records. But they differ from CNAME records in that they are not visible to resolvers.

Resolvers only see the A record and the resulting IP address of the target record. As such, unlike CNAME records, alias resource record sets are available to configure a zone apex (also known as a root domain or naked domain) in a dynamic environment.

So when you create a hosted zone and having a pointer to the load balancer , you need to mark 'yes' for the Alias option as shown below. Then you can choose the Elastic Load balancer which you have defined in aws.

Create Record Set

Name: .demo.com.

Type:

Alias: ☒ Yes ☐ No

Alias Target:

Alias Hosted Zone ID: Z35SXDOTRQ7X7K

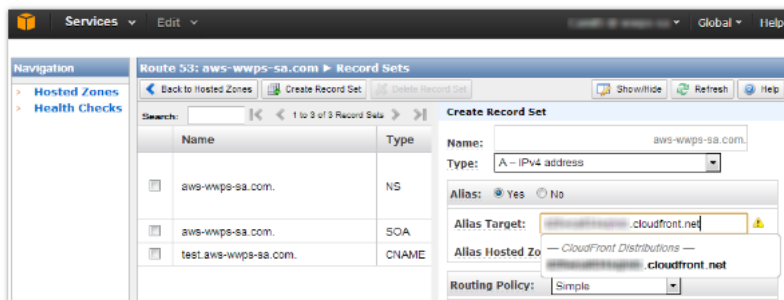
You can also type the domain name for the resource. Examples:

- CloudFront distribution domain name: d1111111abcdef8.cloudfront.net
- Elastic Beanstalk environment CNAME: example.elasticbeanstalk.com
- ELB load balancer DNS name: example-1.us-east-1.elb.amazonaws.com
- S3 website endpoint: example.s3-website-us-east-1.amazonaws.com
- Resource record set in this hosted zone: www.example.com

[Learn More](#)

Step 4: Configure a New Amazon Route 53 Alias Resource Record Set

1. Using your standard AWS account from the previous step, sign in to the [Amazon Route 53 console](#).
2. Create a new alias resource record set for your root domain name. For **Alias**, choose **Yes**. From the **Alias Target** drop-down list, select the CloudFront distribution name you created earlier.



For more information on the zone apex, please visit the link

- <http://docs.aws.amazon.com/govcloud-us/latest/UserGuide/setting-up-route53-zoneapex-elb.html>

The correct answer is: Create an alias for A record to the load balancer DNS name

Feedback about this question and answer

Question 35

Not answered

Mark for Review

To maintain compliance with HIPPA laws, all data being backed up or stored on Amazon S3 needs to be encrypted at rest. What is the best method for encryption for your data, assuming S3 is being used for storing the healthcare-related data? Choose the 2 correct answers from the options given below

Please select :

- ☐ A. Enable SSE on an S3 bucket to make use of AES-256 encryption
- ☐ B. Store the data in encrypted EBS snapshots
- ☐ C. Encrypt the data locally using your own encryption keys, then copy the data to Amazon S3 over HTTPS endpoints
- ☐ D. Store the data on EBS volumes with encryption enabled instead of using Amazon S3

Your answer is incorrect.

Answer – A & C

Data protection refers to protecting data while in-transit (as it travels to and from Amazon S3) and at rest (while it is stored on disks in Amazon S3 data centers). You can protect data in transit by using SSL or by using client-side encryption. You have the following options of protecting data at rest in Amazon S3.

- Use Server-Side Encryption – You request Amazon S3 to encrypt your object before saving it on disks in its data centers and decrypt it when you download the objects.
- Use Client-Side Encryption – You can encrypt data client-side and upload the encrypted data to Amazon S3. In this case, you manage the encryption process, the encryption keys, and related tools.

Protecting Data Using Encryption

Topics

- [Protecting Data Using Server-Side Encryption](#)
- [Protecting Data Using Client-Side Encryption](#)

Data protection refers to protecting data while in-transit (as it travels to and from Amazon S3) and at rest (while it is stored on disks in Amazon S3 data centers). You can protect data in transit by using SSL or by using client-side encryption. You have the following options of protecting data at rest in Amazon S3.

- **Use Server-Side Encryption** – You request Amazon S3 to encrypt your object before saving it on disks in its data centers and decrypt it when you download the objects.
- **Use Client-Side Encryption** – You can encrypt data client-side and upload the encrypted data to Amazon S3. In this case, you manage the encryption process, the encryption keys, and related tools.

For more information on S3 encryption, please refer to the below link


<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingEncryption.html>

The correct answer is: Enable SSE on an S3 bucket to make use of AES-256 encryption, Encrypt the data locally using your own encryption keys, then copy the data to Amazon S3 over HTTPS endpoints

Feedback about this question and answer

Question **36**

Not answered

 Mark for Review

There is a requirement by a company that does online credit card processing to have a secure application environment on AWS. They are trying to decide on whether to use KMS or CloudHSM. Which of the following statements is right when it comes to CloudHSM and KMS. Choose the correct answer from the options given below

Please select :

- ☐ A. It probably doesn't matter as they both do the same thing
- ☐ B. AWS CloudHSM does not support the processing, storage, and transmission of credit card data by a merchant or service provider, as it has not been validated as being compliant with Payment Card Industry (PCI) Data Security Standard (DSS); hence, you will need to use KMS
- ☐ C. KMS is probably adequate unless additional protection is necessary for some applications and data that are subject to strict contractual or regulatory requirements for managing cryptographic keys, then HSM should be used
- ☐ D. AWS CloudHSM should be always be used for any payment transactions

Your answer is incorrect.

Answer – C

AWS Key Management Service (KMS) is a managed service that makes it easy for you to create and control the encryption keys used to encrypt your data, and uses Hardware Security Modules (HSMs) to protect the security of your keys. This is sufficient if you the basic needs of managing keys for security.

For more information on KMS, please refer to the below link

<https://aws.amazon.com/kms/>

For a higher requirement on security one can use CloudHSM. The AWS CloudHSM service helps you meet corporate, contractual and regulatory compliance requirements for data security by using dedicated Hardware Security Module (HSM) appliances within the AWS cloud. With CloudHSM, you control the encryption keys and cryptographic operations performed by the HSM

For more information on CloudHSM, please refer to the below link


<https://aws.amazon.com/cloudhsm/>

The correct answer is: KMS is probably adequate unless additional protection is necessary for some applications and data that are subject to strict contractual or regulatory requirements for managing cryptographic keys, then HSM should be used

Feedback about this question and answer

Question **37**

Not answered

 Mark for Review

You are building a system to distribute confidential training videos to employees. Using CloudFront, what method would be used to serve content that is stored in S3, but not publicly accessible from S3 directly? Choose the correct answer from the options given below

Please select :

- ☐ A. Create an Origin Access Identify (OAI) for CloudFront and grant access to the objects in your S3 bucket to that OAI
- ☐ B. Create an Identity and Access Management (IAM) user for CloudFront and grant access to the objects in your S3 bucket to that IAM user.
- ☐ C. Create a S3 bucket policy that lists the CloudFront distribution ID as the principal and the target bucket as the Amazon Resource Name (ARN)
- ☐ D. Add the CloudFront account security group

Your answer is incorrect.

Answer – A

You can optionally secure the content in your Amazon S3 bucket so users can access it through CloudFront but cannot access it directly by using Amazon S3 URLs. This prevents anyone from bypassing CloudFront and using the Amazon S3 URL to get content that you want to restrict access to. This step isn't required to use signed URLs, but we recommend it.

To require that users access your content through CloudFront URLs, you perform the following tasks:

- Create a special CloudFront user called an origin access identity.
- Give the origin access identity permission to read the objects in your bucket.
- Remove permission for anyone else to use Amazon S3 URLs to read the objects.

For more information on CloudHSM, please refer to the below link


<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-restricting-access-to-s3.html>

The correct answer is: Create an Origin Access Identify (OAI) for CloudFront and grant access to the objects in your S3 bucket to that OAI

Feedback about this question and answer

Question 38

Not answered

 Mark for Review

As part of your application architecture requirements, the company you are working for has requested the ability to run analytics against all combined log files from the Elastic Load Balancer. Which services are used together to collect logs and process log file analysis in an AWS environment? Choose the correct answer from the options given below

Please select :

- ☐ A. Amazon S3 for storing the ELB log files and EC2 for processing the log files in analysis
- ☐ B. Amazon DynamoDB to store the logs and EC2 for running custom log analysis scripts
- ☐ C. Amazon S3 for storing ELB log files and Amazon EMR for processing the log files in analysis
- ☐ D. Amazon EC2 for storing and processing the log files

Your answer is incorrect.

Answer – C

You can use Amazon EMR for processing the jobs

Amazon EMR provides a managed Hadoop framework that makes it easy, fast, and cost-effective to process vast amounts of data across dynamically scalable Amazon EC2 instances. You can also run other popular distributed frameworks such as Apache Spark, HBase, Presto, and Flink in Amazon EMR, and interact with data in other AWS data stores such as Amazon S3 and Amazon DynamoDB.

Amazon EMR securely and reliably handles a broad set of big data use cases, including log analysis, web indexing, data transformations (ETL), machine learning, financial analysis, scientific simulation, and bioinformatics.

For more information on Amazon EMR please refer to the below link


<https://aws.amazon.com/emr/>

The correct answer is: Amazon S3 for storing ELB log files and Amazon EMR for processing the log files in analysis

Feedback about this question and answer

Question 39

Not answered

 Mark for Review

You manage a web application on EC2 instances. Incoming requests are saved as messages in an SQS queue with the maximum message retention period. You return from a week and a half vacation to find that the part of the application which processes received requests is hung and the queue has grown to 1,000 messages. While you are working on resolving the application issue, you need to post an update to end users. What information should that update include? Choose the correct answer from the options below

Please select :

- ☐ A. An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that recently received requests will need to be re-submitted (as the SQS queue has reached the maximum of 1,000 messages, meaning recently submitted requests could not be saved as SQS messages)
- ☐ B. An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that requests greater than ten days old will need to be resubmitted (as they have surpassed the SQS queue's retention period)
- ☐ C. An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that all received requests will be processed at that time.
- ☐ D. An apology for the delay in processing requests , assurance that the application will be operational shortly, and a note that all requests made in the last ten days will need to be resubmitted (because the portion of the application that processes received requests was down)

Your answer is incorrect.

Answer – C

Since the queue is configured with the maximum retention period, messages in the queue can last for 14 days.

For more information on Amazon SQS please refer to the below link


<https://aws.amazon.com/sqs/faqs/>

The correct answer is: An apology for the delay in processing requests, assurance that the application will be operational shortly, and a note that all received requests will be processed at that time.

Feedback about this question and answer

Question 40

Not answered

 Mark for Review

Your company has moved a legacy application from an on-premise data center to the cloud. The legacy application requires a static IP address hard-coded into the backend, which prevents you from deploying the application with high availability and fault tolerance using the ELB. Which steps would you take to apply high availability and fault tolerance to this application?

Select 2 options.

Please select :

- ☐ A. Write a custom script that pings the health of the instance, and, if the instance stops responding, switches the elastic IP address to a standby instance
- ☐ B. Ensure that the instance it's using has an elastic IP address assigned to it
- ☐ C. Do not migrate the application to the cloud until it can be converted to work with the ELB and Auto Scaling
- ☐ D. Create an AMI of the instance and launch it using Auto Scaling which will deploy the instance again if it becomes unhealthy

Your answer is incorrect.

Answer – A and B

The best option is to configure an Elastic IP that can be switched between a primary and failover instance.

Here is a link on using Elastic IP for failover.


<https://aws.amazon.com/articles/2127188135977316>

The correct answer is: Write a custom script that pings the health of the instance, and, if the instance stops responding, switches the elastic IP address to a standby instance, Ensure that the instance it's using has an elastic IP address assigned to it

Feedback about this question and answer

Question 41

Not answered

 Mark for Review

As an IT administrator you have been requested to ensure you create a highly decouple application in AWS. Which of the following help you accomplish this goal? Choose the correct answer from the options below

Please select :

- ☐ A. An SQS queue to allow a second EC2 instance to process a failed instance's job
- ☐ B. An Elastic Load Balancer to send web traffic to healthy EC2 instances
- ☐ C. IAM user credentials on EC2 instances to grant permissions to modify an SQS queue
- ☐ D. An Auto Scaling group to recover from EC2 instance failures

Your answer is incorrect.

Answer – A

Amazon Simple Queue Service (SQS) is a fully-managed message queuing service for reliably communicating among distributed software components and microservices - at any scale. Building applications from individual components that each perform a discrete function improves scalability and reliability, and is best practice design for modern applications.

SQS is the best option for creating a decoupled application.

Amazon Simple Queue Service (SQS) is a fully-managed message queuing service for reliably communicating among distributed software components and microservices - at any scale. Building applications from individual components that each perform a discrete function improves scalability and reliability, and is best practice design for modern applications. SQS makes it simple and cost-effective to decouple and coordinate the components of a cloud application. Using SQS, you can send, store, and receive messages between software components at any volume, without losing messages or requiring other services to be always available. SQS *standard queues* offer maximum throughput, best-effort ordering, and at-least-once delivery. SQS *FIFO queues* are designed to guarantee that messages are processed exactly once, in the exact order that they are sent, with limited throughput. You can get started with SQS in a matter of minutes using the AWS console or SDK of your choice and just three simple commands. SQS lets you eliminate the complexity and overhead associated with managing and operating dedicated messaging software and infrastructure.

For more information on SQS, please refer to the below link

<https://aws.amazon.com/sqs/>

The correct answer is: An SQS queue to allow a second EC2 instance to process a failed instance's job

Feedback about this question and answer

Question 42

Not answered

Mark for Review

A company has resources hosted in AWS and on on-premise servers. You have been requested to create a de-coupled architecture for applications which make use of both types of resources? Which of the below options are valid?

Select 2 options.

Please select :

- ☐ A. You can leverage SWF to utilize both on-premises servers and EC2 instances for your decoupled application
- ☐ B. SQS is not a valid option to help you use on-premises servers and EC2 instances in the same application, as it cannot be polled by on-premises servers
- ☐ C. You can leverage SQS to utilize both on-premises servers and EC2 instances for your decoupled application
- ☐ D. SWF is not a valid option to help you use on-premises servers and EC2 instances in the same application, as on-premises servers cannot be used as activity task workers

Your answer is incorrect.

Answer – A and C

You can use both SWF and SQS to coordinate with EC2 instances and on-premise servers.

Amazon Simple Queue Service (SQS) is a fully-managed message queuing service for reliably communicating among distributed software components and microservices - at any scale. Building applications from individual components that each perform a discrete function improves scalability and reliability, and is best practice design for modern applications.

For more information on SQS, please refer to the below link

<https://aws.amazon.com/sqs/>

The Amazon Simple Workflow Service (Amazon SWF) makes it easy to build applications that coordinate work across distributed components. In Amazon SWF, a task represents a logical unit of work that is performed by a component of your application. Coordinating tasks across the application involves managing intertask dependencies, scheduling, and concurrency in accordance with the logical flow of the application. Amazon SWF gives you full control over implementing tasks and coordinating them without worrying about underlying complexities such as tracking their progress and maintaining their state.

For more information on SWF, please refer to the below link


<http://docs.aws.amazon.com/amazonswf/latest/developerguide/swf-welcome.html>

The correct answer is: You can leverage SWF to utilize both on-premises servers and EC2 instances for your decoupled application, You can leverage SQS to utilize both on-premises servers and EC2 instances for your decoupled application

Feedback about this question and answer

Question 43

Not answered

 Mark for Review

When reviewing the Auto Scaling events, it is noticed that an application is scaling up and down multiple times within the hour. What design change could you make to optimize cost while preserving elasticity? Choose the correct answer from the options below

Please select :

- ☐ A. Change the scale down CloudWatch metric to a higher threshold
- ☐ B. Increase the instance type in the launch configuration
- ☐ C. Increase the base number of Auto Scaling instances for the Auto Scaling group
- ☐ D. Add provisioned IOPS to the instances

Your answer is incorrect.

Answer – A

If the threshold for the scale down is too low then the instances will keep on scaling down rapidly. Hence it is best to keep on optimal threshold for your metrics defined for Cloudwatch.

For more information on scaling on demand, please refer to the below link


<http://docs.aws.amazon.com/autoscaling/latest/userguide/as-scale-based-on-demand.html>

The correct answer is: Change the scale down CloudWatch metric to a higher threshold

Feedback about this question and answer

Question 44

Not answered

 Mark for Review

You are working for a startup company that is building an application that receives large amounts of data. Unfortunately, current funding has left the start-up short on cash, cannot afford to purchase thousands of dollars of storage hardware, and has opted to use AWS. Which services would you implement in order to store a virtually unlimited amount of data without any effort to scale when demand unexpectedly increases? Choose the correct answer from the options below

Please select :

- ☐ A. Amazon S3, because it provides unlimited amounts of storage data, scales automatically, is highly available, and durable
- ☐ B. Amazon Glacier, to keep costs low for storage and scale infinitely
- ☐ C. Amazon Import/Export, because Amazon assists in migrating large amounts of data to Amazon S3
- ☐ D. Amazon EC2, because EBS volumes can scale to hold any amount of data and, when used with Auto Scaling, can be designed for fault tolerance and high availability

Your answer is incorrect.

Answer – A

The best option is to use S3 because you can host a large amount of data in S3 and is the best storage option provided by AWS.

For more information on S3, please refer to the below link


<http://docs.aws.amazon.com/AmazonS3/latest/dev/Welcome.html>

The correct answer is: Amazon S3, because it provides unlimited amounts of storage data, scales automatically, is highly available, and durable

Feedback about this question and answer

Question 45

Not answered

 Mark for Review

A customer is running a multi-tier web application farm in a virtual private cloud (VPC) that is not connected to their corporate network. They are connecting to the VPC over the Internet to manage all of their Amazon EC2 instances running in both the public and private subnets. They have only authorized the bastion-security-group with Microsoft Remote Desktop Protocol (RDP) access to the application instance security groups, but the company wants to further limit administrative access to all of the instances in the VPC. Which of the following Bastion deployment scenarios will meet this requirement?

Please select :

- ☐ A. Deploy a Windows Bastion host on the corporate network that has RDP access to all instances in the VPC.
- ☐ B. Deploy a Windows Bastion host with an Elastic IP address in the public subnet and allow SSH access to the bastion from anywhere.
- ☐ C. Deploy a Windows Bastion host with an Elastic IP address in the private subnet, and restrict RDP access to the bastion from only the corporate public IP addresses.
- ☐ D. Deploy a Windows Bastion host with an auto-assigned Public IP address in the public subnet, and allow RDP access to the bastion from only the corporate public IP addresses.

Your answer is incorrect.

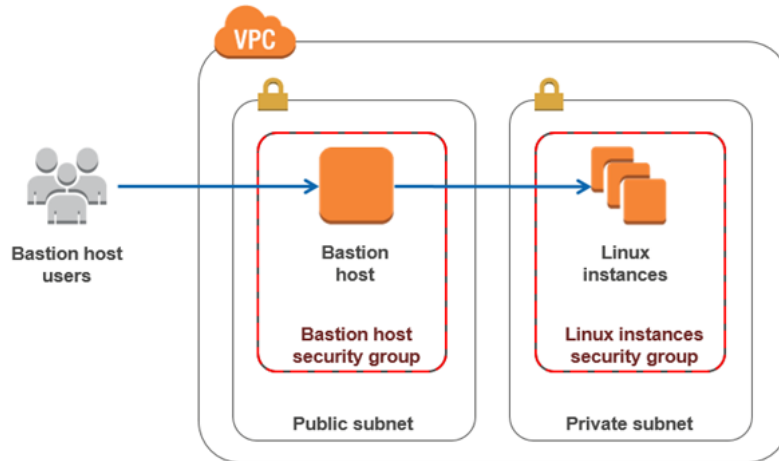
Answer – D

The bastion host should be in a public subnet with either a public or elastic IP and only allow RDP access from one IP from the corporate network.

A bastion host is a special purpose computer on a network specifically designed and configured to withstand attacks. The computer generally hosts a single application, for example a proxy server, and all other services are removed or limited to reduce the threat to the computer.

In AWS, A bastion host is kept on a public subnet. Users log on to the bastion host via SSH or RDP and then use that session to manage other hosts in the private subnets.

This is a security practice adopted by many organization to secure the assets in their private subnets.



The correct answer is: Deploy a Windows Bastion host with an auto-assigned Public IP address in the public subnet, and allow RDP access to the bastion from only the corporate public IP addresses.

Feedback about this question and answer

Question 46

Not answered

Mark for Review

Is it possible to promote a RDS Read replica to a primary state? Choose an answer from the options below.

Please select :

- ☐ A. True
- ☐ B. False

Your answer is incorrect.

Answer – A

Amazon RDS Read Replicas provide enhanced performance and durability for database (DB) instances. This replication feature makes it easy to elastically scale out beyond the capacity constraints of a single DB Instance for read-heavy database workloads. You can create one or more replicas of a given source DB Instance and serve high-volume application read traffic from multiple copies of your data, thereby increasing aggregate read throughput. Read replicas can also be promoted when needed to become standalone DB instances.

For more information on Read Replica, please refer to the below link

<https://aws.amazon.com/rds/details/read-replicas/>

The correct answer is: True

Feedback about this question and answer

Question 47

Not answered

Mark for Review

How does using ElastiCache help to improve database performance? Choose the correct answer from the options below

Please select :

- ☐ A. It can store petabytes of data
- ☐ B. It provides faster internet speeds
- ☐ C. It can store high-taxing queries
- ☐ D. It uses read replicas

Your answer is incorrect.

Answer – C

Amazon ElastiCache is a web service that makes it easy to deploy, operate, and scale an in-memory data store or cache in the cloud. The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases.

For more information on AWS Elastic Cache, please refer to the below link


<https://aws.amazon.com/elasticache/>

The correct answer is: It can store high-taxing queries

Feedback about this question and answer

Question **48**

Not answered

 Mark for Review

The Availability Zone that your RDS database instance is located in is suffering from outages, and you have lost access to the database. What could you have done to prevent losing access to your database (in the event of this type of failure) without any downtime? Choose the correct answer from the options below

Please select :

- ☐ A. Made a snapshot of the database
- ☐ B. Enabled multi-AZ failover
- ☐ C. Increased the database instance size
- ☐ D. Created a read replica

Your answer is incorrect.

Answer – B

The best option is to enable Multi-AZ for the database.

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure, and is engineered to be highly reliable. In case of an infrastructure failure, Amazon RDS performs an automatic failover to the standby (or to a read replica in the case of Amazon Aurora), so that you can resume database operations as soon as the failover is complete. Since the endpoint for your DB Instance remains the same after a failover, your application can resume database operation without the need for manual administrative intervention

For more information on AWS Multi-AZ, please refer to the below link


<https://aws.amazon.com/rds/details/multi-az/>

The correct answer is: Enabled multi-AZ failover

Feedback about this question and answer

Question **49**

Not answered

 Mark for Review

As an AWS administrator you are trying to convince a team to use RDS Read Replica's. What are two benefits of using read replicas?

Choose the 2 correct answers from the options below

Please select :

- ☐ A. Creates elasticity in RDS
- ☐ B. Allows both reads and writes
- ☐ C. Improves performance of the primary database by taking workload from it
- ☐ D. Automatic failover in the case of Availability Zone service failures

Your answer is incorrect.

Answer – A and C

By creating a read replica RDS, you have the facility to scale out the reads for your application, hence increasing the elasticity for your application. Also it can be used to reduce the load on the main database.

Read Replica's don't provide write operations, hence option B is wrong. And Multi-AZ is used for failover so Option D is wrong.

For more information on Read Replica, please refer to the below link


<https://aws.amazon.com/rds/details/read-replicas/>

The correct answer is: Creates elasticity in RDS, Improves performance of the primary database by taking workload from it

Feedback about this question and answer

Question 50

Not answered

 Mark for Review

What is the purpose of an SWF decision task? Choose the correct answer from the options below

Please select :

- ☐ A. It tells the worker to perform a function.
- ☐ B. It tells the decider the state of the work flow execution.
- ☐ C. It defines all the activities in the workflow.
- ☐ D. It represents a single task in the workflow.

Your answer is incorrect.

Answer – B

A decider is an implementation of the coordination logic of your workflow type that runs during the execution of your workflow. You can run multiple deciders for a single workflow type.

Because the execution state for a workflow execution is stored in its workflow history, deciders can be stateless. Amazon SWF maintains the workflow execution history and provides it to a decider with each decision task

For more information on Decider tasks, please refer to the below link


<http://docs.aws.amazon.com/amazonswf/latest/developerguide/swf-dg-dev-deciders.html>

The correct answer is: It tells the decider the state of the work flow execution.

Feedback about this question and answer

Question 51

Not answered

 Mark for Review

What is the best definition of an SQS message? Choose an answer from the options below

Please select :

- ☐ A. A mobile push notification
- ☐ B. A set of instructions stored in an SQS queue that can be up to 512KB in size
- ☐ C. A notification sent via SNS
- ☐ D. A set of instructions stored in an SQS queue that can be up to 256KB in size

Your answer is incorrect.

Answer – D

The maximum size of an SQS message as given in the AWS documentation is given below

Q: How do I configure the maximum message size for Amazon SQS?

To configure the maximum message size, use the console or the `SetQueueAttributes` method to set the `MaximumMessageSize` attribute. This attribute specifies the limit on bytes that an Amazon SQS message can contain. Set this limit to a value between 1,024 bytes (1 KB), and 262,144 bytes (256 KB). For more information, see [Using Amazon SQS Message Attributes](#) in the *Amazon SQS Developer Guide*.

For more information on SQS, please refer to the below link


<https://aws.amazon.com/sqs/faqs/>

The correct answer is: A set of instructions stored in an SQS queue that can be up to 256KB in size

Feedback about this question and answer

Question 52

Not answered

 Mark for Review

CloudTrail can log API calls from? Choose the correct answer from the options below

Please select :

- ☐ A. The command line
- ☐ B. The SDK
- ☐ C. The Console

☐ D. All of the above

Your answer is incorrect.

Answer – D

Cloudtrail can log all API calls which enter AWS.

AWS CloudTrail is a service that enables governance, compliance, operational auditing, and risk auditing of your AWS account. With CloudTrail, you can log, continuously monitor, and retain events related to API calls across your AWS infrastructure. CloudTrail provides a history of AWS API calls for your account, including API calls made through the AWS Management Console, AWS SDKs, command line tools, and other AWS services. This history simplifies security analysis, resource change tracking, and troubleshooting.

For more information on AWS Cloudtrail, please refer to the below link

<https://aws.amazon.com/cloudtrail/>

The correct answer is: All of the above

Feedback about this question and answer

Question 53

Not answered

 Mark for Review

What best describes Recovery Time Objective (RTO)? Choose the correct answer from the options below

Please select :

- ☐ A. The time it takes after a disruption to restore operations back to its regular service level.
- ☐ B. Minimal version of your production environment running on AWS.
- ☐ C. A full clone of your production environment.
- ☐ D. Acceptable amount of data loss measured in time.

Your answer is incorrect.

Answer – A

The recovery time objective (RTO) is the targeted duration of time and a service level within which a business process must be restored after a disaster (or disruption) in order to avoid unacceptable consequences associated with a break in business continuity

Please refer to the below link for more details


https://en.wikipedia.org/wiki/Recovery_time_objective

The correct answer is: The time it takes after a disruption to restore operations back to its regular service level.

Feedback about this question and answer

Question 54

Not answered

 Mark for Review

What AWS service, if used as part of your application's architecture, has an added benefit of helping to mitigate DDoS attacks from hitting your back-end instances? Choose the correct answer from the options below

Please select :

- ☐ A. CloudWatch
- ☐ B. CloudFront
- ☐ C. CloudTrail
- ☐ D. Kinesis

Your answer is incorrect.

Answer – B

The below snapshot from the aws documentation shows the best architecture practises for avoiding DDos attacks.

	AWS Edge Locations			AWS Regions		
	Amazon CloudFront with AWS WAF (BP1, BP2)	Amazon API Gateway (BP4)	Amazon Route 53 (BP3)	Elastic Load Balancing (BP6)	Amazon VPC (BP5)	Amazon EC2 with Auto Scaling (BP7)
Layer 3 (e.g., UDP reflection) attack mitigation	✓	✓	✓	✓	✓	
Layer 4 (e.g., SYN flood) attack mitigation	✓	✓	✓	✓		
Layer 6 (e.g., SSL) attack mitigation	✓	✓	N/A	✓		
Reduce attack surface	✓	✓	✓	✓	✓	
Scale to absorb application layer traffic	✓	✓	✓	✓		✓
Layer 7 (application layer) attack mitigation	✓	✓	✓			
Geographic isolation and dispersion of excess traffic and larger DDoS attacks	✓	✓	✓			

For best practises against DDos attacks , please visit the below link

https://d0.awsstatic.com/whitepapers/DDoS_White_Paper_June2015.pdf

The correct answer is: CloudFront

Feedback about this question and answer

Question 55

Not answered

Mark for Review

Perfect Forward Secrecy is used to offer SSL/TLS cipher suites for which two AWS services? Choose the correct answer from the options below

Please select :

- ☐ A. EC2 and S3
- ☐ B. CloudTrail and CloudWatch
- ☐ C. Cloudfront and Elastic Load Balancing
- ☐ D. Trusted advisor and GovCloud

Your answer is incorrect.

Answer – C

Its currently available for Cloudfront and ELB. Please find the below link for more details

<https://aws.amazon.com/about-aws/whats-new/2014/02/19/elastic-load-balancing-perfect-forward-secrecy-and-more-new-security-features/>

<https://aws.amazon.com/blogs/aws/cloudfront-ssl-ciphers-session-ocsp-pfs/>

The correct answer is: Cloudfront and Elastic Load Balancing

Feedback about this question and answer

Question 56

Not answered

Mark for Review

A customer has a single 3-TB volume on-premises that is used to hold a large repository of images and print layout files. This repository is growing at 500 GB a year and must be presented as a single logical volume. The customer is becoming increasingly constrained with their local storage capacity and wants an off-site backup of this data, while maintaining low-latency access to their frequently accessed data. Which AWS Storage Gateway configuration meets the customer requirements?

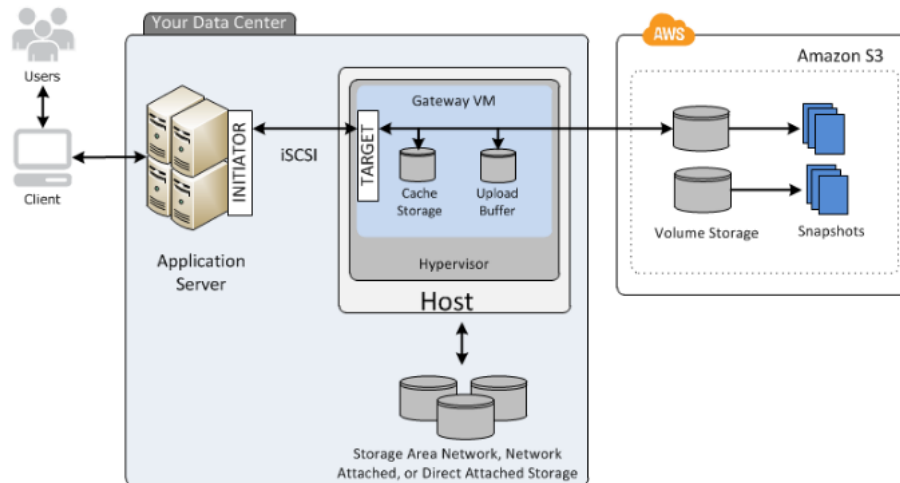
Please select :

- ☐ A. Gateway-Cached volumes with snapshots scheduled to Amazon S3
- ☐ B. Gateway-Stored volumes with snapshots scheduled to Amazon S3
- ☐ C. Gateway-Virtual Tape Library with snapshots to Amazon S3
- ☐ D. Gateway-Virtual Tape Library with snapshots to Amazon Glacier

Your answer is incorrect.

Answer - A

Gateway-cached volumes let you use Amazon Simple Storage Service (Amazon S3) as your primary data storage while retaining frequently accessed data locally in your storage gateway. Gateway-cached volumes minimize the need to scale your on-premises storage infrastructure, while still providing your applications with low-latency access to their frequently accessed data. You can create storage volumes up to 32 TiB in size and attach to them as iSCSI devices from your on-premises application servers. Your gateway stores data that you write to these volumes in Amazon S3 and retains recently read data in your on-premises storage gateway's cache and upload buffer storage.



For more information on Storage gateways, please visit the link

<http://docs.aws.amazon.com/storagegateway/latest/userguide/storage-gateway-cached-concepts.html>

The correct answer is: Gateway-Cached volumes with snapshots scheduled to Amazon S3

Feedback about this question and answer

Question 57

Not answered

Mark for Review

Which of the following best describes what the CloudHSM has to offer? Choose the correct answer from the options given below

Please select :

- ☐ A. An AWS service for generating API keys
- ☐ B. EBS Encryption method
- ☐ C. S3 encryption method
- ☐ D. A dedicated appliance that is used to store security keys

Your answer is incorrect.

Answer - D

The AWS CloudHSM service helps you meet corporate, contractual and regulatory compliance requirements for data security by using dedicated Hardware Security Module (HSM) appliances within the AWS cloud. With CloudHSM, you control the encryption keys and cryptographic operations performed by the HSM.

For more information on CloudHSM, please refer to the below link

<https://aws.amazon.com/cloudhsm/>

The correct answer is: A dedicated appliance that is used to store security keys

Feedback about this question and answer

Question 58

Not answered

Mark for Review

A company wants to launch EC2 instances on aws. For the linux instance, they want to ensure that the Perl language are installed automatically when the instance is launched. In which of the below configurations can you achieve what is required by the customer.

Please select :

- ☐ A. User data
- ☐ B. EC2Config service
- ☐ C. IAM roles

☐ D. AWS Config

Your answer is incorrect.

Answer – A

When you configure an instance during creation, you can add custom scripts to the User data section.

So in Step 3 of creating an instance, in the Advanced Details section, we can enter custom scripts in the User Data section. The below script installs Perl during the instance creation of the EC2 instance.

Step 3: Configure Instance Details

Additional charges apply.

Tenancy ⓘ

Shared - Run a shared hardware instance ▼

Additional charges will apply for dedicated tenancy.

▼ Network interfaces ⓘ

Device	Network Interface	Subnet	Primary IP	Secondary IP addresses
eth0	New network interface ▼	subnet-95ed8dd1 ▼	Auto-assign	Add IP

Add Device

▼ Advanced Details

User data ⓘ

☒ As text ☐ As file ☐ Input is already base64 encoded

```
yum install perl-Switch perl-DateTime perl-Sys-Syslog perl-LWP-Protocol-https  
-y
```

The correct answer is: User data

Feedback about this question and answer

Question 59

Not answered

Mark for Review

Which of the following are true regarding encrypted Amazon Elastic Block Store (EBS) volumes?

Choose 2 answers

Please select :

- ☐ A. Supported on all Amazon EBS volume types
- ☐ B. Snapshots are automatically encrypted
- ☐ C. Available to all instance types
- ☐ D. Existing volumes can be encrypted
- ☐ E. shared volumes can be encrypted

Your answer is incorrect.

Answer - A and B

Option C is wrong because there are some instance types that need to use IOPS storage and not EBS storage.

Option D is wrong because existing volumes cannot be encrypted.

Option E is wrong because Shared volumes cannot be encrypted.

EBS volumes can be applied to all of the below Volume types

	Solid-State Drives (SSD)		Hard disk Drives (HDD)	
Volume Type	General Purpose SSD (gp2)*	Provisioned IOPS SSD (io1)	Throughput Optimized HDD (st1)	Cold HDD (sc1)

For more information on EBS volume types, please visit the link

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>

So if we create a snapshot from a Volume that is encrypted.

Create Volume Actions

Filter by tags

Delete Volume
Attach Volume
Detach Volume
Force Detach Volume
Create Snapshot
Change Auto-Enable IO Setting
Add/Edit Tags

Created	Volume ID	Availability Zone	Alarm Status	Attachment Information	Monitoring	Volume Status	Encrypted
November 27	vol-0e8f9f718d2da2d9d	us-east-1a	None			✓ Okay	Encrypted
04 November 23	vol-0a1b2c3d4e5f6g7h8	us-east-1a	None	i-023260ed8c20507...		✓ Okay	Not Encrypted
04 November 17	vol-0i9j8k7l6m5n4o3p2	us-east-1a	None	i-09f54a79fc36966e...		✓ Okay	Not Encrypted

The Encrypted option will automatically become true for the snapshot.

Create Snapshot

Volume *i* vol-0e8f9f718d2da2d9d

Name *i* Demo

Description *i* Demo

Encrypted *i* Yes

Cancel Create

The correct answer is: Supported on all Amazon EBS volume types, Snapshots are automatically encrypted

Feedback about this question and answer

Question 60

Not answered

Mark for Review

A company is deploying a new two-tier web application in AWS. The company wants to store their most frequently used data so that the response time for the application is improved. Which AWS service provides the solution for the company's requirements?

Please select :

- ☐ A. MySQL Installed on two Amazon EC2 Instances in a single Availability Zone
- ☐ B. Amazon RDS for MySQL with Multi-AZ
- ☐ C. Amazon ElastiCache
- ☐ D. Amazon DynamoDB

Your answer is incorrect.

Answer – C

Amazon ElastiCache is a web service that makes it easy to deploy, operate, and scale an in-memory data store or cache in the cloud. The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases

- **Option A** is wrong because even if MySQL is installed on multiple systems, it will not help to serve the most recently used data.
- **Option B** is wrong because even a Multi-AZ option with an RDS will not suffice the requirement of the customer.
- **Option D** is wrong because this is a pure database option.

For more information on Elastic cache, please visit the link

- <https://aws.amazon.com/elasticache/>

The correct answer is: Amazon ElastiCache

Feedback about this question and answer

Finish review

QUIZ NAVIGATION

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15

16	17	18	19	20
21	22	23	24	25
26	27	28	29	30
31	32	33	34	35
36	37	38	39	40
41	42	43	44	45
46	47	48	49	50
51	52	53	54	55
56	57	58	59	60

Show all incorrect

Show all correct

Show one page at a time

Finish review

